

Optimal searches for a moving target on paths and grids

Angel Ivanov Raychev

September 2026

Research manuscript in preparation

Abstract

An invisible target starts at an unknown vertex of a finite graph. Each day a searcher inspects at most m vertices simultaneously; after a miss, the target must traverse exactly one edge. We study feasibility, optimal strategies, and the minimum number of days guaranteeing capture. We give exact formulas for every budget on paths and rectangular grids with at most five rows. On every odd-sided box, in arbitrary dimension, a proved isoperimetric nesting theorem reduces the unrestricted game to two counts, giving an exact algorithm polynomial in the number of rooms. For odd rectangles of width $2b + 1$, we solve the time recurrence explicitly for every budget $m \geq b^2 + 1$. For any fixed transverse Cartesian box with A rooms and fixed $m > A/2$, put $D = 2m - A$ and $g = \gcd(A, D)$. We prove $T_m(n + 2D/g) = T_m(n) + 4A/g$ for all sufficiently large n , of both parities, with an explicit polynomial sufficient threshold. The argument classifies efficient boundaries of arbitrary optimal searches and replaces their interior crossings by direct height descent. For a fixed deadline t , a separate spatial transfer proves that the minimum budget is eventually An/t plus a periodic correction. We also determine every budget on the $3 \times 3 \times 3$, $4 \times 4 \times 4$, and $3 \times 4 \times 4$ boxes. Five inspections take exactly $18n - 36$ days on every $3 \times 3 \times n$ box with $n \geq 3$. Paths, two-row grids, and the first two cubes have complete Lean proofs quantifying over actual target walks. The account reconstructs the author's October 2019–March 2020 path work and documents subsequent AI-assisted development. A complete explicit classification for arbitrary finite boxes remains open here.

Contents

1	Introduction and results	3
2	The game and its possible positions	6
3	Compressing entire strategies	7
4	Exact neighborhood profiles on odd rectangles	12
5	An exact isoperimetric order in every all-odd box	17
6	Paths: matching geometry and a probe-count potential	22
7	Two rows: an exact pair of counters	25
8	Three rows: a historical rank and the exact time	27
9	Four rows: counting productive inspections	31
10	Five rows	34
11	Seven rows with five inspections per day	47
12	Explicit formulas at large budgets on every odd rectangle	49
13	Minimum-budget searches on seven odd widths	57
14	Eventual affine periods in odd cylinders	61
15	Higher-dimensional boxes	69
16	Eventual affine periods for every fixed transverse box	82
17	A spatial transfer matrix for a fixed deadline	93
18	Formal verification, reproducibility, and open problems	95

1 Introduction and results

A princess moves through a building every night. Each morning a searcher may open several doors. The princess knows the search plan, may start anywhere, and must move to an adjacent room after each unsuccessful day. The problem is to guarantee capture as quickly as possible. Even on a single corridor the answer has a parity correction: dividing the amount of uncertainty by the apparent daily progress is not always enough.

This is a minimum-time version of the Hunters and Rabbit game. A path P_n has n vertices, and $P_{n_1} \square \cdots \square P_{n_d}$ denotes the Cartesian box with those side lengths. Write $T_m(G)$ for the optimal guaranteed number of inspection days, with value ∞ when no finite guarantee exists. Section 2 specifies the order of inspection and compulsory movement precisely.

1.1 The exact narrow-grid formulas

On a path, $T_m(P_n) = 1$ if $n \leq m$. With one inspection per day, $T_1(P_2) = 2$ and $T_1(P_n) = 2n - 4$ for $n \geq 3$. For $n > m \geq 2$,

$$T_m(P_n) = \left\lceil \frac{2n - 4}{2m - 1} \right\rceil + \varepsilon, \quad (1)$$

where $\varepsilon = 1$ exactly when $n \equiv m + 1 \pmod{2m - 1}$ and n, m are not both even. This theorem is the domain-corrected result of the author's original 2019–2020 work. We give a shorter construction and a matching-based lower proof, and formalize the complete statement in Lean.

For two rows and $n \geq 2$, one daily inspection is insufficient, and $m \geq 2n$ gives one day. In the remaining range put $N = n - 1$, $M = m - 1$. Then

$$T_m(P_2 \square P_n) = \left\lceil \frac{2N}{M} \right\rceil + \mathbf{1}_{\{M \text{ even and } N \equiv M/2 \pmod{M}\}}. \quad (2)$$

For three rows and $n \geq 2$, the one-day threshold is $3n$, the two-day threshold is $\lfloor 3n/2 \rfloor$, and one inspection is insufficient. For $2 \leq m < \lfloor 3n/2 \rfloor$, write $3n - 4 = q(2m - 3) + r$, $0 \leq r < 2m - 3$. The answer is

$$\begin{cases} 2q, & r = 0, \\ 2q + 1, & r > 0 \text{ and } m \geq r + 4 - \mathbf{1}_{\{n \text{ even, } q \text{ odd}\}}, \\ 2q + 2, & \text{otherwise.} \end{cases} \quad (3)$$

For four rows and $n \geq 4$, the one-day threshold is $4n$, the two-day threshold is $2n$, and budgets at most two are insufficient. For $3 \leq m < 2n$, write $2n - 2 = q(m - 2) + r$, $0 \leq r < m - 2$. The answer is

$$\begin{cases} 2q, & r = 0, \\ 2q + 1, & r = 1 \text{ or } (r \geq 2 \text{ and } m \geq 2r + 4), \\ 2q + 2, & \text{otherwise.} \end{cases} \quad (4)$$

Rotating a shorter four-row board reduces it to a preceding family. In particular the minimum feasible budgets give $T_2(P_3 \square P_n) = 6n - 8$ and $T_3(P_4 \square P_n) = 4n - 4$ in their stated nondegenerate ranges. For five rows and $n \geq 5$, the minimum feasible budget is three, and

$$T_3(P_5 \square P_n) = 10n - 20. \quad (5)$$

The even-length proof uses a finite symbolic boundary certificate valid for every length, together with a history-dependent potential. The odd-length proof follows from the explicit neighborhood

profiles. For every even $n \geq 6$, putting $h = 5n/2$ also gives

$$T_4(P_5 \square P_n) = 2 \left\lceil \frac{2h - 8}{3} \right\rceil.$$

Theorem 10.5 gives a closed quotient-and-remainder formula for every larger budget on these even-length boards. Theorem 10.7 supplies the closed answer for every odd length as well. In particular, four inspections per day take $2 \lceil (5n - 9)/3 \rceil + 21_{\{3|n\}}$ days when $n \geq 5$ is odd. Together with rotation of shorter boards, these give explicit formulas and attaining strategies for every five-row board and every budget. The wider odd-rectangle theorem extends the closed analysis to every width $2b + 1$ when $m \geq b^2 + 1$.




Figure 1: Optimal time at the smallest feasible daily budget for four narrow families and the five-row three-probe theorem. Each line is a proved formula, not a fit to search data.

1.2 The structural results

The common difficulty is temporal. A small possible-position set may have a particularly favorable next neighborhood, but that neighborhood need not contain an equally favorable set for the following day. We retain a history bit where necessary, and also prove a more general comparison theorem. A monotone, cardinality-preserving set map C with

$$N(C(A)) \subseteq C(N(A))$$

compresses a whole strategy without increasing any day's inspection budget. Idempotent maps supply an exact optimal-strategy normal form. This assertion is stronger than a one-step isoperimetric inequality.

We construct such operators from odd paths and square diagonal shifts. They give exact recurrences on smaller families of states in arbitrary dimension. More strongly, on every box whose

side lengths are all odd, we prove that the restricted simplicial orders minimize open neighborhoods and that neighborhoods of prefixes are prefixes. An induction using fiber compressions and a weighted exchange lemma reduces the unrestricted game to two cardinalities. The resulting exact algorithm is polynomial in the number of rooms and returns an optimal schedule. A recurrence remains distinct from a closed expression: the general box problem is not declared solved by the existence of an exponential search over all possible-position sets.

The three-dimensional example has a particularly concise answer:

m	0:4	5	6	7:8	9:11	12	13:26	27 or more
$T_m(P_3^3)$	∞	18	10	6	4	3	2	1.

Its complete lower and upper proofs are checked in Lean. A separate height construction proves $h(P_3 \square P_3 \square P_n) = 5$ for $n \geq 3$, where $h(G)$ is the minimum feasible daily budget. Moreover, Theorem 15.5 gives the exact time $18n - 36$ at that budget for every $n \geq 3$, of either parity.

The even cube P_4^3 is also completely classified, with physical Lean proofs at every budget. Eight inspections per day require exactly forty days; nine require twenty. Its eight-probe lower bound retains the shape of the possible-position set: using only its cardinality would falsely suggest thirty-two days. The complete tables for P_4^3 and $P_3 \square P_4 \square P_4$ appear in Theorem 15.4.

For a fixed bipartite Hamiltonian cross-section H of order A , every sufficiently long cylinder has hunting number $\lfloor A/2 \rfloor + 1$. For each fixed $m > A/2$, its optimal time is $2An/(2m - A) + O_{A,m}(1)$, with no restriction on side parity. For every transverse Cartesian box we prove more: the exact time has an effective eventual affine period for both longitudinal parities. Efficient searches must have one of finitely many translated frontiers; controlling the remaining days permits insertion and deletion of a physical interval in any optimal strategy. A separate fixed-deadline theorem applies to every finite cross-section: the smallest budget for capture within t days is An/t plus an eventually periodic correction. These are distinct parameter regimes, and neither substitutes for a uniform closed formula for all boxes, budgets, and deadlines.

1.3 Prior work and provenance

Britnell and Wildon [6] studied the princess puzzle with one daily inspection, including the path value $2n - 4$. Haslegrave studied the associated evasion game [9]; related moving- target search appears in Beluhov and Kolev [2]. Abramovskaya, Fomin, Golovach, and Pilipczuk [1] developed Hunters and Rabbit results, including the rectangular-grid feasibility threshold $\lfloor \min(a, b)/2 \rfloor + 1$. Bolkema and Groothuis proved isoperimetric nesting and the hunting number for hypercubes [5]. We use precise versions of these results; minimum-time claims do not follow from feasibility alone. Parity-restricted boundary minimization on the binary cube goes back to Körner and Wei [12]. Compression proofs and local-to-global principles have a substantial earlier literature, including Bezrukov and Serra [4]; their closed-neighborhood theorem is distinct from the open-neighborhood argument given here.

Dmitry Kamenetsky recorded conjectural two- and three-inspection path formulas in March 2018 [10, 11]. Equation (1) proves those special cases and extends them to arbitrary budgets. The author’s independent path work took place between October 2019 and March 2020 and was written in two Bulgarian conference manuscripts [14, 15]. The one-probe restriction omitted from a printed general formula is made explicit here. The original conference documents are preserved.

The time-budget inverse parameter is also considered in recent work [3]; we do not claim the general optimization question as new. Recontamination can matter substantially on general graphs

[7], so a presumed monotone search is never used as an unproved lower-bound assumption. Classical closed-neighborhood isoperimetry for Cartesian grids [13] supplies one of our higher-dimensional applications. Our use of that theorem is distinguished from the open-neighborhood argument for odd rectangles.

Remark 1.1 (Status of this manuscript). This is a research manuscript in preparation, not an arXiv submission. The ordinary theorems have internal independent proof reviews; exact Lean coverage is stated in Section 18. Attribution of the 2026 extensions is still undergoing a focused literature review. The manuscript makes no blanket claim that every derived ingredient is new to the literature.

2 The game and its possible positions

Let $G = (V, E)$ be a finite simple graph without isolated vertices, and let $m \geq 1$ be the daily inspection budget. Every connected board with at least two rooms satisfies this graph assumption; a one-room board takes one inspection and is treated separately. The target chooses an unknown initial vertex. On day $t = 1, 2, \dots$, the searcher inspects a set $S_t \subseteq V$ with $|S_t| \leq m$. If the target occupies a vertex of S_t , it is caught. Otherwise it must move along exactly one edge before the next day. Inspections within a day are simultaneous. The searcher receives no information about an unsuccessful inspection other than the miss.

Write $T_m(G)$ for the minimum number of days in which capture can be guaranteed, and put $T_m(G) = \infty$ if there is no finite guarantee. The associated feasibility threshold is

$$h(G) = \min\{m \geq 1 : T_m(G) < \infty\}.$$

We consider deterministic guarantees. Before capture, every observation is a miss. Thus an adaptive strategy has only one continuing observation history and is represented by a sequence $(S_t)_{t \geq 1}$. This sequence formulation does not restrict the searcher's guaranteed performance.

For $A \subseteq V$, its open neighborhood is

$$N(A) = \{v \in V : \text{some } u \in A \text{ satisfies } uv \in E\}.$$

Define the possible-position sets immediately before inspection by

$$B_1 = V, \quad B_{t+1} = N(B_t \setminus S_t).$$

The following elementary interpretation is used in every subsequent argument.

Lemma 2.1. *For each $t \geq 1$, a vertex v belongs to B_t if and only if there is a walk $v_1, \dots, v_t = v$ in G such that $v_i \notin S_i$ for every $1 \leq i < t$. Consequently capture is guaranteed by day t exactly when $B_t \subseteq S_t$.*

Proof. The assertion about walks follows by induction. At $t = 1$ every vertex is an admissible initial position. For the inductive step, $v \in B_{t+1}$ means that some neighbor $u \in B_t$ avoids S_t . Append v to a walk ending at u supplied by the induction hypothesis. Conversely, the penultimate vertex of any such walk belongs to $B_t \setminus S_t$, so its final vertex belongs to B_{t+1} .

A walk ending outside S_t avoids all inspections through day t , whereas every surviving walk is caught that day if its endpoint lies in S_t . This is exactly the stated inclusion. $\square$

Since G has no isolated vertices, a nonempty set has a nonempty open neighborhood. Hence the capture condition is also equivalent to $B_{t+1} = \emptyset$. We use this equivalent form only under that hypothesis. A board with a single room is handled directly: one inspection on the first day suffices. Thus an inability to make a move is never silently treated as capture.

Two monotonicity observations will be useful. Starting from fewer possible positions cannot make a fixed inspection sequence less successful, and adding inspected vertices cannot make it less successful. Both follow by induction from monotonicity of N and set difference in their respective arguments. In particular, probes outside the current possible-position set may be wasted, but they cause no problem for the model or the lower bounds.

2.1 Two initial parity classes

Suppose G is bipartite, with vertex classes V_0, V_1 . Keep track of the two possible initial colors separately:

$$R_1^s = V_s, \quad R_{t+1}^s = N(R_t^s \setminus S_t), \quad s \in \{0, 1\}.$$

Then

$$B_t = R_t^0 \cup R_t^1, \quad R_t^s \subseteq V_{s+t-1 \pmod{2}}.$$

Indeed, neighborhood and deletion distribute over unions, and every edge reverses color. The two sets therefore stay disjoint while exchanging their current colors each night. Allocations to the two initial classes sum to at most the day's inspection budget. We will state explicitly whether a pair of counts is indexed by initial color or by current color; the latter convention swaps the two counts after every move.

This separates the bookkeeping common to the path, ladder, and grid arguments. The work specific to a graph is to determine how small a neighborhood can be after a given number of inspections, and when a sequence of such bounds can actually be attained.

3 Compressing entire strategies

An isoperimetric inequality bounds one neighborhood. To preserve an optimal search, the replacement sets must also respect containment and remain compatible after successive moves. The following criterion provides that stronger conclusion.

Definition 3.1. A *strategy compression* on G is a map $C : 2^V \rightarrow 2^V$ satisfying

$$A \subseteq B \implies C(A) \subseteq C(B), \tag{6}$$

$$|C(A)| = |A|, \tag{7}$$

$$N(C(A)) \subseteq C(N(A)). \tag{8}$$

Idempotence is an additional property, not part of this definition.

Theorem 3.2 (Strategy comparison and a fixed-state normal form). *Let C be a strategy compression. Every successful strategy from A with prescribed daily budgets $m_1, \dots, m_T$ has a successful strategy from $C(A)$ with the same budgets. If C is idempotent, every starting set fixed by C has an optimal strategy whose possible-position sets and post-inspection survivor sets are all fixed by C .*

Proof. Let A_t be the original possible-position set before day t , and R_t its survivors, so $A_{t+1} = N(R_t)$ and $|A_t| - |R_t| \leq m_t$. Maintain a new actual possible-position set $B_t \subseteq C(A_t)$, initially $B_1 = C(A_1)$. Inspect

$$Q_t = B_t \setminus C(R_t).$$

Monotonicity and equal cardinalities give

$$|Q_t| \leq |C(A_t) \setminus C(R_t)| = |A_t| - |R_t| \leq m_t.$$

The new survivors are $B_t \cap C(R_t)$, and

$$B_{t+1} = N(B_t \cap C(R_t)) \subseteq N(C(R_t)) \subseteq C(N(R_t)) = C(A_{t+1}).$$

When the original survivors are empty, $C(R_t) = \emptyset$, so the new strategy also captures every remaining possibility.

Fixed sets are closed under intersections: if X, Y are fixed, monotonicity gives $C(X \cap Y) \subseteq X \cap Y$, and cardinality forces equality. They are closed under neighborhoods for the same reason: $N(X) \subseteq C(N(X))$ when X is fixed. If C is idempotent, $C(R_t)$ is fixed. Starting the preceding construction from a fixed B_1 , every actual survivor $B_t \cap C(R_t)$ and every next possible-position set is therefore fixed. Restricted strategies are actual strategies, so their optimum equals the unrestricted optimum. $\square$

3.1 Products and simultaneous compression

Lemma 3.3 (Product lift). *A strategy compression on a graph A lifts to $A \square H$ by applying it separately in every H -indexed A -fiber. Idempotence is preserved.*

Proof. Write S_h for the fiber at h . Its neighborhood is

$$N(S)_h = N_A(S_h) \cup \bigcup_{u \sim h} S_u.$$

For the lifted operator $\hat{C}$, the corresponding neighborhood is

$$\begin{aligned} N(\hat{C}(S))_h &= N_A(C(S_h)) \cup \bigcup_{u \sim h} C(S_u) \\ &\subseteq C(N_A(S_h)) \cup \bigcup_{u \sim h} C(S_u) \\ &\subseteq C\left(N_A(S_h) \cup \bigcup_{u \sim h} S_u\right). \end{aligned}$$

The last inclusion uses monotonicity separately on each input subset of the union. Cardinality, monotonicity and idempotence hold fiber by fiber. $\square$

Compositions of strategy compressions are strategy compressions, since

$$N(C(D(S))) \subseteq C(N(D(S))) \subseteq C(D(N(S))).$$

Suppose finitely many such operators strictly decrease a common nonnegative integer potential whenever they change a set. If that potential has a uniform upper bound B , cycling through all operators $B+1$ times reaches a common fixed point for every input. A cycle with a change decreases the potential, and a cycle without a change fixes every operator. The resulting map is a composition with a *uniform* number of factors, hence is monotone and cardinality preserving; it is also idempotent. This avoids assuming that an input-dependent stopping rule preserves monotonicity.

3.2 Odd path fibers

Lemma 3.4 (Odd-path compression). *On P_{2q+1} , $q \geq 1$, number vertices $0, \dots, 2q$. Replace the selected vertices of each parity by the same number of vertices at the beginning of*

$$0, 2, \dots, 2q \quad \text{or} \quad 1, 3, \dots, 2q - 1.$$

This is an idempotent, parity-preserving strategy compression.

Proof. Monotonicity, cardinality and idempotence follow from the definition. An even-parity k -set has at least $\min(k, q)$ odd neighbors. If $k \leq q$, omit an unselected even vertex $2j$ and match $2i$ to $2i + 1$ for $i < j$, and to $2i - 1$ for $i > j$. These matches are distinct and include every selected even vertex. If $k = q + 1$, all q odd vertices are neighbors. A nonempty odd-parity k -set has at least $k + 1$ even neighbors: its consecutive blocks in the spacing-two order each have one more neighbor than their size, and different blocks have disjoint neighbor intervals. The two prefixes attain these bounds, and their neighborhoods are prefixes of the opposite parity. Consequently each part of $N(C(S))$ is contained in the corresponding same-cardinality prefix of $N(S)$, proving (8) even when S contains both colors. $\square$

For several odd coordinates, combine their lifted operators. Every changing operation decreases the sum of coordinates, so simultaneous compression has the normal form of Theorem 3.2. Its fixed sets satisfy $v \in S \implies v - 2e_i \in S$ whenever the latter vertex belongs to the box. These are lower ideals separately in each coordinate-parity pattern. They need not be lower ideals in the whole checkerboard color: for example, $\{(1, 1)\}$ in a 3×3 square is fixed by the two path operators but does not contain $(0, 0)$.

3.3 Exact recurrences on odd cylinders

Let $G = H \square P_{2q+1}$, where H is any finite graph and $q \geq 1$. The product lift and Theorem 3.2 give an exact normal form from the full board. It uses two vectors indexed by $V(H)$: a_v^0 counts even path coordinates in fiber v , and a_v^1 counts odd path coordinates. Their ranges are

$$0 \leq a_v^0 \leq q + 1, \quad 0 \leq a_v^1 \leq q.$$

Thus they refer to the parity of the *path coordinate*, and do not require H to be bipartite. Define

$$\phi_0(k) = \min(k, q), \quad \phi_1(k) = \begin{cases} 0, & k = 0, \\ k + 1, & k > 0. \end{cases}$$

For survivor vectors $b^0 \leq a^0$, $b^1 \leq a^1$, the exact next state $F(b)$ is

$$F(b)_v^0 = \max \left(\phi_1(b_v^1), \max_{u \sim v} b_u^0 \right), \quad (9)$$

$$F(b)_v^1 = \max \left(\phi_0(b_v^0), \max_{u \sim v} b_u^1 \right), \quad (10)$$

with an empty maximum interpreted as zero. Path edges supply the first term, and H -edges the second. All contributions in a fiber are prefixes of the same parity, so their union has exactly the maximum length. The required number of inspections is

$$d(a, b) = \sum_{v \in V(H)} (a_v^0 - b_v^0 + a_v^1 - b_v^1).$$

They are realized by inspecting the removed suffix of each prefix.

For a fixed budget m , form an edge $a \rightarrow F(b)$ whenever $b \leq a$ and $d(a, b) \leq m$. Then $T_m(a)$ is exactly the shortest-path distance to zero in this state graph, or infinity if zero is unreachable. Equivalently,

$$T_m(a) = 1 + \min_{\substack{b \leq a \\ d(a, b) \leq m}} T_m(F(b)), \quad T_m(0) = 0, \quad (11)$$

where the shortest-path interpretation specifies the solution of the possibly cyclic equations. For all budgets simultaneously, let $W_t(a)$ be the least constant budget that wins within t days. Then

$$W_{t+1}(a) = \min_{b \leq a} \max\{d(a, b), W_t(F(b))\}, \quad (12)$$

with $W_0(0) = 0$, $W_0(a) = \infty$ for $a \neq 0$. Backpointers in either recurrence give actual inspection sets.

If $A = |V(H)|$, the number of states is exactly

$$K = ((q+1)(q+2))^A \leq (2q+2)^{2A}.$$

There are at most K^2 state-survivor pairs. Explicit construction of their transitions takes $O((A + |E(H)|)K^2)$ elementary operations, apart from integer bit costs. For fixed H , this is polynomial in the odd longitudinal length. Any successful path can have cycles removed, so a finite optimum is at most $K - 1$. This proves an exact algorithm, not a uniform polynomial bound in dimension or a closed time formula.

When H is bipartite, one may instead index vectors by global color. For a coloring χ , put $\varepsilon_p(v) = p - \chi(v) \pmod{2}$. A single current color has transition

$$[\Phi_p(b)]_v = \max\left(\phi_{\varepsilon_p(v)}(b_v), \max_{u \sim v} b_u\right).$$

The two current-color vectors then become $(\Phi_1(b^1), \Phi_0(b^0))$. This convention is useful when comparing the cylinder recurrence with the two-color grid arguments below.

3.4 Square fibers and equal-sided boxes

Abramovskaya, Fomin, Golovach and Pilipczuk [1, Section 3.1, Lemma 2] prove that diagonal shifts do not increase the neighborhood size of a one-color set in a square. The following fiber calculation strengthens the scalar comparison to (8) and applies it to both colors. That stronger statement is what allows whole strategies to be compressed.

Lemma 3.5 (Square diagonal compression). *On $\{0, \dots, n-1\}^2$, compress each line $x+y=k$ toward increasing priority for smaller y , keeping its selected cardinality. This is an idempotent, parity-preserving strategy compression. The same is true on lines $x-y=k$, again toward smaller y .*

Proof. Only the neighborhood inclusion needs proof. Consecutive square diagonals have lengths differing by one. Index both from the common increasing- y end. From a diagonal of length L to one of length $L+1$, a selected index j has neighbors $j, j+1$. A nonempty a -set therefore has at least $a+1$ neighbors; a prefix attains this bound with a prefix. From length L to length $L-1$, the indices are $j-1, j$, clipped to the valid range. A proper nonempty a -set has at least a neighbors,

and the full set has $L - 1$. Indeed, if it has c consecutive blocks and occupies e endpoints, its neighbor count is $a + c - e$; for a proper set $c \geq e$. Prefixes again attain the bounds with prefixes.

Fix an output diagonal. Its two input diagonals, after compression, contribute prefixes from the same end. Their union is the longer prefix. Each contributing length is no larger than its original contribution, hence no larger than the original union's size. The compressed union is therefore contained in the prefix selected by compressing that original neighborhood. This argument works on either parity and on their union. Reflecting x proves the other diagonal direction. $\square$

Theorem 3.6 (Normal form on equal-sided boxes). *On P_n^d , $n \geq 2$, an optimal strategy from the full board can be chosen so that its possible-position sets and survivors are lower ideals for the legal moves*

$$v \mapsto v + e_i - e_j, \quad v \mapsto v - e_i - e_j, \quad i < j.$$

The conclusion respects arbitrary prescribed daily budgets. The same operators are available within every group of equal-length coordinates in an arbitrary Cartesian box. Odd-coordinate path compressions can be included simultaneously.

Proof. Lift both square operators to every pair $i < j$ of equal coordinates. They compress toward smaller x_j on lines $x_i + x_j = k$ and $x_i - x_j = k$. Use the potential

$$\mathcal{E}(S) = \sum_{v \in S} \sum_{i=1}^d i v_i.$$

A nontrivial plus-diagonal transfer by $\delta > 0$ decreases it by $(j - i)\delta$, and a minus-diagonal transfer by $(i + j)\delta$. An odd-coordinate prefix transfer also strictly decreases it. The potential is uniformly bounded; on P_n^d the bound $n^d(n - 1)d(d + 1)/2$ suffices. Uniform cycling therefore gives an idempotent strategy compression. Its fixed sets are precisely the ideals for the displayed local moves, with the additional $-2e_i$ conditions when odd-coordinate operators are included. Apply Theorem 3.2. $\square$

In dimension two the two-diagonal fixed sets are the downward pyramidal sets of the square isoperimetric argument. They differ from ordinary checkerboard corner ideals. In any dimension they give an exact finite recurrence over their fixed family, since both neighborhoods and intersections remain fixed. The theorem does not bound the number of these ideals by a polynomial in d or in n .

Remark 3.7 (Why even path factors require another idea). Every parity-preserving strategy compression on an even path is the identity. Each endpoint is the unique degree-one vertex in its color. Applying (8) to its singleton forces that singleton to be fixed. Fixed sets are closed under neighborhoods, so every N^t of either endpoint is fixed. These are the spacing-two prefixes from the two ends. Intersections of suitable prefixes isolate every singleton; fixed intersections follow from monotonicity and cardinality, without idempotence. Thus all singletons, and consequently all sets, are fixed. This rules out a nontrivial even-path factor reduction within this specific comparison framework. It does not rule out richer state labels or different comparison theorems.

4 Exact neighborhood profiles on odd rectangles

Let $G = P_{2a+1} \square P_{2b+1}$, where $a \geq b \geq 1$, with coordinates (x, y) ; thus x is the long coordinate. Write V_0, V_1 for its even and odd checkerboard classes, and

$$E = |V_0| = \frac{|G| + 1}{2}, \quad M = |V_1| = E - 1.$$

Define the two open-neighborhood profiles by

$$g_p(k) = \min\{|N(S)| : S \subseteq V_p, |S| = k\}.$$

It is convenient to use the integer functions

$$R(k) = \lceil \sqrt{k} \rceil, \quad Q(k) = \min\{r \geq 1 : k \leq r(r-1)\}.$$

In particular $R(0) = 0$ and $Q(0) = 1$.

Theorem 4.1 (Odd-rectangle profiles and compatible orders). *Both profiles have $g_p(0) = 0$. For nonempty sets,*

$$g_0(k) = k + \min\{R(k), b, R(E - k) - 1\}, \quad 1 \leq k \leq E, \quad (13)$$

$$g_1(k) = k + \min\{Q(k), b + 1, Q(M - k)\}, \quad 1 \leq k \leq M. \quad (14)$$

Order each parity by increasing $(x + y, x)$, and let $I_p(k)$ be its prefix of size k . Every prefix attains the corresponding minimum, and

$$N(I_p(k)) = I_{1-p}(g_p(k)).$$

The order fills the short direction first within a diagonal. Its direction is significant on an unequal rectangle. For example, the opposite tie order on 7×5 gives the five-element even prefix $\{(0, 0), (0, 2), (1, 1), (2, 0), (4, 0)\}$, with eight neighbors; the set $\{(0, 0), (0, 2), (0, 4), (1, 1), (1, 3)\}$ has seven.

4.1 Quadrant bounds

Apply the odd-coordinate compressions of Lemma 3.4 in both directions. Their common fixed sets are lower closed under decreasing either coordinate by two, and compression does not increase the neighborhood size. We may therefore assume that each coordinate-parity pattern of S is a lower Ferrers diagram. Let e_i and o_i be the selected counts in rows $y = 2i$ and $y = 2i + 1$, respectively. These are separately nonincreasing sequences. For source color 0 they count even- x and odd- x vertices, whereas for source color 1 they count odd- x and even- x vertices.


Lemma 4.2 (Quadrant area and row bounds). *Regard a finite set with this Ferrers property as a subset of the infinite nonnegative quadrant. Put $\delta = |N(S)| - |S|$. If S has color 0, then*

$$e_i > 0 \implies e_i + i \leq \delta, \quad o_i > 0 \implies o_i + i + 1 \leq \delta, \quad |S| \leq \delta^2.$$

If S has color 1, then

$$e_i > 0 \implies e_i + i + 1 \leq \delta, \quad o_i > 0 \implies o_i + i + 1 \leq \delta, \quad |S| \leq \delta(\delta - 1).$$

A nonempty set has surplus at least one in color 0 and at least two in color 1.



Filled circles: source set. Open circles: its possible neighbors.

Figure 2: Minimizing even prefixes on a 9×7 rectangle. The three regimes are controlled by a corner, the short width, and the complement.

Proof. The sequences eventually vanish. A nonempty odd- x prefix of length t has $t + 1$ even- x neighbors in the quadrant; an even- x prefix has t odd- x neighbors. Contributions from neighboring rows are also prefixes, so their unions are given by maxima. Subtracting the source row sizes therefore yields, in color 0,

$$\delta = (o_0 - e_0)_+ + \sum_{i \geq 0} \left[\max\{e_i - o_i, \mathbf{1}_{\{o_i > 0\}}\} + (o_i - e_{i+1})_+ \right]. \quad (15)$$

For color 1 the identity is

$$\delta = \max\{\mathbf{1}_{\{e_0 > 0\}}, o_0 - e_0\} + \sum_{i \geq 0} \left[(e_i - o_i)_+ + \max\{o_i - e_{i+1}, \mathbf{1}_{\{e_{i+1} > 0\}}\} \right]. \quad (16)$$

All terms displayed are nonnegative.

Suppose first that S has color 0 and $e_i > 0$. In (15), each of the first i paired summands is at least one. If its odd count is positive, the first term is at least one; otherwise its even count is positive. The tail from i onward is at least the telescoping sum

$$\sum_{j \geq i} ((e_j - o_j) + (o_j - e_{j+1})) = e_i.$$

Thus $\delta \geq i + e_i$. If $o_i > 0$, every preceding pair contributes at least one, as does the first term of pair i . The second term of pair i and the subsequent pairs have sum at least o_i , by the same telescoping argument. This gives $\delta \geq i + 1 + o_i$.

For color 1 and $e_i > 0$, the initial term of (16) is at least one, and the first i pairs each contribute at least one because $e_{j+1} > 0$ for $j < i$. The remaining tail is at least e_i . If $o_i > 0$, the initial term is again at least one. Every preceding pair contributes at least one: either $e_{j+1} > 0$, or the positive o_j appears in its second term. The tail after the first term of pair i is at least o_i . These are the two claimed bounds for color 1.

Finally sum the row bounds. In color 0 the maximum possible total is

$$\sum_{i \geq 0} \max(\delta - i, 0) + \sum_{i \geq 0} \max(\delta - i - 1, 0) = \delta^2.$$

In color 1 both sums have the second form, giving $\delta(\delta - 1)$. The row inequalities also give the stated positive surplus bounds for nonempty sets. $\square$

4.2 The two far edges

Lemma 4.3 (Corner, strip, or complement). *Let $S \subseteq V_p$ have the Ferrers property. Put $\delta = |N_G(S)| - |S|$, and let F and Z count its vertices on the far edges $x = 2a$ and $y = 2b$, respectively. If exactly one of F, Z is positive, then $\delta \geq b + p$. If neither is positive, the quadrant bounds apply to S without change. If both are positive, $N_G(S)$ covers both near edges in color $1 - p$, and the set*

$$V_{1-p} \setminus N_G(S),$$

reflected in both coordinate midlines, has the Ferrers property and no far-edge vertices.

Proof. Viewing the same S in the quadrant adds exactly one neighbor beyond the rectangle for every selected far-edge vertex. These $F + Z$ vertices are distinct, even when the far corner is selected. Hence

$$\delta_{\text{quadrant}} = \delta + F + Z. \tag{17}$$

If $F = 0$ and $Z > 0$, the last even row has $e_b = Z$. The row bound in Lemma 4.2 gives $\delta + Z \geq Z + b + p$. If $F > 0$ and $Z = 0$, swap the axes to get the stronger bound $\delta \geq a + p$. If $F = Z = 0$, there is no clipping.

Suppose now that both are positive. In color 0, the Ferrers property forces every even- x vertex of row $y = 0$ and every even- y vertex of column $x = 0$ into S . They cover both opposite-color near edges. In color 1, far- x occupancy forces the full even- x row $y = 1$, covering the even- x near edge $y = 0$. Far- y occupancy forces $(1, 2i) \in S$ for every $0 \leq i \leq b$, covering the even- y near edge $x = 0$.

The neighborhood of a fixed set is fixed under the two path operators, so its complement is upper closed within each coordinate-parity pattern. The two reflections convert it to a lower Ferrers set. They preserve checkerboard color because both side lengths are odd. The near-edge coverage just proved means the reflected complement has no far-edge vertices, as required. $\square$

4.3 Lower bounds in Theorem 4.1

Let $S \subseteq V_0$, $|S| = k > 0$, and put $\delta = |N(S)| - k$. If $\delta \geq b$ the desired lower bound is immediate. Otherwise Lemma 4.3 leaves two possibilities. With no clipping, $k \leq \delta^2$. With both far edges occupied, let T be the reflected opposite-color complement of $N(S)$. It has size $t = M - k - \delta$. Before reflection, its neighborhood is contained in $V_0 \setminus S$, so its quadrant surplus is at most $\delta + 1$. For $\delta \geq 1$, Lemma 4.2 gives

$$M - k - \delta = t \leq \delta(\delta + 1), \quad k \geq E - (\delta + 1)^2.$$

If $\delta = -1$ or 0 , T must be empty, since a nonempty quadrant color-1 set has surplus at least two. This yields respectively $k = E$ or $k = M$, the same conclusion. No smaller δ is possible: pairing

along each row, then vertically in the unpaired last column, gives a matching covering all vertices except the far even corner, and hence $|N(S)| \geq k - 1$. In all cases,

$$\delta \geq \min\{R(k), b, R(E - k) - 1\}.$$

For $S \subseteq V_1$, the same matching gives $\delta \geq 0$. Below $b + 1$, the unclipped case gives $k \leq \delta(\delta - 1)$. In the complement case, T has size $E - k - \delta$ and quadrant color-0 surplus at most $\delta - 1$. The case $\delta = 0$ is impossible: T is nonempty because $k \leq M < E$, but would have negative surplus. For $\delta \geq 1$,

$$E - k - \delta \leq (\delta - 1)^2, \quad k \geq M - \delta(\delta - 1).$$

The unclipped nonempty case also has $\delta \geq 1$. Thus

$$\delta \geq \min\{Q(k), b + 1, Q(M - k)\}.$$

These prove the lower bounds in both formulas.

For comparison, the two exact profiles of any bipartite graph obey the complement-inverse identity

$$g_1(y) = E - \max\{x : g_0(x) \leq M - y\}. \quad (18)$$

Indeed, the existence of nonadjacent sets of sizes x and y in the two colors is equivalent to $g_0(x) \leq M - y$, and also to $g_1(y) \leq E - x$. Thus the second formula can alternatively be checked from the first by exact inversion.

4.4 Attainment and neighborhood nesting

The vertices on diagonal $x + y = d$ have x in the interval

$$L_d = \max(0, d - 2b) \leq x \leq \min(2a, d),$$

whose length is

$$\ell(d) = \min(2a, d) - \max(0, d - 2b) + 1, \quad 0 \leq d \leq 2a + 2b.$$

A nonempty prefix ends after j vertices of diagonal d , for some $1 \leq j \leq \ell(d)$, and has size

$$k = \sum_{\substack{s < d \\ s \equiv d \pmod{2}}} \ell(s) + j. \quad (19)$$

All earlier opposite-color diagonals belong to its neighborhood. On diagonal $d + 1$, each selected x contributes x or $x + 1$ when the corresponding step stays inside the board. Their union is the initial interval there, of length

$$f(d, j) = j + 1 - \mathbf{1}_{\{d \geq 2b\}} - \mathbf{1}_{\{L_d + j - 1 = 2a\}}. \quad (20)$$

The first indicator removes the unavailable step past the short-direction boundary, and the second removes the step past the long-direction boundary. At the far corner both are present and $f = 0$. The initial minority diagonal $d = 1$ additionally covers the corner on diagonal zero, already included among the earlier opposite diagonals. Consequently

$$|N(I_p(k))| = \sum_{\substack{s < d \\ s \not\equiv d \pmod{2}}} \ell(s) + f(d, j),$$

and this neighborhood is exactly a prefix of the other color.

To evaluate its surplus, let $\Delta(d)$ be the earlier opposite-color sum minus the earlier same-color sum. The recurrence $\Delta(0) = 0$, $\Delta(d+1) = \ell(d) - \Delta(d)$ gives

$$\Delta(d) = \begin{cases} \lceil d/2 \rceil, & d \leq 2b, \\ b + (d \bmod 2), & 2b \leq d \leq 2a, \\ a + b - \lfloor d/2 \rfloor + (d \bmod 2), & d \geq 2a. \end{cases}$$

The formulas agree at their shared endpoints. The exact surplus is $\Delta(d) + 1$ minus the two indicators in (20). Complete even diagonals through sum $2r - 2$ contain r^2 vertices, and complete odd diagonals through sum $2r - 1$ contain $r(r+1)$. Substitution into (19) yields

Color	Size range	Prefix surplus
0	$1 \leq k \leq b^2$	$R(k)$
0	$b^2 < k < E - b^2$	b
0	$E - b^2 \leq k \leq E$	$R(E - k) - 1$
1	$1 \leq k \leq b(b+1)$	$Q(k)$
1	$b(b+1) < k < M - b(b+1)$	$b+1$
1	$M - b(b+1) \leq k \leq M$	$Q(M - k)$

For a square, the shared minority endpoint has surplus $b+1$ in both rows; empty ranges are omitted. The table equals the two lower bounds. It proves attainment and the compatible-neighborhood assertion, completing the proof of Theorem 4.1.

4.5 An exact two-count game

Corollary 4.4. *For every daily budget on every odd-by-odd rectangle, the unrestricted minimum capture time is exactly the shortest-path distance from (E, M) to $(0, 0)$ in the following two-count state graph:*

$$(u, v) \longrightarrow (g_1(j), g_0(i)), \quad 0 \leq i \leq u, \quad 0 \leq j \leq v, \quad u + v - i - j \leq m.$$

A shortest path gives an optimal physical strategy by deleting suffixes of the two current-color prefix orders. The reduction also respects arbitrary prescribed daily budgets.

Proof. Define

$$C(S) = I_0(|S \cap V_0|) \cup I_1(|S \cap V_1|).$$

It is monotone, cardinality preserving and idempotent. By the exact profiles and nesting,

$$N(C(S)) \cap V_{1-p} = I_{1-p}(g_p(|S \cap V_p|)) \subseteq I_{1-p}(|N(S \cap V_p)|).$$

The last set is the corresponding part of $C(N(S))$. Apply Theorem 3.2. Every displayed transition is also the exact neighborhood of the indicated survivors, so it has a physical realization. Shortest paths, with infinity for an unreachable zero state, settle both feasibility and minimum time. $\square$

The corollary is an exact recurrence for every parameter, rather than a closed expression for its shortest-path distance. It avoids imposing an unproved sequential allocation to the two initial colors: both may receive inspections on any day.

5 An exact isoperimetric order in every all-odd box

The preceding rectangle theorem extends to every dimension. The extension uses a weighted exchange argument: compression in lower-dimensional faces first makes neighborhood size a sum of local costs, and an order comparison then identifies exchanges that cannot increase that sum. Parity-restricted isoperimetry on hypercubes was studied by Körner and Wei [12], and local-to-global methods for Cartesian-product vertex boundaries by Bezrukov and Serra [4]. Here the neighborhood is the open neighborhood of a single color class. We give the face compressions and their exchange argument explicitly, rather than infer the result from a theorem about closed vertex boundaries.

Delete any factors of side length one, and write the remaining box as

$$B = \prod_{i=1}^d \{0, \dots, L_i\}, \quad 2 \leq L_1 \leq \dots \leq L_d, \quad L_i \text{ even.}$$

Thus coordinates are ordered from the shortest side to the longest. Let B_p be its color- p class, where the color of x is the parity of $|x| = \sum_i x_i$. Within either color, order vertices by increasing $|x|$ and then by *decreasing* lexicographic coordinates. Write $x \prec y$ for this order and $I_p(k)$ for its first k vertices. The direction of the lexicographic tie is part of the definition.

Theorem 5.1 (Compatible isoperimetric nesting in all-odd boxes). *For every $S \subseteq B_p$,*

$$|N(S)| \geq |N(I_p(|S|))|.$$

Moreover, $N(I_p(k))$ is a prefix in the opposite color. Consequently the minimum capture time for every inspection budget on B is given by an exact recurrence on $O(|B|^2)$ states, and optimal inspection sets can be recovered from that recurrence.

For $d = 1$ the minimizing order is the spacing-two prefix order of Lemma 3.4. For $d = 2$, putting the shorter coordinate first and ordering it decreasingly within a weight layer is equivalent to the increasing-long-coordinate convention of Theorem 4.1. These provide the induction bases. In particular, the higher-dimensional assertion does not assume a general Cartesian-product isoperimetric principle.

5.1 Prefix neighborhoods and a local cost identity

For a nonzero vertex z , let $j(z)$ be its last positive coordinate and put

$$\ell(z) = z - e_{j(z)}.$$

This is its earliest lower neighbor in decreasing lexicographic order.

Lemma 5.2 (Nesting of prefix neighborhoods). *For the stated weight and lexicographic order, the neighborhood of a prefix is a prefix of the opposite color. This statement holds even without the assumptions that the side lengths are odd and sorted.*

Proof. Within a fixed weight layer, ℓ preserves order, allowing ties. Indeed, suppose $z \prec w$ first differ in coordinate i , so $z_i > w_i$. Equal weights imply that w has a positive coordinate after i . If z does too, neither predecessor operation changes the first difference. Otherwise $j(z) = i$ and

$z_i - 1 \geq w_i$. A strict inequality preserves the order. In the equality case the tail of w has total weight one; removing its unique unit gives $\ell(z) = \ell(w)$.

Consider a nonempty prefix whose last, possibly partial, layer has weight r . Every smaller opposite-color layer is completely covered: its nonzero vertices have lower neighbors in complete source layers. When weight zero belongs to the opposite color, it is covered by the first odd layer. No vertex above layer $r + 1$ is covered. On layer $r + 1$, a vertex is covered precisely when its earliest lower neighbor belongs to the selected prefix in layer r . The order preservation just proved makes these covered vertices an initial interval of layer $r + 1$. The empty prefix is immediate. $\square$

Call a set *pair-compressed* if, after fixing all but any two coordinates, its selected vertices in either free-coordinate parity form a prefix of the inherited order. Define

$$a(0) = d, \quad a(x) = d - j(x) + 1 - \mathbf{1}_{\{x_{j(x)} = L_{j(x)}\}} \quad (x \neq 0). \quad (21)$$

Lemma 5.3 (Neighborhood cost of a pair-compressed set). *If $d \geq 2$ and $\emptyset \neq S \subseteq B_p$ is pair-compressed, then*

$$|N(S)| = \sum_{x \in S} a(x) + \mathbf{1}_{\{p=1\}}. \quad (22)$$

Proof. For every nonzero z , we claim

$$z \in N(S) \iff \ell(z) \in S.$$

Only the forward implication needs proof. Choose a neighbor $x \in S$ of z . The vertices x and $\ell(z)$ differ in at most two coordinates. If x is a lower neighbor, then $\ell(z)$ is no later than x in their common weight layer. If x is an upper neighbor, $\ell(z)$ has weight two less than x . Pair compression therefore puts $\ell(z)$ in S . If only one coordinate differs, any second coordinate can be used; one exists because $d \geq 2$.

For $x \neq 0$ with last positive coordinate J , the preimages of x under ℓ are obtained by adding one in coordinate J , if it is not full, or by adding one in any coordinate after J . Their number is $a(x)$. The origin has the d unit vectors as preimages. Thus the sum in (22) counts every nonzero neighbor exactly once.

If $p = 1$, a nonempty pair-compressed set contains a unit vector. Starting with any of its vertices of weight greater than one, decrease by two in a coordinate that is at least two, or decrease by one in two positive coordinates. Each move stays in a pair fiber and goes earlier in its order, so it preserves membership. The process ends at weight one. The origin is therefore an additional neighbor when $p = 1$, and it can never be a neighbor when $p = 0$. $\square$

5.2 Face compression and its partial order

Suppose $d \geq 3$ and the isoperimetric theorem has been proved in dimension $d - 1$. On each fixed color define a partial order $\leq_P$ as the reflexive transitive closure of

$$x <_P y \quad \text{if } x \prec y \text{ and } x_i = y_i \text{ for at least one } i. \quad (23)$$

All generating relations increase the total order, so this is indeed a partial order. Its ideals are exactly the sets whose restriction to every codimension-one fiber is a prefix.

The inductive isoperimetric theorem and compatible neighborhood nesting make the prefix map on each $(d - 1)$ -dimensional box a strategy compression: it is monotone, preserves cardinality, and

sends each color's neighborhood into the corresponding prefix of the original neighborhood's size. The remaining side lengths are still sorted, and their order is the restriction of the global order to a fixed-coordinate fiber. By Lemma 3.3, compressing all such fibers for one coordinate index is a strategy compression of B .

Cycle through the d indices. Every change strictly decreases the sum of the selected vertices' positions in the global weight and lexicographic order. This nonnegative integer is bounded by $|B|(|B| - 1)/2$. The simultaneous-compression argument in Section 3 therefore gives a uniform finite composition whose output is fixed by every face compression. In particular, every $S \subseteq B_p$ can be replaced by a P -ideal of the same size without increasing its neighborhood.

Every P -ideal is pair-compressed. A two-coordinate fiber lies inside a codimension-one fiber obtained by fixing a coordinate outside the pair; such a coordinate exists because $d \geq 3$. The restriction of a prefix to a subfiber is a prefix of its inherited order. Thus Lemma 5.3 applies to the resulting ideals.

5.3 The exchange comparison

The following elementary comparison is the step that links the local cost identity to the global order.

Lemma 5.4 (Terminal-coordinate comparison). *Let $d \geq 3$. If same-parity vertices satisfy $x \prec y$ and $x_d = 0 < y_d$, then $x \leq_P y$. Also, if $x \prec y$ and $x_d < y_d = L_d$, then $x \leq_P y$.*

Proof. First, if $u \leq v$ coordinatewise and their weights have the same parity, then $u \leq_P v$. Perform all required increments of size two within individual coordinates, and then pair the remaining increments of size one. Each step increases weight by two and changes at most two coordinates. Since $d \geq 3$, it leaves a coordinate unchanged and is a generator of (23).

For the first assertion, a coordinate agreement between x and y already gives a generator. Assume henceforth that all coordinates differ. If $|x| < |y|$ and some $i < d$ has $x_i > y_i$, transfer $t = x_i - y_i$ units from coordinate i to coordinate d of x , obtaining z . This is legal because $0 < t \leq L_i \leq L_d$ and $x_d = 0$. We have $x \prec z$ within their common weight layer, while $z \prec y$ follows from $|z| = |x| < |y|$. The vertex z retains $d - 2 \geq 1$ coordinates of x and agrees with y in coordinate i . Hence $x <_P z <_P y$. If there is no such excess coordinate, then $x \leq y$ coordinatewise, and the preceding observation applies.

It remains to consider $|x| = |y|$. Since all coordinates differ, $x_1 > y_1$. If a middle coordinate $2 \leq i \leq d - 1$ has $x_i > y_i$, use the same transfer to coordinate d . Again $x \prec z$, and $z \prec y$ because $z_1 = x_1 > y_1$. The same coordinate agreements give the two generators. Otherwise all middle coordinates increase: $x_i < y_i$ for $2 \leq i \leq d - 1$. Weight equality gives

$$x_1 - y_1 = y_d + \sum_{i=2}^{d-1} (y_i - x_i).$$

Transfer y_d units from the first to the last coordinate of x . The resulting vertex z is valid, since

$$z_1 = y_1 + \sum_{i=2}^{d-1} (y_i - x_i) > y_1 \geq 0, \quad z_d = y_d \leq L_d.$$

It satisfies $x \prec z \prec y$, shares the middle coordinates with x , and shares the last coordinate with y . This proves the first assertion.

The coordinate complement $\rho(x)_i = L_i - x_i$ reverses the weight and lexicographic order, and it preserves coordinate agreements. It therefore reverses the generated partial order. Applying the first assertion to $\rho(y) \prec \rho(x)$ proves the second assertion. $\square$

Lemma 5.5 (Cost comparison). *If $x \prec y$ have the same parity and $a(x) > a(y)$, then $x \leq_P y$.*

Proof. If $x_d = y_d$, they share a coordinate. If $x_d = 0 < y_d$, use the first part of Lemma 5.4. The case $y_d = 0 < x_d$ cannot give a strict cost decrease, since $a(x) \leq 1$ and $a(y) \geq 1$. Finally, if both last coordinates are positive, both costs are zero or one. A strict decrease forces $x_d < L_d = y_d$, so the second part of that lemma applies. $\square$

Proof of Theorem 5.1. The path and rectangle bases were noted above. In dimension $d \geq 3$, the inductive face compressions replace an arbitrary $S \subseteq B_p$ by a P -ideal of the same size and with no larger neighborhood. The empty set needs no further argument.

If a nonempty P -ideal is not a global prefix, let x be its earliest missing vertex and y its latest included vertex. Then $x \prec y$. They are incomparable in P : $x \leq_P y$ would contradict ideality, and $y \leq_P x$ would contradict the global order. By Lemma 5.5, $a(x) \leq a(y)$.

Replace y by x . Removing y preserves ideality because every proper successor of y is later in the global order and hence absent. Adding x preserves ideality because every proper predecessor of x is earlier and hence present; y is not such a predecessor. The new set has the same positive cardinality and color. The neighborhood cost identity (22) shows that its neighborhood does not increase. The sum of global positions strictly decreases, so iteration ends at the global prefix of that size. This proves the isoperimetric inequality. Lemma 5.2 gives compatible nesting, completing the induction. $\square$

5.4 Profiles, exact capture times, and optimal inspections

Write $V = |B|$, $E = (V + 1)/2$, and $M = (V - 1)/2$. The theorem supplies the complete profiles by a direct cumulative-sum construction:

$$g_p(0) = 0, \quad g_p(k) = \sum_{x \in I_p(k)} a(x) + \mathbf{1}_{\{p=1\}} \quad (k > 0). \quad (24)$$

For $d = 1$ these expressions also agree with the odd-path profiles. Enumerate the vertices, order each color as specified, evaluate (21), and take cumulative sums. Using a mixed-radix lexicographic index, sorting requires $O(V \log V)$ integer comparisons; computing weights, indices and costs directly requires $O(dV)$ additional operations. This is a constructive profile algorithm, without an isoperimetric optimization subroutine.

Corollary 5.6 (Feasibility by one profile maximum). *For every nontrivial all-odd box, the minimum feasible daily budget is*

$$h(B) = 1 + \max_{0 \leq k \leq E} (g_0(k) - k) = 1 + \max_{0 \leq k \leq E} \sum_{x \in I_0(k)} (a(x) - 1). \quad (25)$$

The one-room box has threshold one by direct inspection.

Proof. We verify the hypotheses of the general hunter-number criterion of Bolkema and Groothuis [5, Theorem 16]. That theorem states that a bipartite graph with compatible isoperimetric nesting

has hunter number $1 + \min(u_0, u_1)$ when the maximum neighborhood surpluses $u_p = \max_k(g_p(k) - k)$ differ by at most one.

For completeness, the exact profile duality gives the required relation between these maxima. Define

$$J(z) = \max\{x : 0 \leq x \leq E, g_0(x) \leq z\}, \quad 0 \leq z \leq M.$$

Then

$$g_1(y) = E - J(M - y). \quad (26)$$

Indeed, a majority set of size x with at most $M - y$ neighbors leaves a minority y -set with no edges to it, giving $g_1(y) \leq E - x$. Conversely, the majority complement of the neighborhood of a minimizing minority y -set has size $E - g_1(y)$ and at most $M - y$ neighbors. These two choices prove (26). Since $E = M + 1$, it follows that

$$u_1 = 1 + \max_{0 \leq z \leq M} (z - J(z)).$$

This last maximum equals u_0 . If $J(z) < E$, let $x = J(z) + 1$. Then $g_0(x) > z$, so $z - J(z) \leq g_0(x) - x \leq u_0$. If $J(z) = E$, the absence of isolated vertices implies $g_0(E) = M$; thus $z = M$ and $z - J(z) = -1 \leq u_0$. For the reverse inequality choose a positive x attaining u_0 . Such an x exists even when $u_0 = 0$, because $g_0(1) \geq 1$ then forces $g_0(1) = 1$. For $z = g_0(x) - 1$ we have $0 \leq z < M$ and $J(z) \leq x - 1$, hence $z - J(z) \geq g_0(x) - x = u_0$. Therefore $u_1 = u_0 + 1$.

Theorem 5.1 supplies compatible nesting, so the cited criterion yields $h(B) = 1 + u_0$. Its general lower-bound theorem is being used here; profile duality alone is not a lower-bound proof. The cumulative-sum formula follows from (24). $\square$

The upper bound also admits a direct description. Spend all $m = u_0 + 1$ inspections on one current prefix cohort. If it is not yet captured, a majority step decreases its size by at least one, while a minority step does not increase it, because $u_1 = u_0 + 1$. Thus it clears after finitely many steps. Clear the other initial cohort next. This proves feasibility without asserting that serial allocation minimizes the time.

For example, the partial-sum scan gives

B	3×3	5×5	$3 \times 3 \times 3$	$3 \times 5 \times 5$	$5 \times 5 \times 5$	$3 \times 3 \times 3 \times 3$
$h(B)$	2	3	5	8	12	12

Feasibility therefore needs only the profile scan. Minimum capture time uses the following two-count recurrence.

Define

$$C(S) = I_0(|S \cap B_0|) \cup I_1(|S \cap B_1|).$$

Isoperimetry and nesting imply that C is an idempotent strategy compression. Theorem 3.2 therefore gives an optimal strategy from the full board in which both current-color parts and both survivor parts are prefixes.

A state is the pair (u, v) of counts in the *current* colors. Choosing survivor counts $0 \leq i \leq u$ and $0 \leq j \leq v$ costs $u + v - i - j$ inspections and gives the exact next state

$$(u, v) \longrightarrow (g_1(j), g_0(i)). \quad (27)$$

The inspection sets are the suffixes $I_0(u) \setminus I_0(i)$ and $I_1(v) \setminus I_1(j)$. Thus every transition is physically realized, and every unrestricted successful search is matched by a path in this state graph.

For a fixed daily budget m , the optimum is the shortest-path distance from (E, M) to $(0, 0)$, with value infinity when there is no such path. There are $(E + 1)(M + 1) = O(V^2)$ states. If $u + v \leq m$, one final day captures every possibility. Otherwise, additional inspections cannot hurt, so it suffices to allocate exactly m useful inspections. For

$$\max(0, m - v) \leq p \leq \min(m, u),$$

the corresponding transition is

$$(u, v) \longrightarrow (g_1(v - m + p), g_0(u - p)). \quad (28)$$

There are at most $m + 1$ transitions per state. Constructing the graph and performing breadth-first search therefore take $O((m + 1)V^2)$ operations after constructing the profiles. Recording a shortest path recovers an explicit optimal inspection sequence. A finite optimum is at most $(E + 1)(M + 1) - 1$, since a shortest path repeats no state. The bounds are polynomial in the number of rooms, rather than in the logarithms of the side lengths.

Equivalently, all budgets can be considered together. Let $W_t(u, v)$ be the smallest daily budget sufficient in at most t days. Then

$$\begin{aligned} W_0(0, 0) &= 0, & W_0(u, v) &= \infty \quad ((u, v) \neq (0, 0)), \\ W_{t+1}(u, v) &= \min_{\substack{0 \leq i \leq u \\ 0 \leq j \leq v}} \max\{u + v - i - j, W_t(g_1(j), g_0(i))\}. \end{aligned} \quad (29)$$

This exact recurrence includes arbitrary interleaving of the two initial parity cohorts; it makes no serial-allocation assumption.

Remark 5.7. Every weight and lexicographic prefix is a checkerboard corner ideal: a distinct coordinatewise smaller vertex of the same parity has smaller weight. Hence the theorem also proves a corner-ideal normal form for the full-board game on every all-odd box. It does not assert that a global prefix can replace an arbitrary prescribed partial starting set without altering its optimum, and it does not settle boxes with even side lengths. The all-odd-box result gives exact feasibility, minimum time, and strategies through a polynomial recurrence; a uniform closed arithmetic expression for that time is a further question.

6 Paths: matching geometry and a probe-count potential

We prove the path formula stated in (1). The geometric lower proof applies to arbitrary possible-position sets. An interval description is used only for the explicit strategy attaining the bound.

6.1 Neighborhood inequalities

Number the vertices $0, \dots, n - 1$. Each set below is contained in one color class. If n is even, the edges $(0, 1), (2, 3), \dots$ give a perfect matching and therefore $|N(R)| \geq |R|$. We need the following information about equality.

Lemma 6.1. *On an even path, if $0 < |R| < n/2$ and $|N(R)| = |R|$, then $N(R)$ contains neither endpoint. Every nonempty set containing neither endpoint has strictly more neighbors than vertices. On an odd path with larger class E and smaller class O ,*

$$|N(R)| \geq \min\{|R|, |O|\} \quad (R \subseteq E), \quad |N(R)| \geq |R| + 1 \quad (\emptyset \neq R \subseteq O).$$




Figure 3: An optimal four-day search of eight rooms with two daily probes. The rows show every position still possible before that day's inspection.

Proof. Let M send each vertex of an even path to its matching partner. Equality implies $N(R) = M(R)$. If R consists of even vertices and contains $2i > 0$, its neighbor $2i - 1$ forces $2i - 2 \in R$. Thus equality makes R an initial segment of the even vertices. Similarly an odd equality set is a final segment of the odd vertices. These statements prove both assertions about endpoints: a proper initial even segment has an odd neighborhood missing the far endpoint, and a proper final odd segment has an even neighborhood missing the other endpoint. Conversely, a nonempty equality set must itself contain an endpoint.

When n is odd, for every even vertex g there is a matching covering all vertices except g : match consecutive vertices separately to its left and right. If $R \subsetneq E$, choose $g \notin R$ to inject R into $N(R)$. For $R = E$ its neighborhood is O . Finally, for nonempty $R \subseteq O$, the matching exposing vertex 0 injects R into $N(R)$, and equality is impossible: the least occupied odd vertex has a left neighbor that would force a smaller occupied odd vertex, or the unmatched vertex 0. Thus there is at least one extra neighbor. $\square$

Track separately the targets that started in the two colors. Their current colors are opposite each day, so their allocations p, q always satisfy $p + q \leq m$. Write b for a cohort's size before inspection. On an even path retain a bit z : it is one just after a nonempty proper zero-growth neighborhood, and zero otherwise. Such a set is endpointless by the lemma. A surviving subset of it therefore expands strictly. The rank $u = 2b + z$, with rank zero for the empty set, is at most n and each physical step dominates

$$F_C(u, p) = \begin{cases} 0, & \lfloor u/2 \rfloor \leq p, \\ \min\{C, u - 2p + c\}, & \lfloor u/2 \rfloor > p, \end{cases} \quad (30)$$

with $C = n, c = 1$. Indeed, a zero-growth step raises the bit from zero to one; a step following bit one gains a vertex and may discard the bit. Both changes give at least one unit after subtracting $2p$.

For odd n , assign rank $2b$ in E , rank $2b + 1$ in nonempty O , and zero to the empty set. The same comparison holds with $c = 1$ and alternating new caps $n + 1$ in E and n in O . Initially the ranks are $n + 1, n$. Replacing both caps and initial ranks by n gives a valid lower comparison, since F_C is nondecreasing in C and u .

6.2 A general potential for the counting argument

Lemma 6.2 (Affine-rank potential). *Let $0 \leq c < 2m$, $D = 2m - c$, and $u \leq C$. Define*

$$J(u) = \left\lfloor \frac{\max\{u - c - 2, 0\}}{D} \right\rfloor, \quad \Phi(u) = \left\lfloor \frac{u + cJ(u)}{2} \right\rfloor.$$

For $0 \leq p \leq m$, the constant-cap transition (30) satisfies $\Phi(u) \leq p + \Phi(F_C(u, p))$. Consequently two such ranks sharing budget m need at least $\lceil (\Phi(u_0) + \Phi(v_0))/m \rceil$ rounds to become zero.

Proof. Both J and Φ are nondecreasing. If capture occurs, then $u \leq 2p + 1 \leq 2m + 1$, whence $J(u) = 0$ and $\Phi(u) \leq p$. If $2p \leq c$ and capture does not occur, the next rank is at least u , because $u \leq C$. Otherwise put $v = u - 2p + c$. It is positive and at most $u \leq C$, so the cap does not act. Since $2p - c \leq D$, $J(u) \leq J(v) + 1$. Therefore

$$u + cJ(u) \leq v + 2p + cJ(v).$$

Divide by two and take floors. Summing the resulting inequalities over both ranks and all rounds proves the last assertion. $\square$

For paths $c = 1$. Suppose $n > m \geq 2$, put $D = 2m - 1$, and use positive remainder coordinates

$$n = qD + r + 2, \quad 1 \leq r \leq D.$$

Here $n \geq 3$ and $q \geq 0$. Direct substitution gives

$$\Phi(n) = qm + \lfloor (r + 2)/2 \rfloor.$$

The resulting lower bound is $2q + 1$ when $r \leq m - 2$, or when $r = m - 1$ and m is even; otherwise it is $2q + 2$. For even n this is the required answer. For odd n it misses just the case m even and $r = m - 1$, which we now settle.

6.3 The equality obstruction on odd paths

Here $n = q(2m - 1) + m + 1$ with m, q even. Assume a strategy succeeds in $T = 2q + 1$ days. Use the exact alternating-cap comparison counters. For each counter, select its first zero and its last full state before that zero. If this interval starts at time a , has length L , and uses P probes, all intermediate transitions are unsaturated. Telescoping the affine transitions, including the final capture inequality, gives

$$C(a) + L \leq 2P + 2, \quad P \leq Lm.$$

Since $C(a) \geq n$, these imply $L \geq q + 1$ and $P \geq A := qm + m/2$. The two cohorts together have only $Tm = 2A$ probes, so equality holds for both. In particular $L = q + 1$ and $C(a) = n$: both intervals start when their cohort is in the smaller color class.

Their starting times have opposite parity and are at most q . The odd one, a , satisfies $1 \leq a \leq q - 1$ because q is even. If the even starting time is zero, both intervals lie in the first $2q$

days; otherwise both lie in the last $2q$ days. Their total probe count $2A = (2q + 1)m$ exceeds the $2qm$ available in either window. This contradiction proves that one extra day is necessary. The argument permits overlapping active intervals and arbitrary wasted inspections outside them.

For $m = 1$, the same potential has $D = 1$ and $\Phi(n) = n - 2$ for $n \geq 3$, giving $2n - 4$ days. On P_2 one probe cannot inspect both initial possibilities, and two consecutive inspections of one endpoint suffice. A single room is inspected directly.

6.4 An explicit strategy attaining the bound

For this construction number rooms $1, \dots, n$. Put

$$A(r) = \{r, r - 2, r - 4, \dots\} \cap \{1, \dots, n\}.$$

For $1 \leq r < n$, $N(A(r)) = A(r + 1)$. Inspecting the b largest vertices of $A(r)$ leaves $A(r - 2b)$; if nonempty, movement changes the frontier to $r - 2b + 1$. Reflection in the middle of the path gives the same rule in the opposite direction.

Assume $n \geq 3$ and $n > m$. Set $d = 2m - 1$ and

$$e = \left\lceil \frac{n - 2}{d} \right\rceil, \quad \rho = n - 2 - (e - 1)d, \quad c = \lfloor \rho/2 \rfloor + 1, \quad \ell = m - c.$$

Thus $1 \leq \rho \leq d$. Start with the cohort occupying $A(n - 1)$ and inspect its largest m possible rooms each day. Its frontier before day t is $n - 1 - (t - 1)d$, so it finishes on day e , using only c inspections that day. If $\ell > 0$, assign the remaining ℓ inspections to the other cohort immediately; if $\ell = 0$, start that cohort on day $e + 1$.

Before its first inspection the second cohort is a whole color class. For odd n , use the same orientation. Its frontier is n on an odd starting day and $n - 1$ on an even starting day. For even n , choose the orientation whose frontier is $n - 1$: reflect when the starting day is odd. If its first batch has size b , the surviving frontier is $R_0 - 2b + 1$. Continue with batches of m until capture. A nonempty frontier $A(R)$ with $R \geq 2$ requires $\lceil (R - 1)/d \rceil$ such days; a singleton is inspected once.

For completeness the arithmetic determining the joining day is as follows. If $\ell = 0$ the total is $2e$. If $\ell > 0$, the totals are

$$2e - 2 + \left\lceil \frac{2\rho + 2}{d} \right\rceil \quad (n \text{ odd}), \quad 2e - 2 + \left\lceil \frac{\rho + 2c}{d} \right\rceil \quad (n \text{ even}).$$

They equal $2e - 1$ exactly when $\rho \leq m - 2$, or when $\rho = m - 1$ and n, m are both even; otherwise they equal $2e$. Comparing with $\lceil (2n - 4)/d \rceil$ proves (1). For $m = 1$ it gives $2n - 4$. If $n \leq m$, inspect every vertex on the first day. This completes both the construction and the arbitrary-strategy lower bound.

7 Two rows: an exact pair of counters

A two-row rectangle $P_2 \square P_n$ is also called a *ladder*: the two horizontal paths are joined by one vertical edge in each column. Its graph terminology does not impose any restriction on the searcher's movements between inspections.

For a cohort of fixed initial color, at any one time exactly one room per column has the right color. Identify its possible rooms with their column set $B \subseteq \{1, \dots, n\}$. A vertical move keeps

the column, and a horizontal move changes it by one. Consequently the next column set after inspections is the *closed* path neighborhood

$$N_{P_n}[B \setminus S] = (B \setminus S) \cup N_{P_n}(B \setminus S).$$

Every nonempty proper column set gains at least one column. An initial or final interval attains the bound and stays an interval. Therefore the exact one-cohort size transition with p probes is

$$F(b, p) = \begin{cases} 0, & p \geq b, \\ \min\{n, b - p + 1\}, & p < b. \end{cases}$$

The two counters start at (n, n) and share the budget m . This is an exact lower and upper model: the inequality holds for arbitrary supports, and two boundary intervals attain every chosen sequence of allocations.

Theorem 7.1. *For $n \geq 2$ the minimum-time function is (2), with value ∞ for $m \leq 1$ and value one for $m \geq 2n$.*

Proof. For $2 \leq m < 2n$, put $N = n - 1$ and $M = m - 1$. Assign a counter of size b the weight $w(b) = \max\{b - 1, 0\}$. With a positive allocation p , its weight drops by at most $p - 1$; with no allocation its weight cannot drop. Thus the combined daily drop is at most M , and $T \geq \lceil 2N/M \rceil$.

Suppose $2N = TM$. Equality in this bound requires every day to allocate all m probes to exactly one cohort and to reduce its weight by exactly M . Each initial weight N must consequently be divisible by M . If $2N$ is divisible by M but N is not, one extra day is necessary. This is exactly the exceptional condition that M is even and $N \equiv M/2 \pmod{M}$.

To attain the answer, write $N = qM + s$, $0 \leq s < M$. If $s = 0$, clear the first interval in q full days and the second in q more. Suppose $s > 0$. After q full days on the first interval its size is $s + 1$. On the next day finish it with $s + 1$ probes and assign the remaining $M - s$ to the other, still full, interval. Its next size is

$$b = n - (M - s) + 1 = (q - 1)M + 2s + 2.$$

It then needs $\lceil (b - 1)/M \rceil$ days. The total is

$$2q + \left\lceil \frac{2s + 1}{M} \right\rceil,$$

which is the claimed ceiling plus precisely the equality correction. The second interval always survives the shared day. If $q = 0$, then $s = n - 1$ and its allocation is $M - s = m - n < n$, because $m < 2n$. If $q \geq 1$, then $M \leq n - 1$ and $M - s < n$ as well. Thus $b \geq 2$ in both cases, and the displayed solo-day formula counts at least one remaining day. The algebra also holds for $q = 0$, by the integer translation rule for ceilings.

For one probe, an uninspected cohort remains full. Inspecting one room of a full color class also leaves the other color class full after movement, since every room on a ladder with $n \geq 2$ has at least two neighbors. Hence capture is impossible. The one-day assertion is immediate from the number of initial rooms. For $n = 1$ the graph is P_2 and the path result applies. $\square$

8 Three rows: a historical rank and the exact time

We prove the three-row classification in (3), including optimal inspection schedules. Put $G = P_3 \square P_n$. The length $n = 1$ is a three-vertex path. Henceforth $n \geq 2$; a four-cycle contained in G precludes capture with one inspection per day. Indeed, a target may stay on that cycle, where a nonempty parity cohort always has two possible vertices after movement. The constructions below show that two inspections suffice.

8.1 Even lengths: the efficient survivor sets

Let $n = 2s$ and $h = 3s$, so both parity classes have h vertices. In coordinates (column, row), label one parity by

$$a_i = (2i, 0), \quad b_i = (2i, 2), \quad c_i = (2i + 1, 1) \quad (0 \leq i < s).$$

Match these vertices respectively to $(2i + 1, 0)$, $(2i + 1, 2)$, $(2i, 1)$ in the other parity. Contracting an edge followed by the inverse matching map gives a directed graph D on the labelled parity. Besides loops, its arcs are

$$a_i \leftrightarrow c_i \leftrightarrow b_i, \quad a_i \rightarrow a_{i-1}, \quad b_i \rightarrow b_{i-1}, \quad c_i \rightarrow c_{i+1},$$

with out-of-range subscripts omitted. Consequently, for every survivor set R in this parity,

$$|N_G(R)| - |R| = |N_D^+(R) \setminus R|. \quad (31)$$

Each triple is strongly connected; its middle chain moves right and its outer chains move left. Thus D is strongly connected, and every nonempty proper parity set has surplus at least one.

Lemma 8.1. *If R is nonempty, $|R| \leq h - 2$, and $|N(R)| = |R| + 1$, then*

$$R = \{a_i, b_i, c_i : 0 \leq i < j\} \cup U, \quad U \subseteq \{a_j, b_j\}, \quad 0 \leq j < s.$$

For the opposite parity the analogous sets start at the other end. Moreover, if $R' \subseteq N(R)$ is nonempty, also has surplus one, and $|R'| \leq h - 2$, then $|R| = h - 2$ and $|R'| = 1$.

Proof. By (31), the outside boundary is a single vertex v , and R is closed under outgoing arcs of $D - v$. Deleting an outer vertex leaves a strongly connected graph: the other outer chain still moves left, the middle chain moves right, and all remaining outer vertices connect to their middle vertices. A nonempty closed set would then have all $h - 1$ remaining vertices, contrary to the hypothesis.

Deleting c_j leaves the prefix triples, the two singletons a_j, b_j , and the suffix triples as strongly connected components, omitting empty components. Arcs go from the suffix to each singleton and from each singleton to the prefix. A nonempty outgoing-closed set of size at most $h - 2$ is therefore precisely the stated prefix with some of the two singletons. Reflection in the long coordinate exchanges the physical parities and gives the opposite description.

Every nonempty classified set contains an outer corner at its starting end, and a set of size at least two contains both outer corners there. The neighborhood of R contains no outer corner at the opposite end unless $j = s - 1$. In that case the size bound permits at most one of a_j, b_j . Containing an opposite corner therefore forces $|R| = h - 2$ and supplies exactly one such corner. Hence R' must be a singleton. $\square$

8.2 A rank valid for arbitrary supports

We first assume $m < h - 2$. Track the two initial-parity cohorts separately. For a physical support A , set $z = 1$ if the preceding survivor set was nonempty, had surplus one, and had size at most $h - 2$; otherwise set $z = 0$. In particular full and empty supports have label zero. Define $\rho(A) = 2|A| + z$ for nonempty A , and $\rho(\emptyset) = 0$. This label records history and imposes no shape restriction on A .

For a numerical rank $0 \leq u \leq 2h$ and an allocation p , put

$$F_h(u, p) = \begin{cases} 0, & \lfloor u/2 \rfloor \leq p, \\ \min\{2h, u - 2p + 3\}, & \lfloor u/2 \rfloor > p. \end{cases} \quad (32)$$

It is nondecreasing in u and nonincreasing in p .

Lemma 8.2. *If $p \leq m < h - 2$ useful probes leave $R \subseteq A$, then $\rho(N(R)) \geq F_h(\rho(A), p)$.*

Proof. Write $r = |R| = |A| - p$. Empty survivors are immediate. If $r = h$ or $r = h - 1$, the neighborhood is full: every opposite-parity vertex has degree at least two. Thus its rank is the cap $2h$.

Suppose $1 \leq r \leq h - 2$. If the surplus is one, the next label is one. The old label cannot be one: Lemma 8.1 would require a support of size $h - 1$ to be reduced to one vertex, using $h - 2$ probes. Hence the new rank is $2(r + 1) + 1 = \rho(A) - 2p + 3$. If the surplus is at least two, the new rank is at least $2r + 4 \geq \rho(A) - 2p + 3$, since the old label is at most one. Capping can only weaken these inequalities. $\square$

Starting from $(2h, 2h)$, compare any actual schedule to the two-counter process

$$(u, v) \mapsto (F_h(u, p), F_h(v, m - p)), \quad 0 \leq p \leq m.$$

Its terminal inspection condition is $\lfloor u/2 \rfloor + \lfloor v/2 \rfloor \leq m$. Use the actual two useful allocations; assign unused budget arbitrarily. Monotonicity and Lemma 8.2 keep the abstract ranks below the actual ranks throughout, so the numerical process captures no later than the actual schedule.

Conversely this process is physically realizable. Order each parity by increasing coordinate sum, breaking ties by increasing column. The two orders are

$$\begin{aligned} E &:: (0, 0), \quad ((2i - 2, 2), (2i - 1, 1), (2i, 0))_{i=1}^{s-1}, \quad (2s - 2, 2), (2s - 1, 1), \\ O &:: (0, 1), (1, 0), \quad ((2i - 1, 2), (2i, 1), (2i + 1, 0))_{i=1}^{s-1}, \quad (2s - 1, 2). \end{aligned}$$

Empty repeated blocks are omitted. Reading the neighbors of successive vertices in these lists gives, for every $k > 0$,

$$N(E_k) = O_{\min(h, k+1)}, \quad N(O_k) = E_{\min(h, k+2)}.$$

Here E_k, O_k denote prefixes of size k ; empty prefixes have empty neighborhood. Deleting a suffix of p vertices and moving therefore realizes (32), with phases $z = 0, 1$ respectively. A full cohort may choose either reflected orientation freely; reflection in the long coordinate exchanges physical parities. The two cohorts choose these orientations independently. Proper cohorts retain their orientation until they fill a parity class again or disappear. Thus the numerical process gives the exact time for $m < h - 2$.

8.3 Closed formula for the even case

Assume $2 \leq m < h$, put $W = 2h - 4$, $D = 2m - 3$, and write $W = qD + r$, $0 \leq r < D$. A solo cohort receiving m inspections every day clears in $\lceil W/D \rceil$ days: before capture its rank decreases by D , and a final rank of at most $2m + 1 = D + 4$ can be inspected. Sequential solo sweeps therefore give $2q$ days if $r = 0$, and $2q + 2$ otherwise.

For $r > 0$, an overlap saves one day precisely in the case appearing in (3). Give q full-budget days to the first cohort, whose rank becomes $r + 4$. On the shared day clear it with $p = \lfloor (r + 4)/2 \rfloor$ probes and use the other $m - p$ probes on the second full cohort. Its new rank minus four is

$$(q - 1)D + r + 2p.$$

It clears in q further days if $r + 2p \leq D$. Since W is even and D is odd, $r \equiv q \pmod{2}$; this inequality is exactly $m \geq r + 4 - (q \bmod 2)$. It also ensures $m - p \geq 2$, so the shared day starts a proper sweep without hidden saturation.

For the matching lower bound, take any successful numerical schedule. For each cohort, retain only its interval from the last full state before first capture to that capture. Probes before this interval can be discarded: an uninspected full cohort remains full, and the retained trace starts in exactly the same state. There is no later saturation within the retained interval. If it has ℓ inspection days and uses P probes, the $\ell - 1$ nonterminal moves each add three rank units. The final capture inequality yields

$$P \geq \left\lceil \frac{W + 3\ell}{2} \right\rceil, \quad \ell D \geq W. \quad (33)$$

The intervals may overlap arbitrarily; their probe counts still sum to at most the budget of the whole schedule. If $r = 0$, each has $\ell \geq q$ and $P \geq qm$, giving $T \geq 2q$. If $r > 0$, each has $\ell \geq q + 1$, whence

$$P \geq qm + \lceil (r + 3)/2 \rceil.$$

This rules out $2q$ days. A $(2q + 1)$ -day schedule requires $m \geq 2 \lceil (r + 3)/2 \rceil = r + 4 - (q \bmod 2)$, exactly the condition used by the shared-day construction. Extra interleaving cannot evade it.

8.4 The even high-budget exceptions

The lower simulation used $m < h - 2$, so its boundary must be checked separately. A perfect matching gives $|N(B)| \geq |B|$ for every room set, including mixed parities. If $m < h$, at least $2h - m > m$ rooms remain after the first movement; two-day capture is impossible. If $h \leq m < 2h$, inspect one fixed physical parity on each of two consecutive days. One day requires $m \geq 2h$.

For $m = h - 1$, $h \geq 6$, and for $m = h - 2$, $h \geq 9$, there is a three-day construction. Start a phase-zero sweep on one cohort with m probes; its next size is $h - m + 1$. Clear it on day two, giving the remaining probes to the independently oriented other full cohort. The latter then has $2h - 2m + 2 \leq m$ rooms, and day three clears it. The preceding two-day obstruction proves optimality.

The remaining case is $h = 6, m = 4$. After the first movement at least nine rooms remain: neither initial cohort can be eliminated, and any touched proper survivor cohort expands by at least one. With no useful probe the bound is immediate. The perfect matching then leaves at least $9 - 4 = 5$ rooms after the second movement, precluding capture on day three. Two consecutive two-day cohort sweeps, using four probes and then three, give four days. Finally $h = 3$ is the

two-row board $P_2 \square P_3$: its two-probe time is four, by (2). These boundary values agree directly with (3).

8.5 Odd lengths and the extra parity constraint

Let $n \geq 3$ be odd and write $3n = 2h + 1$. By Theorem 4.1 and Corollary 4.4, the diagonal minimizing prefixes reduce the unrestricted game exactly to their two counts. Their neighborhood sizes are $\min(h, k + 1)$ on nonempty majority prefixes and $\min(h + 1, k + 2)$ on nonempty minority prefixes. Assign rank $2b$ to majority prefixes and $2b + 1$ to nonempty minority prefixes; empty supports have rank zero. Put $W = 3n - 4$. The full ranks are $W + 5$ and $W + 4$, and every nonterminal unsaturated transition is again $u \mapsto u - 2p + 3$. Saturation uses the full rank of the next physical parity, so the caps alternate.

Suppose $2 \leq m < h$ and write $W = qD + r$ with $D = 2m - 3$ and $0 \leq r < D$. Begin with the minority cohort, of rank $W + 4$. If $r = 0$, then q is odd; after its q -day solo sweep the untouched cohort is also in minority parity. Two such sweeps give $2q$ days. If $r > 0$, either full rank clears in at most $q + 1$ days, because $W + 1 \leq (q + 1)D$, giving $2q + 2$ days sequentially.

For a shared day, give the first, initially minority, cohort q full-budget days and then $p = \lfloor (r + 4)/2 \rfloor$ probes to finish it. At that day the other cohort's full rank is $W + 4 + e$, where $e = 1 - (q \bmod 2)$. Its next rank minus four is

$$(q - 1)D + r + e + 2p.$$

Since W is odd, $r \equiv e \pmod{2}$. Thus $r + e + 2p \leq D$ is exactly $m \geq r + 4$, giving the $(2q + 1)$ -day construction. This condition also ensures a proper, unsaturated departure of the second cohort on the shared day.

For the lower bound, extract each cohort's last-full active interval as above. If it starts in minority parity put $e = 0$, and otherwise put $e = 1$. Its length and probe count satisfy

$$P \geq \left\lceil \frac{W + e + 3\ell}{2} \right\rceil, \quad \ell D \geq W + e. \quad (34)$$

When $r = 0$, $\ell \geq q$ gives $P \geq qm$ for each cohort, proving $T \geq 2q$. When $r > 0$, $\ell \geq q + 1$ gives $P \geq qm + \lceil (r + 3)/2 \rceil$, ruling out $2q$ days. A total of $2q + 1$ days therefore requires $m \geq 2 \lceil (r + 3)/2 \rceil$. For even r this is the claimed $m \geq r + 4$. For odd r the only remaining possibility is $m = r + 3$.

That possibility is impossible by equality and parity. Here m and q are even. Equality in the total probe budget forces both active intervals to have length $q + 1$ and to start in minority parity; a longer interval or a majority start increases the right side of (34). Number days from zero. The two initial cohorts have opposite parity, so their starting times consist of an odd a and an even b , both at most q . Thus $1 \leq a \leq q - 1$. If $b = 0$, both intervals lie within $[0, 2q]$; if $b > 0$, both lie within $[1, 2q + 1]$. Either way their total probe count is at most $2qm$, contradicting the necessary $(2q + 1)m$. This establishes the extra parity constraint and proves the odd case of (3).

For completeness, a Hamilton path supplies a matching missing only one vertex, so $|N(B)| \geq |B| - 1$ for every room set. When $m < h$, the first movement leaves at least $2h - m > m$ rooms, excluding two-day capture. For $h \leq m < 3n$, inspect the entire minority physical parity on each of two consecutive days. One day requires all $3n$ rooms. Together with the four-cycle obstruction and the already treated lengths one and two, this completes every three-row parameter case.

9 Four rows: counting productive inspections

Let $G = P_4 \square P_n$ with $n \geq 4$, and put $h = 2n$. We prove (4). Smaller lengths are the path, two-row, and three-row families after exchanging coordinates. The key simplification is that all surplus-one survivor sets have exceptional cardinalities.

9.1 A directed matching contraction

Match rows zero and one, and rows two and three, in each column. Identify one parity with two layers of matched pairs. Contracting an edge followed by the inverse matching gives a directed graph with loops, bidirectional horizontal edges in both layers, and one directed vertical rung in each column. The rung directions alternate with the column. The other physical parity reverses the rungs. As before,

$$|N_G(R)| - |R| = |N_D^+(R) \setminus R|.$$

Both rung directions occur, so D is strongly connected.

Lemma 9.1. *Every nonempty proper parity set has positive surplus. If its surplus is one, its size is in $\{1, h-2, h-1\}$. Consequently every parity set R satisfies $|N(R)| \geq g_h(|R|)$, where*

$$g_h(k) = \begin{cases} 0, & k = 0, \\ 2, & k = 1, \\ h-1, & k = h-2, \\ \min\{h, k+2\}, & \text{otherwise.} \end{cases} \quad (35)$$

Proof. If the outside boundary consists of a single vertex v , the survivor set is outgoing-closed in $D - v$. Deleting v splits one horizontal layer into two segments while the other layer stays intact. Every segment of length at least two contains both rung directions and is strongly connected to the intact layer. These vertices form a single core. A segment of length one can only be an endpoint singleton; there cannot be two such segments because $n \geq 4$. If no singleton remains, the graph after deletion is strongly connected, so a nonempty closed set has size $h-1$. Otherwise the singleton is joined to the core by one directed rung. It is either a sink, allowing closed-set sizes $1, h-1$, or a source, allowing sizes $h-2, h-1$.

This proves the surplus-one classification. Strong connectivity gives positive surplus for all other nonempty proper sets, hence surplus at least two there. A full parity has full neighborhood, establishing (35). The exceptional values are attainable: use a corner singleton, the complement of the two neighbors of an opposite corner, or a whole parity minus one vertex. Attainment of the bulk values follows from the diagonal prefixes used below. $\square$

Two inspections cannot succeed. A cohort of size at least $h-1$ leaves at least $h-3$ survivors, whose neighborhood still has at least $h-1$ vertices by (35). Starting from h , this barrier persists for both cohorts.

9.2 The three-inspection endpoint

Give a cohort of size b the weight $w(b) = \max\{b-3, 0\}$. With zero or one useful probe this weight cannot decrease, by (35). With two probes it decreases by at most one; the only positive possibility is a full cohort becoming size $h-1$. With three probes it again decreases by at most one: survivor

sizes from two through $h - 3$ gain at least two neighbors, and the near-full cases satisfy the same bound directly. If only one survivor remains, the old size is four and the old weight is one.

Every nonempty support after movement has at least two vertices, since the grid has minimum degree two. Thus three probes cannot eliminate both cohorts on the same day. A successful schedule has two distinct elimination days. On either day the eliminated cohort already had weight zero; the other receives at most one useful probe and cannot lose weight. Every other day reduces total weight by at most one. Since the initial total is $2(h - 3)$, a T -day schedule has $2(h - 3) \leq T - 2$, or

$$T_3(G) \geq 2h - 4. \quad (36)$$

To attain this, use diagonal prefixes, ordered by coordinate sum and then increasing column. Their neighborhoods are opposite-parity prefixes. Both phases have surplus two at every prefix size $2, \dots, h - 3$. A corner-starting singleton has two neighbors; the other phase's singleton has three. Reflection in the short coordinate exchanges physical parities, permitting either full cohort to choose the corner-starting phase independently.

With three probes daily, a solo prefix cohort decreases from h to four in $h - 4$ rounds. All those survivor sizes lie between two and $h - 3$. Since $h - 4$ is even, the four-vertex support is now in the corner phase. Leave its corner singleton on the next day and inspect its two neighbors on the last. This takes $h - 2$ days per cohort, attaining (36) in $2h - 4$ days.

9.3 A lower game with free holding

For larger budgets, pointwise minimum neighborhoods need not have compatible minimizing shapes. Instead make the lower comparison game strictly easier. Define

$$H_h(b, p) = \min\{b, g_h(\max\{b - p, 0\})\}. \quad (37)$$

The searcher may hold a cohort's count unchanged for free. The profile g_h is nondecreasing, so H_h is nondecreasing in b and nonincreasing in p . Comparing the same useful allocations with arbitrary physical supports, Lemma 9.1 inductively keeps the relaxed counts no larger. Every successful physical schedule therefore clears the relaxed game within the same time.

Replace any inspection which does not decrease a relaxed count by zero probes on that cohort; $H_h(b, 0) = b$ makes this valid. Trim a final allocation to the cohort's actual size. Consider its k remaining *productive inspections*; these need not be consecutive, and may overlap the other cohort's productive inspections. Its count strictly decreases on each, and the last clears it.

Every nonterminal productive transition pays surplus two, with only two possible one-unit savings:

- (A) The first productive inspection uses two probes and sends h to $h - 1$ with surplus one.
- (B) The penultimate productive inspection leaves a singleton, whose neighborhood has size two. The final useful allocation is then exactly two.

There are no others. Survivor size $h - 1$ returns a full support and is unproductive. Survivor size $h - 2$ produces $h - 1$, productive only from the initial count h . A singleton produces two, and the next productive inspection clears it. The profile never produces count three. Let $A, B \in \{0, 1\}$ record the savings and P be the total useful probes on this cohort. Telescoping gives

$$P \geq h + 2(k - 1) - A - B. \quad (38)$$

If either saving occurs, some productive allocation is exactly two; if both occur, the first and last are distinct allocations of two. Hence respectively

$$P \leq (k-1)m + 2, \quad P \leq (k-2)m + 4. \quad (39)$$

Put $D = m - 2$, with $m \geq 3$. By monotonicity, spending m each day is best for a single relaxed cohort. The largest initial count it can clear in $j \geq 1$ inspections is

$$K_j = \min\{h, jD + 2\}. \quad (40)$$

For $j = 1$ this is direct capture. For the induction, if $b \geq m + 2$, its survivor size $b - m$ lies between two and $h - 3$ and the next count is $b - D$. If $b = m + 1$, it leaves a singleton and produces two rooms, already below the next threshold because $K_j \geq m \geq 3$. This proves (40). Thus each cohort needs at least $\lceil (h-2)/D \rceil$ productive inspections, regardless of interleaving.

9.4 The residue lower bound

Assume $3 \leq m < h$, put $W = h - 2$, and write $W = qD + r$, $0 \leq r < D$. The case $m = 3$ is (36), so assume $m \geq 4$.

If $r = 0$, each cohort has $k \geq q$. If $k \geq q + 1$, (38) gives $P \geq qm$ even allowing both savings. If $k = q$, one saving would require $P \geq qm - 1$ but $P \leq qm - m + 2 < qm - 1$; two would require $P \geq qm - 2$ but $P \leq qm - 2m + 4 < qm - 2$. Therefore no savings occur and again $P \geq qm$. The two cohorts require $2qm$ probes, so $T \geq 2q$.

If $r = 1$, each has $k \geq q + 1$. For $k \geq q + 2$, (38) gives $P \geq qm + 3$. For $k = q + 1$, both savings would require $P \geq qm + 1$ and $P \leq qm + 4 - m \leq qm$, a contradiction. At most one occurs, so $P \geq qm + 2$. The combined $2qm + 4$ probes rule out $2q$ days, yielding $T \geq 2q + 1$.

Finally suppose $r \geq 2$. Again $k \geq q + 1$. If $k \geq q + 2$, (38) gives $P \geq qm + r + 2$. If $k = q + 1$, exactly one saving would require $P \geq qm + r + 1 > qm + 2$, contradicting the first inequality in (39). Two would require $P \geq qm + r \geq qm + 2$, contradicting the second, whose right side is $qm + 4 - m \leq qm$. Neither saving is possible, so again $P \geq qm + r + 2$. The combined probe total is at least $2qm + 2r + 4$. A $(2q + 1)$ -day schedule therefore requires $m \geq 2r + 4$; if this fails, $T \geq 2q + 2$. These are exactly the lower cases of (4), proved even with free holding.

9.5 Attaining schedules for every budget

For an actual diagonal prefix, when its survivor size is between two and $h - 3$, movement adds exactly two. A singleton has at most three neighbors, all inspectable next day because $m \geq 3$. Thus the solo capacity calculation (40) also proves an attaining sweep from a full cohort. Two sequential sweeps take $2q$ days for $r = 0$, and $2q + 2$ otherwise.

When $r = 1$, choose the first cohort's initial reflected orientation so that the singleton survivor on its q th full-budget day is a corner. Its next support has two rooms. On day $q + 1$, clear it with two probes and spend $m - 2 = D$ on the second full cohort. Since $W = h - 2$ is even, the remainder one cannot occur with $D = 2$; therefore $D \geq 3$. The second support now has size

$$h - D + 2 = (q - 1)D + 5 \leq qD + 2,$$

so q further full-budget days clear it. This attains $2q + 1$ days.

If $r \geq 2$ and $m \geq 2r + 4$, the first cohort has $r + 2$ rooms after q full-budget days. Clear it on the shared day with $r + 2$ probes and give $m - r - 2 = D - r \geq r + 2 \geq 4$ to the second cohort. This transition lies in the bulk, and its next size is

$$h - (D - r) + 2 = (q - 1)D + 2r + 4 \leq qD + 2.$$

It clears in q more days, again attaining $2q + 1$. Only the explicit compatible prefix families are used; there is no assumption that arbitrary minimum-neighborhood shapes can be chained.

The extreme budgets are immediate. A perfect matching gives $|N(B)| \geq |B|$ for arbitrary mixed-parity room sets. When $m < h$, the first movement leaves at least $2h - m > m$ rooms, ruling out two-day capture. For $h \leq m < 2h$, inspect one fixed physical parity on each of two consecutive days. One day requires all $2h = 4n$ rooms. The two-probe barrier proved above establishes infeasibility for all $m \leq 2$. This completes (4) and supplies an explicit optimal schedule in every feasible case.

10 Five rows

Five rows admit a complete minimum-time classification at every budget. The two parities of the long side require different lower arguments.

Theorem 10.1 (Five rows, three inspections). *For every $n \geq 5$,*

$$T_3(P_5 \square P_n) = 10n - 20. \quad (41)$$

The odd-length proof uses the profiles of Theorem 4.1. The even-length proof uses an exhaustive finite symbolic certificate for arbitrary supports. Both arithmetic potential inequalities are verified in Lean; the complete five-row geometric argument is not presently formalized.

10.1 Odd lengths: exact profiles and a parity potential

Suppose $n \geq 5$ is odd, and put $h = (5n - 1)/2 \geq 12$. The even physical parity has $h + 1$ vertices and the odd physical parity has h . Denote these parities by $z = 0, 1$, respectively. Specializing Theorem 4.1 to short-side parameter $b = 2$ gives $g_0(0) = g_1(0) = 0$ and

$$g_0(k) = \begin{cases} h, & h \leq k \leq h + 1, \\ k + 1, & k = 1 \text{ or } h - 3 \leq k \leq h - 1, \\ k + 2, & 2 \leq k \leq h - 4, \end{cases} \quad (42)$$

$$g_1(k) = \begin{cases} h + 1, & k = h, \\ k + 2, & 1 \leq k \leq 2 \text{ or } h - 2 \leq k \leq h - 1, \\ k + 3, & 3 \leq k \leq h - 3. \end{cases} \quad (43)$$

For example, the complementary corner term in the majority profile is at most one exactly when $h + 1 - k \leq 4$. The minority corner term $Q(k)$ is at most two exactly when $k \leq 2$, and $Q(0) = 1$ handles its full set.

Define

$$W_h(b, z) = \begin{cases} 0, & b = 0, \\ 1, & 1 \leq b \leq 2, \\ b - 2, & 3 \leq b \leq 5, \\ \min\{2b + z - 8, b + h + z - 10\}, & b \geq 6. \end{cases} \quad (44)$$

A reachable nonempty support has at least three vertices in parity zero and at least two in parity one: these are the minimum degrees of vertices in the opposite parities. Thus its size satisfies

$$b = 0 \quad \text{or} \quad 3 - z \leq b \leq h + 1 - z. \quad (45)$$

Lemma 10.2 (Odd-length potential inequality). *For $h \geq 12$, $z \in \{0, 1\}$, a size satisfying (45), and $0 \leq p \leq 3$,*

$$W_h(b, z) \leq W_h(g_z((b - p)_+), 1 - z) + \mathbf{1}_{\{p \geq 2\}}. \quad (46)$$

For fixed h, z , the function $W_h(b, z)$ is nondecreasing in b .

Proof. Substitute (42)–(44), split source and target sizes at 0, 2, 5, select the smaller affine expression for sizes above five, and split the survivor size at the displayed profile endpoints. Every resulting branch is a linear integer inequality under $h \geq 12$, $p \leq 3$, and (45). This complete case split, including truncated subtraction, is checked by `potential_step` and `value_monotone` in `FiveRowOddRankArithmetic.lean`. These proofs use Lean’s ordinary kernel-checked integer-arithmetic tactic and have no admitted cases. $\square$

Take p to be the number of useful probes actually meeting one cohort. Its survivor support has size $b - p$, and its actual next support has at least $g_z(b - p)$ vertices. Monotonicity makes (46) valid for this actual transition. The cohorts occupy disjoint parities each day. Their useful allocations sum to at most three, so at most one receives two or more inspections. The sum of their potentials decreases by at most one per day. Initially

$$W_h(h + 1, 0) = W_h(h, 1) = 2h - 9;$$

at capture both vanish. Thus every strategy takes at least $4h - 18 = 10n - 20$ days.

For an attaining strategy, sweep the initially minority cohort first. Use the compatible diagonal prefixes of Theorem 4.1; each day delete its last three vertices, or all of it if fewer remain. The first two transitions are

$$(h, 1) \longrightarrow (h, 0) \longrightarrow (h - 2, 1).$$

For every minority size $6 \leq b \leq h - 2$, two rounds give

$$(b, 1) \longrightarrow (b, 0) \longrightarrow (b - 1, 1).$$

After $h - 7$ such pairs the support has size five in the minority parity. The final three transitions are

$$(5, 1) \longrightarrow (4, 0) \longrightarrow (2, 1) \longrightarrow 0.$$

The solo duration is $2 + 2(h - 7) + 3 = 2h - 9$, which is odd. The untouched cohort stays the full alternating parity, and is consequently minority when its own sweep begins. Repeat the same sweep, for a total of $4h - 18$ days. For $n = 5$, each solo sweep takes fifteen days.

10.2 Even lengths: a finite symbolic boundary classification

Suppose $n = 2q \geq 6$, and put $h = 5q$. Match horizontal pairs of columns. In row y and matched column j , the physical parity-zero vertex is $(2j + (y \bmod 2), y)$. Use the matching to identify the opposite parity with the same $5 \times q$ array. The resulting directed adjacency has a loop at every vertex, bidirectional vertical edges, leftward horizontal edges on rows 0, 2, 4, and rightward edges on rows 1, 3. The other physical parity reverses the horizontal arrows. For a support R , write

$\sigma(R) = |N(R)| - |R|$. The matching identifies this surplus with its directed outside-boundary size. The directed graph is strongly connected: vertical travel reaches a row of either horizontal orientation. Thus a nonempty proper support has positive surplus.

Represent a column by a subset $B \subseteq \{0, 1, 2, 3, 4\}$, and let $V(B)$ be its vertical neighbors. For successive column masks A, B, C , define

$$c(A, B, C) = \left| \overline{B} \cap (V(B) \cup (A \cap \{1, 3\}) \cup (C \cap \{0, 2, 4\})) \right|. \quad (47)$$

The complement is within the five rows. With zero sentinel columns, the surplus of a word $B_1 \cdots B_q$ is exactly

$$\sum_{j=1}^q c(B_{j-1}, B_j, B_{j+1}).$$

For a bound $s \leq 2$, use states (A, B, e) , $0 \leq e \leq s$. Every $(0, B, 0)$ is initial. For each of the 32 choices of C , retain

$$(A, B, e) \longrightarrow (B, C, e + c(A, B, C))$$

if the new weight is at most s . A state is terminal for boundary exactly s when $e + c(A, B, 0) = s$. Explore every reachable state. Remove only self-loops, verifying that each repeats an empty or full column at zero additional weight. The remaining graph is acyclic. Its complete enumeration gives

	boundary one	boundary two
Reduced terminal paths	59	1057
Maximum reduced word length	5	8

At bound two the reduced reachable graph has 1124 states and 3874 edges. These counts follow from the stated finite rule; they are not a cutoff on physical grid length.

Every accepted word reduces to one of these terminal paths by deleting loop repetitions. Conversely, arbitrary nonnegative repetitions at the recorded positions preserve its boundary. A reduced word with k present and a absent cells therefore represents counts $k + 5u$ and $a + 5v$, where $u, v \geq 0$ count allowed full and empty repetitions. A parameter is zero if its loop type is absent. Repetition preserves whether each row is a prefix.

Lemma 10.3 (Five-row small-boundary geometry). *For $q \geq 3$, in the directed orientation just specified:*

1. *A surplus-one support has size 1, $h - 2$, or $h - 1$.*
2. *A surplus-two support of size $3 \leq k \leq h - 5$ has left-prefix rows.*
3. *A surplus-two support with prefix rows and an empty row has at most five vertices.*
4. *A surplus-two prefix support of size four, five, or six has no neighbor in the last column.*

Finite symbolic verification. Enumerate every terminal path in the specified finite graph, recording every removed-loop position. For a reduced word let k, a be its present and absent counts, and let F, E indicate a full or empty repetition position. Discard only words shorter than three columns with no repetition positions. Every remaining word satisfies these checks:

1. At boundary one, either $k = 1$ and F is false, or $a \in \{1, 2\}$ and E is false.

2. At boundary two, F implies $k \geq 10$. If $(k \geq 3$ or $F)$ and $(a \geq 5$ or $E)$, all rows are prefixes.
3. At boundary two, prefix rows and an empty row imply that F is false and $k \leq 5$.
4. At boundary two, prefix rows and $4 \leq k \leq 6$ imply that F is false and

$$B_q \cup V(B_q) \cup (B_{q-1} \cap \{1, 3\}) = \emptyset.$$

These checks cover 59 and 1057 reduced words, respectively. They are implemented directly from (47) in `five_row_boundary_automaton.py`; the full words and repetition positions are included in the accompanying certificate. No interval assumption is made about arbitrary supports.

The first three conclusions persist for all repetition counts by $k \mapsto k + 5u$ and $a \mapsto a + 5v$. In the fourth size range, no full repetition is possible. Empty repetitions occur after the nonempty part of a prefix support, and cannot create a neighbor in the new last column. This proves the assertions for every $q \geq 3$. $\square$

Lemma 10.4 (Incompatible efficient transitions). *Suppose $\sigma(R) = 2$, $4 \leq k = |R| \leq h - 5$, and $R' \subseteq N(R)$ belongs to the opposite physical parity. If*

$$k - 1 \leq |R'| \leq k + 2, \quad |R'| \leq h - 5,$$

then $\sigma(R') \neq 2$.

Proof. If both surpluses were two, the preceding lemma would give left-prefix rows for R and right-prefix rows for R' , since the second parity reverses horizontal arrows. The neighborhood of a left-prefix support also has left-prefix rows. Its size is $k + 2 \leq h - 3$, so some row misses its last cell. The corresponding right-prefix row of R' must be empty. Part 3 gives $|R'| \leq 5$, hence $k \leq 6$. Part 4 now says $N(R)$ misses the entire last column, forcing every right-prefix row of R' to be empty, a contradiction. $\square$

10.3 The even-length historical potential and attaining sweeps

Mark a current cohort by $z = 1$ exactly when its immediately preceding survivor support had surplus two and size in $[4, h - 5]$; otherwise use $z = 0$. A marked support has $6 \leq b \leq h - 3$. Every nonempty reachable support has size at least two. Define

$$V_h(b, z) = \begin{cases} 0, & b = 0, \\ 1, & 1 \leq b \leq 2, \\ b - 2, & 3 \leq b \leq 5, \\ \min\{2b + z - 8, b + h - 10\}, & b \geq 6. \end{cases} \quad (48)$$

Here z records history, unlike the physical-parity variable in (44).

For $p \leq 3$ useful inspections put $r = b - p$, and denote the next state by (b', z') . These alternatives exhaust every physical transition:

1. $r = 0$, giving $(b', z') = (0, 0)$;
2. $r = h$, necessarily $b = h, p = 0$, giving $(h, 0)$;
3. $\sigma = 1$, giving $r \in \{1, h - 2, h - 1\}$ and $(b', z') = (r + 1, 0)$;

4. $\sigma = 2$, giving $b' = r + 2$ and $z' = \mathbf{1}_{\{4 \leq r \leq h-5\}}$; if $z = 1$ and $3 \leq r \leq h - 5$, this is forbidden by Lemma 10.4;
5. $\sigma \geq 3$, giving $b' \geq r + 3$ and $z' = 0$.

In the forbidden case, the previous survivor size is $k = b - 2$ and $r = k + 2 - p$, so $k - 1 \leq r \leq k + 2$ follows from $0 \leq p \leq 3$.

Substitution into (48) gives

$$V_h(b, z) \leq V_h(b', z') + \mathbf{1}_{\{p \geq 2\}}. \quad (49)$$

This is a finite split of linear integer inequalities under $h \geq 15$, $b \leq h$, $b = 0$ or $b \geq 2$, and $z = 1 \Rightarrow 6 \leq b \leq h - 3$. Its universal proof is `FiveRowRankArithmetic.potential_step` in the Lean artifact; the geometric alternatives remain the ordinary argument just given. The two-cohort potential decreases by at most one per day. Initially it is $2V_h(h, 0) = 4h - 20 = 10n - 20$, proving the lower bound.

For the upper bound use ascending $(x + y, x)$ prefixes, choosing a fixed reflected orientation independently for each initial cohort. Their surpluses are as follows; these describe the stated orders, rather than claiming that they simultaneously minimize all even-length profiles:

Phase	survivor sizes	surplus
0	$1, h - 2, h - 1$	1
0	all other proper nonempty sizes	2
1	$h - 1$	1
1	$1, 2, h - 4, h - 3, h - 2$	2
1	$3, \dots, h - 5$	3

Empty and full supports have empty and full neighborhoods. Neighborhoods of prefixes are opposite-phase prefixes. To derive the tables, let $L(d)$ be the length of diagonal $x + y = d$, and set $A(0) = 0$, $A(d + 1) = L(d) - A(d)$. For a prefix with j cells in its final diagonal, the surplus is

$$A(d) + 1 - \mathbf{1}_{\{d \geq 4\}} - \mathbf{1}_{\{\max(0, d-4) + j - 1 = n-1\}}.$$

The diagonal lengths are 1, 2, 3, 4, 5, then five through diagonal $n - 1$, then 4, 3, 2, 1. The recurrence gives $A(d) = \lceil d/2 \rceil$ for $d \leq 4$, $A(d) = 2 + (d \bmod 2)$ for $4 \leq d \leq n - 1$, and final values 2, 2, 1, 1. Substitution gives the tables and the neighborhood-prefix assertion.

Starting in phase zero, three rounds give $h \rightarrow h - 1 \rightarrow h - 2 \rightarrow h - 3$. Then $h - 8$ pairs of a stationary round and a one-cell decrease reach size five. The last three rounds are $5 \rightarrow 4 \rightarrow 2 \rightarrow 0$. The solo duration is

$$3 + 2(h - 8) + 3 = 2h - 10.$$

Reflection $x \mapsto n - 1 - x$ swaps physical parities because n is even, so each initial cohort can start its sweep in favorable phase zero. The untouched cohort remains full until its sweep begins. Two consecutive sweeps take $4h - 20$ days, completing Theorem 10.1.

10.4 Even lengths at every budget

The one-bit obstruction can be strengthened enough to resolve every budget on the even-length boards. We treat these first, then solve the odd-length recurrence explicitly.

Theorem 10.5 (Five rows, even length, every budget). *Let $n \geq 6$ be even and $h = 5n/2$. Budgets one and two are impossible; $T_3 = 4h - 20$; and*

$$T_4 = 2 \left\lceil \frac{2h - 8}{3} \right\rceil. \quad (50)$$

For $5 \leq m < h$, write $2h - 6 = q(2m - 5) + r$, $0 \leq r < 2m - 5$, and set

$$J(0) = 0, \quad J(1) = 2, \quad J(2) = 3, \quad J(3) = J(4) = 4, \quad J(r) = \left\lceil \frac{r + 5}{2} \right\rceil \quad (r \geq 5).$$

Then

$$T_m = \begin{cases} 2q, & r = 0, \\ 2q + 1, & r > 0 \text{ and } 2J(r) \leq m, \\ 2q + 2, & \text{otherwise.} \end{cases} \quad (51)$$

Finally, $T_m = 2$ for $h \leq m < 2h$, and $T_m = 1$ for $m \geq 2h$.

We first strengthen the geometry in Lemma 10.3. Call the two endpoints in the last matched column *outward corners* when considering the reversed directed orientation.

Lemma 10.6 (Endpoint compatibility). *Suppose $\sigma(R) = 2$ and $3 \leq k = |R| \leq h - 5$. If $N(R)$ contains an outward corner, then $k \geq h - 6$. If $R' \subseteq N(R)$ has opposite physical parity, surplus two, and $3 \leq \ell = |R'| \leq h - 5$, then*

$$(k, \ell) \in \{(h - 6, 3), (h - 5, 3), (h - 5, 4)\}. \quad (52)$$

Proof. Write $n = 2s$. For $k \geq 6$, the prefix and empty-row conclusions of Lemma 10.3 imply that every row of R is a nonempty left prefix. Let its lengths be t_i and set $d_i = s - t_i$. The missing lengths in the rows of $N(R)$ are exactly

$$\begin{aligned} e_0 &= \min(d_0, d_1), & e_1 &= \min(d_0, (d_1 - 1)_+, d_2), & e_2 &= \min(d_1, d_2, d_3), \\ e_3 &= \min(d_2, (d_3 - 1)_+, d_4), & e_4 &= \min(d_3, d_4). \end{aligned}$$

The surplus identity is $\sum d_i = \sum e_i + 2$. Splitting the minima gives these three elementary integer consequences:

$$\begin{aligned} e_0 e_4 = 0 &\implies \sum d_i \leq 6, \\ (e_0 = e_1 = e_2 = 0) \text{ or } (e_2 = e_3 = e_4 = 0) &\implies \sum d_i \leq 5, \\ e_0 = e_2 = e_4 = 0 &\implies \sum d_i \leq 2. \end{aligned}$$

Their complete universal arithmetic proofs are the theorems `endpoint`, `three_rows`, and `alternating_rows` in `FiveRowDeficitArithmetic.lean`. The first proves the corner claim.

The rows of R' are right prefixes. Every occupied row must therefore be full in $N(R)$, that is, have $e_i = 0$. Since $|N(R)| \leq h - 3$, some row is not full, so the empty-row conclusion implies $\ell \leq 5$. The surplus-two prefix row vectors of sizes three through five, up to vertical reflection, are

$$\begin{array}{l|l} 3 & (2, 1, 0, 0, 0), (1, 1, 1, 0, 0), (1, 0, 1, 0, 1) \\ 4 & (2, 1, 1, 0, 0), (1, 1, 1, 0, 1) \\ 5 & (2, 1, 1, 0, 1), (1, 1, 1, 1, 1). \end{array}$$

This finite list follows by substituting the row-neighborhood maxima for vectors with total at most five. For $s \geq 6$ the far boundary cannot affect the check; direct substitution for $s = 3, 4, 5$ gives the same list. A size-three vector forces a zero endpoint deficit; a size-four vector forces three consecutive zeros at an end or the stronger alternating zeros; a size-five vector forces alternating zeros. The three inequalities above now give respectively $k \geq h - 6$, $k \geq h - 5$, and the contradiction $k \geq h - 2$.

For the source sizes $3 \leq k \leq 5$, use the same finite list directly. Every neighborhood row has length at most two, hence misses the last column because $s \geq 3$. It contains neither an outward corner nor a nonempty right prefix. This separate check is needed because the displayed deficit formula assumes that every source row is nonempty. $\square$

For this all-budget argument, mark a cohort by $z = 1$ when its previous survivor had surplus two and size in $[3, h - 5]$. Thus $z = 1$ implies $5 \leq b \leq h - 3$, and the previous survivor size was $b - 2$. Every physical transition, for p useful inspections and $a = b - p$, belongs to the following relation:

1. $a = 0$ gives $(0, 0)$, and $a = h$ gives $(h, 0)$;
2. surplus at least three gives any $(b', 0)$ with $a + 3 \leq b' \leq h$;
3. surplus one gives $(a + 1, 0)$ for $a \in \{1, h - 2, h - 1\}$, with $a = 1$ forbidden when $z = 1$ and $b < h - 4$;
4. surplus two gives $(a + 2, \mathbf{1}_{\{3 \leq a \leq h - 5\}})$; when $z = 1$ and $3 \leq a \leq h - 5$, the pair $(b - 2, a)$ must satisfy (52).

The last two restrictions are precisely Lemma 10.6. In particular, the large-surplus branch includes every possible target size, so this is a lower relaxation for arbitrary supports.

10.4.1 Uniform lower bound for $5 \leq m \leq h - 8$

Put $D = 2m - 5$, and for integer $u \geq 0$ define

$$j(u) = \left\lfloor \frac{(u - 7)_+}{D} \right\rfloor, \quad s(u) = ((u - 7)_+ \bmod D) + 1,$$

$$F_m(u) = \left\lfloor \frac{u + 5j(u)}{2} \right\rfloor - \mathbf{1}_{\{j(u) \geq 1, s(u) \in \{1, 2, 4\}\}}. \quad (53)$$

Its residue values give the following properties:

1. F_m is nondecreasing and $F_m(u + 2) - F_m(u) \leq 3$;
2. if $\lfloor u/2 \rfloor \leq m$, then $F_m(u) = \lfloor u/2 \rfloor$;
3. if $p \leq m$ and $\lfloor u/2 \rfloor > p$, then $F_m(u) \leq p + F_m(u - 2p + 5)$.

For completeness, the last inequality without the indicator is the movement-five affine-ammunition inequality of Lemma 6.2. If $2p \leq 5$, monotonicity suffices. Otherwise the argument decreases by $2p - 5 \leq D$, so its quotient decreases by at most one. With no quotient decrease, the uncorrected inequality has slack at least two. With a quotient decrease, a discounted target has positive quotient

and $s(u - 2p + 5) = s(u) + 2(m - p)$. Target remainder one, two, or four therefore forces source remainder one, two, or two/four, respectively. The source is also discounted, proving the inequality. The first two properties follow by substitution, including the first positive quotient and the wrap from remainder D to one.

Set

$$C = \min\{F_m(2h), 2 + F_m(2h - 2), 4 + F_m(2h - 4)\}, \quad \Phi(b, z) = \min\{C, F_m(2b + z)\}.$$

Every permitted transition lowers Φ by at most its useful allocation p . Here are all exceptional cases in that verification. A bulk surplus-two transition from $z = 0$ has target rank $u - 2p + 5$; from $z = 1$ it would require $p \geq h - 7 > m$ by (52). Surplus at least three has target rank at least $u - 2p + 5$. The three properties of F_m therefore handle these branches; a common cap preserves the inequality.

Capture follows from property 2. A surplus-one singleton is impossible from a marked state because it would require $p \geq h - 5$. From an unmarked state, $b \leq m + 1$; the only extra endpoint value is $F_m(2m + 2) = m + 2$, paid by m inspections and the target pair. For surplus two at survivor sizes one or two, states $b \leq m$ follow from cardinality. The four extra source values at $(b, z) = (m + 1, 0), (m + 1, 1), (m + 2, 0), (m + 2, 1)$ are $m + 2, m + 3, m + 4, m + 4$, respectively, and the target size is at most four.

Finally consider survivor sizes $h - 4, \dots, h - 1$. Rank increases need only monotonicity. The sole proper-source rank decrease is unmarked $h - 1 \rightarrow h - 2$ with three inspections, paid by property 1. The improvements from the full state to $h - 1$ with two inspections and to $h - 2$ with four inspections are exactly the last two terms defining C . The full, zero-inspection transition is tautological. This exhausts the relation.

Both initial cohorts have potential C and both final potentials vanish, so $T_m \geq \lceil 2C/m \rceil$. Direct substitution in the quotient and remainder of (51) gives, for $m \geq 6$,

$$C = qm + J(r) - \mathbf{1}_{\{r=2\}}.$$

When $m = 5$, the physical value h is divisible by five, so $r = 4$ and the same identity holds. The otherwise exceptional residues one and three at this budget cannot occur. At $r = 2$, $m \geq 6$ ensures that $\lceil 2C/m \rceil = 2q + 1$; all other residues immediately give the right side of (51).

10.4.2 The four-inspection lower bound

Define

$$H(b, z) = \begin{cases} b, & b \leq 4, \\ 6, & (b, z) = (5, 0), \\ 4 \lfloor (2b + z)/3 \rfloor - 8 + 3\mathbf{1}_{\{(2b+z) \bmod 3=2\}}, & \text{otherwise,} \end{cases}$$

and

$$A_1 = \min\{H(h - 1, 0), 3 + H(h - 2, 0)\}, \quad C_4 = \min\{H(h, 0), 2 + A_1, 4 + H(h - 2, 0)\}.$$

Use potential C_4 at the full state, A_1 at the unmarked state $h - 1$, and H elsewhere. Allocating $p \leq 4$ lowers it by at most p . This assertion is a finite residue split with the unbounded parameter $h \geq 15$: the theorem `potential_step` in `FiveRowFourProbeArithmetic.lean` checks every branch of the physical relaxation above. In this budget range a marked source forbids all surplus-one transitions

and all bulk surplus-two transitions; the endpoint restrictions prove these exclusions. The checker includes all larger targets in the surplus-at-least-three branch.

Writing $h = 3t, 3t + 1, 3t + 2$ gives respectively

$$(A_1, C_4) = (8t - 12, 8t - 10), \quad (8t - 9, 8t - 8), \quad (8t - 5, 8t - 4).$$

For $L = \lceil (2h - 8)/3 \rceil$, this means $C_4 = 4L$ except when $h \equiv 0 \pmod{3}$, where $C_4 = 4L - 2$. The total-ammunition inequality already proves $T_4 \geq 2L$ outside that residue.

In the exceptional residue, $2L - 1$ days would use exactly $2C_4$ probes, so every cohort transition must be tight and every day must use four useful probes. From a full state, the only positive tight allocation is two inspections, producing size $h - 1$. Thus the first day must split $2 + 2$. From an unmarked state of size $h - 1$, every tight allocation requires four inspections. Two such transitions cannot share day two. This proves $T_4 \geq 2L$. The two local tightness assertions are universally checked by `full_tight` and `penultimate_tight` in `FiveRowFourProbeEquality.lean`.

10.4.3 High budgets and explicit attaining schedules

For $h \geq 25$ and $h - 7 \leq m < h$, write $s = h - m \in \{1, \dots, 7\}$. Then $q = 1$, $r = 2s - 1$, and the seven thresholds $2J(r)$ are 4, 8, 10, 12, 14, 16, 18, all at most $h - s$. Thus the proposed value is three. The matching forbids two days for every $m < h$: after day one at least $2h - m > m$ positions remain.

Only $h = 15, 20$ require additional lower checks. The universal cardinality profile is

$$g(0) = 0, \quad g(1) = 2, \quad g(k) = k + 2 \quad (2 \leq k \leq h - 3), \quad g(h - 2) = h - 1, \quad g(h - 1) = g(h) = h.$$

Iterate $(a, b) \mapsto (g((a - p)_+), g((b - m + p)_+))$, allowing every $0 \leq p \leq m$, sorting the pairs and retaining Pareto-minimal members. Monotonicity of g justifies every deletion. From (h, h) the required final antichains are

(h, m)	days played	remaining pairs
(15, 8)	3	(0, 11), (2, 10), (4, 9), (7, 8)
(15, 11)	2	(0, 12), (2, 11), (4, 10), (5, 9), (6, 8), (7, 7)
(20, 14)	2	(0, 16), (2, 15), (4, 14), (5, 13), (6, 12), (7, 11), (8, 10), (9, 9).

Every remaining total exceeds its budget, proving lower times five, four, and four. These finite universal allocation checks are kernel checked in `FiveRowHighBudgetArithmetic.lean`. Budget monotonicity covers $h = 15, m = 9, 10$ and $h = 20, m = 13$. All other high-budget lower bounds are three.

For the matching upper bounds, use the previously displayed physical prefix tables. With four inspections, start each solo sweep in phase one. Its first round gives $(h - 2, 0)$. At phase-zero size $b \geq 9$, two rounds give $(b - 2, 1)$ and $(b - 3, 0)$. Repeat until size 7, 8, 6, for h modulo three equal to 0, 1, 2, respectively. The final tails are $7 \rightarrow 5 \rightarrow 3 \rightarrow 0$, $8 \rightarrow 6 \rightarrow 4 \rightarrow 0$, and $6 \rightarrow 4 \rightarrow 0$. The solo duration is L , so two independently oriented sweeps attain $2L$.

For $m \geq 5$, full allocations decrease the bulk rank by $D = 2m - 5$. If $r = 0$, each solo sweep lasts q rounds. If $r > 0$, after q full allocations its remaining size is $J(r)$; substitution in the last two prefix transitions gives the special residues one, two, and four. A fresh cohort can be placed within q further rounds of capture by an initial allocation of $J(r)$. Indeed, the optimal oriented initial rank decreases at allocations $p = 0, 1, 2, 3, 4$ are 0, 0, 1, 2, 4, and the decrease is $2p - 5$ for $5 \leq p < h/2$. At $p = J(r)$ it is at least r , leaving rank at most $qD + 6$, which clears within q further full allocations

by the same prefix table. The shared-round construction only needs $p = J(r) \leq m/2 < h/2$, within this range. Share the finishing and starting rounds exactly when $2J(r) \leq m$; otherwise take two separate solo sweeps. This attains (51).

The impossibility at budgets below three follows from the 4×4 subgrid obstruction. At budget at least h , inspect one full parity then the other cohort's full parity; one day is possible exactly at budget $2h$. Together with Theorem 10.1, these facts complete the proof of Theorem 10.5. The geometry here is an ordinary proof with a finite symbolic certificate; the cited Lean modules verify its universal scalar consequences, not the complete physical five-row classification.

10.5 Odd lengths at every budget

Theorem 10.7 (Five rows, odd length, every budget). *Let $n \geq 5$ be odd and $h = (5n - 1)/2$. Budgets at most two are impossible, $T_3 = 10n - 20$, and*

$$T_4 = 2 \left\lceil \frac{5n - 9}{3} \right\rceil + 2\mathbf{1}_{\{3|n\}}, \quad T_5 = 2n - 2.$$

For $5 \leq m < h$, write $5n - 6 = q(2m - 5) + r$, $0 \leq r < 2m - 5$. Use the function J from Theorem 10.5, and put

$$\delta(r) = \mathbf{1}_{\{r \in \{1, 2, 4\} \text{ or } (r \geq 5 \text{ and } r \text{ odd})\}}.$$

Then

$$T_m = \begin{cases} 2q, & r = 0, \\ 2q + 1, & r > 0 \text{ and } 2J(r) + \delta(r) \leq m, \\ 2q + 2, & \text{otherwise.} \end{cases} \quad (54)$$

Finally, $T_m = 2$ for $h \leq m < 2h + 1$, and $T_m = 1$ for $m \geq 2h + 1$.

All transitions below use the exact profiles (42)–(43). An arbitrary physical search dominates the corresponding count search, and compatible prefixes attain it. Thus a lower proof may either use a nondecreasing potential directly on arbitrary supports, or rule out an equally fast prefix search. No serial-search assumption is needed.

10.5.1 Four inspections

Use the interior function $H(b, z)$ from the four-inspection even-length proof. For $h = 3t, 3t+1, 3t+2$, respectively, put $C(h) = 8t - 8, 8t - 5, 8t - 4$, and define

$$V_h(b, z) = \begin{cases} C(h) - 1, & z = 0, b = h, \\ \min\{H(b, 0), C(h) - 2\}, & z = 0, b = h - 1, \\ \min\{C(h), H(b, z)\}, & \text{otherwise.} \end{cases}$$

For $h \geq 12$, $b = 0$ or $3 - z \leq b \leq h + 1 - z$, and $0 \leq p \leq 4$,

$$V_h(b, z) \leq p + V_h(g_z((b - p)_+), 1 - z). \quad (55)$$

Here is a finite reduction that verifies the assertion for unbounded h . The identities $C(h + 3) = C(h) + 8$ and $H(b + 3, z) = H(b, z) + 8$ for $b \geq 6$ give the same translation for both exceptional caps. Also $g_{h+3,z}(k + 3) = g_{h,z}(k) + 3$ for $k \geq 3$ in the domain. For $h \geq 24, b \geq 13$,

reduce both h, b by three; survivors are at least six before reduction, and both sides of (55) decrease by eight. For $b \leq 12$, reduce only h by three: target sizes are at most fifteen, no high-end clause is reached, and the caps are inactive since $C(h-3) \geq 48 > H(b', z')$ for $b' \leq 15$. Only $12 \leq h \leq 23$ remain. All states and all five allocations in this finite range are checked by `RecurrenceBridgeOddFiveFour.finite_certificate` in Lean. The translations are ordinary proofs, separate from this finite check.

The initial potential sum is $2C(h)$, so summation gives $T_4 \geq \lceil C(h)/2 \rceil$. For attainment start the majority cohort. The first two rounds give $(h+1, 0) \rightarrow (h-2, 1) \rightarrow (h-3, 0)$. From majority size $b \geq 9$, two rounds give $(b-2, 1) \rightarrow (b-3, 0)$. The final majority sizes 6, 7, 8 have tails $6 \rightarrow 4 \rightarrow 0$, $7 \rightarrow 5 \rightarrow 3 \rightarrow 0$, and $8 \rightarrow 6 \rightarrow 4 \rightarrow 0$. The solo durations are $L = 2t-2, 2t-1, 2t-1$, respectively. When L is odd the second cohort starts in majority parity. When L is even, its minority solo has the same duration: the first round reaches majority size $h-1$, followed by the same pairs and the size-eight tail. Thus $2L = \lceil C(h)/2 \rceil$ days attain the formula.

10.5.2 The ammunition bound for larger budgets

Let $m \geq 5$, $D = 2m - 5$, and use $F = F_m$ in (53). In addition to its three proved properties, its residue values give $F(u+1) - F(u) \leq 2$. Define

$$C = \min\{F(2h+1), 4 + F(2h-3)\}, \quad A = \min\{C, F(2h), 3 + F(2h-3)\},$$

$$\Phi(b, z) = \begin{cases} A, & z = 0, b = h, \\ \min\{C, F(2b+z)\}, & \text{otherwise.} \end{cases}$$

The value A lies between $\min\{C, F(2h-2)\}$ and $\min\{C, F(2h)\}$, so Φ is nondecreasing in each parity's count. Every useful allocation $p \leq m$ satisfies

$$\Phi(b, z) \leq p + \Phi(g_z(b-p), 1-z). \quad (56)$$

We give the endpoint verification to specify this lower bound completely. Bulk transitions change the rank $u = 2b + z$ to $u - 2p + 5$ and use the proved inequality for F . Capture uses $F(2b+z) = b$. For small survivors use

$$F(2m+2), F(2m+3), F(2m+4), F(2m+5) = m+2, m+3, m+4, m+4.$$

These values handle the majority singleton and the minority singleton or pair; smaller sources follow from cardinality. At the high end:

1. From the corrected majority size h , allocations zero and one reach full minority. Allocation two uses $F(2h) \leq 2 + F(2h-1)$; allocation three uses $A \leq 3 + F(2h-3)$; larger allocations are bulk or low-end.
2. A corrected target h has minority survivor size $h-2$ or $h-3$. For survivor $h-2$ the source ranks are $2h-3, 2h-1, 2h+1$, and the needed bounds are $F(2h-3) \leq A$, $F(2h-1) \leq 1+A$, and $C \leq 2+A$. They follow from the three terms defining A , the increment bounds, and $C \leq 4 + F(2h-3)$. Survivor $h-3$ requires no stronger inequalities.
3. A proper majority survivor of size $h-1$ or $h-2$ increases rank. At survivor size $h-3$, the only proper-source decrease is $(h-1, 0), p=2 \rightarrow (h-2, 1)$; the one-rank decrease costs at most two.

4. Full-majority allocations zero through two reach full minority. Allocations three and four are paid by $C \leq 3 + F(2h - 1)$ and $C \leq 4 + F(2h - 3)$. Larger allocations are bulk or low-end.
5. The remaining high minority survivor $h - 1$ reaches full majority; survivor $h - 2$ was handled above. Full minority at zero allocation stays full.

Clipping target values at C preserves each inequality because the source value is at most C . This exhausts both profiles.

For the physical residue $h \equiv 2 \pmod{5}$, substitution gives

$$C = F(2h + 1) = qm + J(r), \quad 2h - 5 = qD + r. \quad (57)$$

For $m \geq 6$, shifting the residue by four gives $4 + F(2h - 3) \geq F(2h + 1)$. At $m = 5$ the only possible failure has $r = 3$, whereas physical lengths have $r = 4$; the first quotient creates no additional exception. Both full cohorts have potential C . Consequently

$$T_m \geq \lceil 2C/m \rceil = \begin{cases} 2q, & r = 0, \\ 2q + 1, & r > 0, 2J(r) \leq m, \\ 2q + 2, & \text{otherwise.} \end{cases}$$

10.5.3 When equality costs an extra day

The remaining cases have $m = 2J(r)$ and $\delta(r) = 1$, namely

$$(m, r) = (6, 2), (8, 4), \quad \text{or} \quad m \geq 10 \text{ even, } r = m - 5.$$

At these parameters $2C = (2q + 1)m$. A successful $(2q + 1)$ -day prefix schedule would use the full useful budget every day, and every individual inequality (56) would be tight. Direct substitution gives the positive tight full-state allocations:

Case	Majority full	Minority full
$m = 6, r = 2$	3, 4	5, 6
$m = 8, r = 4$	4	8
even $m \geq 12, r = m - 5$	none	$m/2, \dots, m$

The case $m = 10, r = 5$ cannot occur because $5n - 6 \equiv 4 \pmod{5}$. For clarity, the unbounded row can be checked without a finite cutoff. A majority bulk start with $p \leq m/2$ keeps quotient q and has positive remainder $m + 1 - 2p$, so its cost exceeds C ; with $p > m/2$ the quotient is $q - 1$ and the even remainder is $3m - 4 - 2p \geq 8$, giving cost $C + 1$. A minority start with $p < m/2$ likewise costs more than C ; at equality the remainder is zero, and above it the odd remainder $3m - 5 - 2p \geq 7$ gives cost C . Allocations at most four use the explicit profile endpoints. The two small rows have fixed denominators seven and eleven and follow by their finitely many allocations.

Every positive tight start lowers $j(2b + z)$ from q to $q - 1$. After such a start the potential is below C and cannot return to C along tight moves. Every subsequent noncapture tight move lowers j by exactly one, and a tight capture starts at quotient zero. To verify this, a bulk move without a quotient decrease has uncorrected slack at least two, while its discount can change by at most one. Its quotient decreases by at most one because $2p - 5 \leq D$. Allocations zero, one, and two cannot be tight at a proper bulk state. Small noncapture sources of size at most m have strict

cardinality slack; the larger low-end exceptions have source quotient one and target quotient zero. The only additional high-end tight moves below the cap are

$$\begin{array}{l} m = 6 \mid (h, 0), p = 2 \rightarrow (h - 1, 1); \quad (h, 0), p = 3 \rightarrow (h - 2, 1), \\ m = 8 \mid (h - 1, 0), p = 2 \rightarrow (h - 2, 1); \quad (h, 0), p = 3 \rightarrow (h - 2, 1). \end{array}$$

Their quotient also drops from q to $q - 1$. For $m = 6$, $A = F(2h)$; for $m = 8$, $A = C - 1$ while $F(2h) = C$. In the generic critical case $A = C$, so the corrected state cannot occur after a positive tight move. There is no proper tight move into it. Zero allocation is never tight at a nonempty proper state.

Each cohort is therefore active on exactly $q + 1$ consecutive days. The two intervals must cover all $2q + 1$ days, so they overlap once and start at days zero and q . In the generic case r is odd and q even: both starts would have to be minority, which the opposite initial colors and equal starting-day parity forbid. In the two small cases q is odd. Day zero's full allocation can only go to the minority cohort; the other cohort is also minority when it starts on day q . Its required start is at least five of six probes, or all eight probes. The first cohort's last allocation is at least $m/2$, since its total charge is $qm + m/2$. The shared-day budget is exceeded. This proves the extra day in every case.

10.5.4 A capacity recurrence and attaining schedules

Consider the uncapped low-corner profiles, with surplus two in majority parity except at a singleton, and surplus three in minority parity except at a singleton or pair. The greatest count of phase z clearable by $t \geq 1$ successive full allocations has the exact recurrence and solution

$$K_z(1) = m, \quad K_z(t + 1) = m + K_{1-z}(t) - (2 + z), \quad K_z(t) = \left\lfloor \frac{t(2m - 5) + 6 - z}{2} \right\rfloor. \quad (58)$$

Indeed, each target capacity is at least $m \geq 5$, beyond the small inverse-profile exceptions. Invert the bulk profile and add the first allocation. The finite-board profiles are no larger, so these capacities also give sufficient conditions on the actual board.

Start with the minority cohort and use q full allocations. For $r = 0$ it is captured; then q is odd and the other cohort starts minority too, giving $2q$ days. If $r > 0$, the remaining count is $J(r)$. Immediately before the last movement its survivor count is $r/2$ in minority parity for even r , and $(r + 1)/2$ in majority parity for odd r . Substitution gives precisely the small values $J(1), J(2), J(4)$ and the stated bulk values.

On a shared day the other full cohort is minority when r is even and majority when r is odd. An initial allocation of $J(r) + \delta(r)$ puts it within q further full allocations by (58). For the used minority cases this allocation is four at $r = 2$, five at $r = 4$, and $J(r)$ at even $r \geq 6$. For the used majority cases it is three at $r = 1$, four at $r = 3$, and $\lceil (r + 6)/2 \rceil$ at odd $r \geq 5$. The shared day therefore costs $2J(r) + \delta(r)$, proving the middle branch of (54). These partial allocations are at most $m - J(r) \leq m - 2 < h - 2$, excluding the unlisted near-total deletions. If sharing is unavailable, either solo parity clears in $q + 1$ days: its residual after q full allocations is $J(r)$ or $J(r) + \delta(r)$, both at most m . Two solos prove the last branch. At $m = 5$ the formula has $q = n - 2, r = 4$, giving $T_5 = 2n - 2$. The established feasibility and two-day thresholds finish the theorem.

The independent audit and direct room-coordinate construction are included in the research archive. The finite Lean arithmetic bases do not amount to a complete physical Lean proof of this five-row theorem.

11 Seven rows with five inspections per day

The following low-budget family illustrates how a periodic potential and a short equality obstruction can solve the two-count recurrence explicitly. The proof is an ordinary argument with finite rational certificates; it is not presently a complete physical Lean theorem.

Theorem 11.1. *For every odd $n \geq 7$,*

$$T_5(P_7 \square P_n) = 2 \left\lceil \frac{7n - 16}{3} \right\rceil.$$

Put $H = (7n + 1)/2$, so the two checkerboard classes have sizes $H, H - 1$. Theorem 4.1, with $b = 3$, supplies exact profiles g_0, g_1 and compatible prefixes. Thus $g_p(0) = 0$, and

$$g_0(k) = k + \min\{R(k), 3, R(H - k) - 1\}, \quad g_1(k) = k + \min\{Q(k), 4, Q(H - 1 - k)\}$$

on the respective nonempty domains. We use the profiles for the lower bound on arbitrary supports and the compatible prefixes for construction.

Finite potentials and the exceptional residue. Assign charges

$$c_0 = (0, 0, \frac{1}{2}, \frac{3}{4}, 1, 1), \quad c_1 = (0, 0, \frac{1}{4}, \frac{1}{2}, 1, 1).$$

They satisfy $c_0(r) + c_1(s) \leq 1$ whenever $r + s \leq 5$. For each finite base below, define $f_p(k)$ as the least total charge of a path from (p, k) to an empty state in the one-cohort graph

$$(p, k) \longrightarrow (1 - p, g_p(\max\{k - r, 0\})), \quad \text{edge charge } c_p(r), \quad 0 \leq r \leq 5.$$

All quantities are multiples of $1/4$. The accompanying certificate computes these finite shortest-path distances exactly and verifies their monotonicity, zero terminal values, and every inequality

$$f_p(k) \leq c_p(r) + f_{1-p}(g_p(\max\{k - r, 0\})). \quad (59)$$

The complete base data give the following initial sums $V = f_0(H) + f_1(H - 1)$ and attaining serial times U :

n	7	9	11	13	15	17	19	21	23
V	22	63/2	41	50	119/2	69	78	175/2	97
U	22	32	42	50	60	70	78	88	98

There are 5,778 one-cohort inequalities in these nine certificates. Summing (59) across the current physical colors bounds the total potential loss by one per day. Monotonicity handles neighborhoods larger than the profile minimum. Hence $T_5 \geq \lceil V \rceil$.

For $n = 11, 17, 23$, this leaves one day. If capture occurred in $V = U - 1$ days, every day would have to lose exactly one potential unit. In these three bases the only tight first allocation is three inspections in color zero and two in color one. Its minimum successor counts are $(H - 1, H - 2)$. The second-day losses from that state, for allocations $r = 0, \dots, 5$ in color zero, are

$$\left(\frac{3}{4}, -\frac{1}{4}, \frac{1}{2}, 0, 0, \frac{1}{4}\right). \quad (60)$$

None is one, a contradiction. An actual first successor with larger counts cannot evade this calculation: if its first day is tight, its potential equals that of the minimum successor, while monotonicity makes its next potential at least as large. Filling unused quota is harmless. This proves all nine base lower bounds.

A periodic extension valid at every larger length. Use the three base triples

$$(n_0, H_0, c) = (19, 67, 33), (21, 74, 36), (23, 81, 39).$$

Their exact tables satisfy, for both colors,

$$\begin{aligned} f_p(k+3) &= f_p(k) + 2, & c-10 \leq k \leq c+10, \\ g_p(k) &= k+3+p, & c-15 \leq k \leq c+15. \end{aligned} \tag{61}$$

For $n = n_0 + 6j$, write $\Delta = 21j$, so $H = H_0 + \Delta$. Keep f_p unchanged through c . On the inserted interval define

$$\tilde{f}_p(c+3\ell+r) = f_p(c+r) + 2\ell, \quad 1 \leq r \leq 3, \quad 0 < 3\ell+r \leq \Delta.$$

Above $c + \Delta$, put $\tilde{f}_p(k) = f_p(k - \Delta) + 2\Delta/3$. The period identity makes this a nondecreasing extension.

We check (59) by three exhaustive source ranges. For $k \leq c - 5$, every output is at most $c - 1$ and the distant upper corner has no effect; the old inequality is unchanged. For $k \geq c + \Delta + 5$, subtract Δ from source and output. The lower corner is saturated, and both profile and potential translate exactly. In the remaining range $c - 4 \leq k \leq c + \Delta + 4$, residuals are at least $c - 9$ and outputs are at most $c + \Delta + 8$. All relevant profiles are on their plateau. Reduce the source modulo three to a representative in $\{c, c+1, c+2\}$. The output changes by the same multiple of three, and both potentials change by the same multiple of two. The ten-place collar in (61) contains all needed old values, so this too is a verified base inequality. The explicit profile formulas justify the corner invariance and translation at every Δ .

The initial sum increases by $4\Delta/3 = 28j$. For the residue $n \equiv 5 \pmod{6}$, the first two transitions remain in the translated upper part, so the tight first allocation and the loss table (60) are unchanged. The extra-day obstruction therefore persists. Since the stated formula also increases by $28j$ when n increases by $6j$, all lower bounds follow. The six smaller bases cover the lengths below these three starting points.

An attaining physical strategy. Use the compatible prefixes from Theorem 4.1. Allocate five inspections to one cohort until it is captured, assigning any spare quota on its last day to the other cohort, and then finish that cohort. The base times in the table are direct evaluations of this strategy. At bases $n_0 = 19, 21, 23$, successful first colors are respectively 1, 0, 0. The two active middle visits, recorded as (count, current color), are

n_0	first cohort	second cohort
19	(34, 1)	(34, 1)
21	(37, 1)	(38, 1)
23	(41, 1)	(41, 1)

All cuts exceed five. The companion is consequently globally full at the first visit and empty at the second; the finite certificates check these conditions directly.

On the common plateau, two consecutive five-inspection days reduce an active count by three and restore its color, from either starting color. Insert $2\Delta/3 = 14j$ days at each visit to reduce the translated count from $k + \Delta$ to k . Before a visit, the profile trajectory translates by Δ ; afterwards its lower part is unchanged. The active count decreases under five inspections, so the old segments

stay on their respective sides. The inserted segment lies wholly in the plateau. Its even duration preserves the companion's color and the original last-day spare allocation. Both insertions are therefore compatible physical prefix strategies. Their total added time is $28j$, matching the lower bound and proving Theorem 11.1.

The finite certificates and their independent replay include the full rational base tables, all inequalities, the strict second-day losses, the collar identities, and actual coordinate-neighborhood replays of the nine base searches. The three-range proof and the physical insertion establish every larger length; the result does not rely on extrapolating the finite time table.

12 Explicit formulas at large budgets on every odd rectangle

The exact recurrence of Corollary 4.4 admits a closed solution in a uniform budget range, with no restriction on the length. Let

$$w = 2b + 1 \geq 3, \quad n \geq w \text{ odd}, \quad N = wn = 2h + 1, \quad B = b^2.$$

The majority and minority classes have sizes $K_0 = h + 1$ and $K_1 = h$. All neighborhood profiles and prefixes below are those of Theorem 4.1.

Define $L_0(0) = L_1(0) = 0$ and, for $k > 0$,

$$L_0(k) = k + \min\{R(k), b\}, \quad L_1(k) = k + \min\{Q(k), b + 1\}. \quad (62)$$

These are the corner profiles before the far boundary clips a neighborhood. In particular $g_z(k) \leq L_z(k)$. For $r > 0$ put

$$J(r) = \begin{cases} L_1(r/2), & r \text{ even}, \\ L_0((r+1)/2), & r \text{ odd}, \end{cases} \quad (63)$$

$$\delta(r) = \begin{cases} L_0(r/2 + 1) - L_1(r/2), & r \text{ even}, \\ L_1((r+1)/2) - L_0((r+1)/2), & r \text{ odd}. \end{cases} \quad (64)$$

Set $J(0) = \delta(0) = 0$. The two corner profiles interlace, so $\delta(r) \in \{0, 1\}$. These explicit square-root functions describe the small set left by the first sweep and the inspection cost for starting the second sweep on the same day.

Theorem 12.1 (All odd rectangles above a quadratic budget). *Suppose $m \geq b^2 + 1$. If $m \geq N$, then $T_m(P_w \square P_n) = 1$; if $h \leq m < N$, then $T_m(P_w \square P_n) = 2$. For $m < h$, write*

$$D = 2m - w, \quad N - w - 1 = qD + r, \quad 0 \leq r < D.$$

Then

$$T_m(P_w \square P_n) = \begin{cases} 2q, & r = 0, \\ 2q + 1, & r > 0 \text{ and } 2J(r) + \delta(r) \leq m, \\ 2q + 2, & r > 0 \text{ and } 2J(r) + \delta(r) > m. \end{cases} \quad (65)$$

The constructions use two prefix sweeps, with at most one shared day. The lower bound applies to arbitrary inspection schedules.

Corollary 12.2 (A formula using only integer division). *For $m \geq 2b^2 + 1$ and $m < h$, put $c_m = m - w - 1 + \mathbf{1}_{\{m=2b^2+1\}}$. With the same q, r ,*

$$T_m(P_w \square P_n) = \begin{cases} 2q, & r = 0, \\ 2q + 1, & 1 \leq r \leq c_m, \\ 2q + 2, & r > c_m. \end{cases}$$

12.1 An inspection potential and a solved partial-state recurrence

Write $\rho_z(k) = L_z(k) - k$, with $\rho_z(0) = 0$. For $k > 0$,

$$0 \leq L_0(k+1) - L_1(k) \leq 1, \quad 0 \leq L_1(k) - L_0(k) \leq 1.$$

Both L_z and their surpluses ρ_z are nondecreasing. Thus J is nondecreasing, $J(r+2) \geq J(r) + 1$, and $J(1) = 2$. Throughout the proof we assume $m \geq B + 1$. Then $D > 0$ and $J(D) = m$.

Assign a nonempty parity- z count k the rank $u = 2k + z$. Define

$$F(u) = \begin{cases} \lfloor u/2 \rfloor, & \lfloor u/2 \rfloor \leq m, \\ am + J(s), & u - w - 1 = aD + s, \quad 0 \leq s < D, \quad \lfloor u/2 \rfloor > m. \end{cases} \quad (66)$$

In particular $F(0) = F(1) = 0$. Empty supports always have potential zero. The function measures a lower bound on the total number of inspections still required against one cohort.

Lemma 12.3 (Interior inequality and its equality cases). *The function F is nondecreasing, $F(u) \geq \lfloor u/2 \rfloor$, and $F(u+2) \geq F(u) + 1$. For $0 \leq p \leq m$,*

$$F(2k + z) \leq p + F(2L_z((k - p)_+) + 1 - z). \quad (67)$$

Put $j(u) = \lfloor (u - w - 2)_+ / D \rfloor$. Every noncapture equality in (67) lowers j by exactly one. A capture has source quotient zero.

Proof. For $j \geq 1$, use the positive-remainder convention $u = w + 1 + jD + s$, $1 \leq s \leq D$. Then

$$F(u) = jm + J(s) = \left\lfloor \frac{u + jw}{2} \right\rfloor - \epsilon(s), \quad 0 \leq \epsilon(s) = \left\lfloor \frac{s + w + 1}{2} \right\rfloor - J(s) \leq b - 1.$$

At $j = 0$, $F(u) = \lfloor u/2 \rfloor$, without a correction. These formulas prove the stated monotonicity properties, including the period boundaries. A capturable source has potential k . Every nonempty survivor has $L_z(k - p) > k - p$, so a noncapture transition from $k \leq m$ is strict.

Suppose $k > m$ and put $s = k - p > 0$. On the plateau of L_z , the next rank is $v = u - 2p + w$. If $p \leq b$, then $v \geq u + 1$: monotonicity makes the inequality strict for $p > 0$, while $p = 0$ is strict because $v \geq u + 2$. If $p > b$, the quotient cannot increase and decreases by at most one. When it stays unchanged, the floor term, after including the p inspections, increases by at least b ; the correction can save at most $b - 1$. When the quotient falls once, its positive remainders obey $t = s + 2(m - p)$, with s, t now denoting those remainders. Hence $J(t) - J(s) \geq m - p$, as required. A target of quotient zero has its actual cardinality as potential, at least the corresponding $J(t)$.

Off the plateau, the survivor size is at most $b(b - 1)$ in phase one or $(b - 1)^2$ in phase zero. Its next count is at most $B < m$. Since $m > B$, the noncapturable source has quotient exactly one. Write $a = k - m > 0$ and $e = m - p$, so the survivor size is $a + e$. Substitution gives $F(2k + z) = m + L_z(a)$, and $L_z(a + e) - L_z(a) \geq e$ proves the inequality. This step goes from quotient one to zero. Finally a capture source has $k \leq p \leq m$, hence quotient zero. $\square$

Lemma 12.4 (Full-budget capacity). *Using the unbounded profiles L_z , the largest starting count in phase z that clears in $t \geq 1$ consecutive m -inspection days is*

$$A_z(t) = \left\lfloor \frac{tD + w + 1 - z}{2} \right\rfloor.$$

In particular an actual finite-board prefix of rank at most $tD + w + 1$ clears in t such days.

Proof. All queried capacities are at least $m \geq B+1$, beyond the last nonplateau output. The inverse of L_z there is subtraction by $b + z$. Thus $A_z(1) = m$ and $A_z(t+1) = m + A_{1-z}(t) - (b + z)$. The displayed formula solves these two affine recurrences. Finite-board neighborhoods are no larger, and compatible prefixes realize every step. $\square$

12.2 A finite correction for the far corner

For a count k in phase z , let $d = K_z - k$ be its deficit from a full cohort. For a noncapture step with p useful inspections, write $d' = K_{1-z} - g_z(k - p)$. The profiles imply

$$d' \leq d + p. \quad (68)$$

For majority survivors this uses $g_0(s) \geq s - 1$; for nonempty minority survivors it uses the strict Hall bound $g_1(s) \geq s + 1$. The inequality need not hold at capture, which will be treated separately. If $g_z(s) \neq L_z(s)$, the survivor deficit $x = K_z - s = d + p$ is at most B . Define the finite constant

$$C = \min_{\substack{z \in \{0,1\} \\ 0 \leq x \leq B}} [x + F(2g_z(K_z - x) + 1 - z)] \quad (69)$$

and the monotone potential

$$\Phi_z(k) = \min\{F(2k + z), (C - K_z + k)_+\}, \quad \Phi_z(0) = 0. \quad (70)$$

All sets in (69) are nonempty because $n \geq w$. The $x = 0$ cases give $C \leq F(N), F(N + 1)$.

Lemma 12.5 (Boundary envelope). *For every exact-profile transition,*

$$\Phi_z(k) \leq p + \Phi_{1-z}(g_z((k - p)_+)).$$

Consequently $T_m(P_w \square P_n) \geq \lceil 2C/m \rceil$.

Proof. For a noncapture step, (68) gives the inequality against the target's linear branch. If the far corner does not clip, Lemma 12.3 gives the inequality against its F branch. Otherwise $x = d + p \leq B$, and the defining minimum gives $C - d \leq p + F(\text{next rank})$. Thus the source is at most p plus each target branch. At capture, $k \leq p \leq m$ and $\Phi_z(k) \leq F(2k + z) = k \leq p$. Both initial potentials equal C and both terminal potentials are zero. Summing over the two cohorts gives the bound. Monotonicity of Φ and the isoperimetric lower bounds give the same comparison for arbitrary physical supports; no special shape of an uncompressed belief is assumed. $\square$

Lemma 12.6 (Corner correction). *For $m \geq B + 1$ and $m < h$,*

$$F(N) - 1 \leq C \leq F(N).$$

For $m \geq B + b$, the stronger equality $C = F(N)$ holds.

Proof. For $x > 0$, the two full-cohort departures in (69) have next ranks

$$N - d_E(x), \quad d_E(x) = 2(x - R(x)); \quad N - d_O(x), \quad d_O(x) = 2(x - Q(x)) + 1.$$

For $x \leq B$, the high-end term in the finite profile indeed gives these expressions. A negative gain, possible at $x = 1$ in the minority case, is harmless by monotonicity. Every nonnegative gain is at most $2B - 2b < D$, so at most one quotient boundary is crossed.

Without a quotient crossing, the potential drop is at most x . One can use the strict inequalities proved below in (73) when the target remainder is positive. If that remainder is zero, direct substitution gives $J(d_E(x)) \leq x$ and $J(d_O(x)) \leq x$ instead.

It remains to bound a crossing. Put $a = R(x)$ for an E departure, $a = Q(x)$ for an O departure, and $d = b - a$. Let Δ be $F(N)$ minus the departure cost. The following table writes its periodic endpoint expression; using the actual cardinality branch of F can only decrease Δ . Here the source remainder is $2k$ or $2k + 1$.

Case	Target argument ℓ	c	Δ
E , even	$m - b - x + a + k$	$\rho_1(k)$	$d + c - \rho_0(\ell)$
E , odd	$m - b - x + a + k$	$\rho_0(k + 1)$	$d + 1 + c - \rho_1(\ell)$
O , even	$m - b - x + a + k - 1$	$\rho_1(k)$	$d + 1 + c - \rho_1(\ell)$
O , odd	$m - b - x + a + k$	$\rho_0(k + 1)$	$d + 1 + c - \rho_0(\ell)$

All target arguments are positive. Suppose first that $c > 0$. If $\Delta \geq 2$, its bound on the target surplus is below that surplus's plateau, so the defining square or consecutive-product inequality applies without clipping. Substitute $x \leq a^2$ for E and $x \leq a(a - 1)$ for O , together with $k \geq (c - 1)(c - 2) + 1$ in an even-source row and $k \geq (c - 1)^2$ in an odd-source row. Write $v = 2d(c - a)$ and $I = \mathbf{1}_{\{a=b+1\}}$. Rearrangement gives the bounds in the middle column below. Assuming only $\Delta \geq 1$ increases the allowed target root by one and gives the last column.

Case	Upper bound on $m - B$ if $\Delta \geq 2$	Upper bound if $\Delta \geq 1$
E , even	$v - 3d - c + 1$	$v - d + c - 2$
E , odd	$v - 2d - c + 1$	$v + c - 1$
O , even	$v - d - b - bI$	$v + d + 2c - b - 2 - bI$
O , odd	$v - b - bI$	$v + 2d + 2c - b - 1 - bI$

The indicator comes from $x \leq B$, which improves $x \leq a(a - 1)$ by b when $a = b + 1$.

For E we have $d \geq 0$ and $c \leq a \leq b$, with $c \geq 2$ in the even row. The middle-column bounds are nonpositive; the last-column bounds are at most $b - 2$ and $b - 1$. For O with $a \leq b$, the same conclusions hold, with last-column bounds at most $b - 2$ and $b - 3$; the odd row has $c \leq a - 1$. For $a = b + 1$, the middle-column bounds become $3 - 2c$ and $2 - 2c$, and the last-column bounds both become -1 . Thus $\Delta \geq 2$ contradicts $m \geq B + 1$, while $\Delta \geq 1$ contradicts $m \geq B + b$.

If $c = 0$, only even-source rows occur, with $k = 0$. The assumption $\Delta \geq 2$ forces $d \geq 3$ and gives, respectively, $m - B \leq -2ad - 3d + 4$ and $m - B \leq -2ad - d - b + 3$, both nonpositive. The assumption $\Delta \geq 1$ forces $d \geq 2$ and gives $m - B \leq -2ad - d + 1$ and $m - B \leq -2ad + d - b + 1$, both below b . This completes both lower bounds on C . The $x = 0$ majority departure always supplies the upper bound $C \leq F(N)$. $\square$

12.3 The attaining sweeps

The upper construction works throughout $m \geq B + 1$. Suppose $m < h$. Start with the minority cohort. Every noncapture full-budget shot decreases its rank by at least D . If $r = 0$, its rank is $qD + w + 1$, so Lemma 12.4 clears it in q days. Here q is odd, since $N - w - 1$ and D are odd. The other untouched cohort is therefore also minority when its q -day sweep starts.

If $r > 0$, after $q - 1$ full shots the first rank is at most $2m + r + 1$. Its next shot leaves at most $J(r)$ rooms: the last survivor has phase one and size at most $r/2$ for even r , or phase zero and size at most $(r + 1)/2$ for odd r . Clear those rooms on day $q + 1$.

On the same day, the other cohort is minority for even r and majority for odd r . To put its next rank at most $qD + w + 1$, create a deficit of $r/2 + 1$ majority rooms in the even case, or $(r + 1)/2$ minority rooms in the odd case. Inspecting the neighbors of a terminal corner prefix achieves this at costs, respectively,

$$L_0(r/2 + 1) = J(r) + \delta(r), \quad L_1((r + 1)/2) = J(r) + \delta(r).$$

These neighborhoods are terminal prefixes too; the remaining initial prefix is a compatible next belief. The next q full shots suffice by Lemma 12.4. Thus a shared day works precisely at the stated sufficient budget $2J(r) + \delta(r) \leq m$. Otherwise either full phase has rank at most $(q + 1)D + w + 1$, because $r + 1 \leq D$, and two $(q + 1)$ -day sweeps give the other upper bound. Early capture during a sweep only reduces its actual inspection costs.

For $m \geq N$ one inspection of the entire board suffices. For $h \leq m < N$, inspect the minority class on day one; after movement the other cohort occupies at most the h minority rooms and clears on day two. One day is impossible when $m < N$.

12.4 The equality case and the missing shared-day inspection

For $m \geq B + b$, Lemmas 12.5 and 12.6 give the lower bound $\lceil 2F(N)/m \rceil$. It matches the construction except when

$$2J(r) = m, \quad \delta(r) = 1. \tag{71}$$

We now rule out $2q + 1$ days in this case. Assume first $b \geq 2$. Then $m > B$, $r \geq 2$, $r < D$, and $C = qm + m/2$. Moreover

$$F(N + 1) = C + 1, \quad C - h = \frac{(q - 1)w + m - r}{2} \geq 2. \tag{72}$$

For the latter inequality, $m - r \geq 4$ for even r and $m - r \geq 3$ for odd r . Every linear branch of Φ is consequently positive on nonempty supports; its zero cutoff cannot introduce a false equality.

We first record the strict inequalities used when a full-cohort corner departure leaves the same positive quotient. For $x > 0$, put $a = R(x)$ in the first two rows and $a = Q(x)$ in the last two. For positive displayed arguments and $x \leq B$,

$$\begin{aligned} x + L_1(k - x + a) &> L_1(k), & E, \ r = 2k, \\ x + L_0(k - x + a) &> L_0(k), & E, \ r = 2k - 1, \\ x + L_0(k - x + a) &> L_1(k), & O, \ r = 2k, \\ x + L_1(k - x + a - 1) &> L_0(k), & O, \ r = 2k - 1. \end{aligned} \tag{73}$$

Here E and O indicate the phase of the full cohort, while the source potential is $qm + J(r)$ in both cases.

To verify these inequalities, write c for the source surplus. In the first row, the cost difference is $a + \rho_1(\text{target}) - c$. It is positive if $a \geq c - 1$. Otherwise the bounds $k \geq (c - 1)(c - 2) + 1$ and $x \leq a^2$ place the target above $(c - a)(c - a - 1)$: their difference is at least $2c(a - 1) - 2a^2 + 3 > 0$. Its surplus is therefore at least $c - a + 1$. For the second row, only $a < c$ needs checking; the target exceeds $(c - a)^2$ by at least $2c(a - 1) - 2a^2 + a + 2 > 0$. For the third row, use $a \geq 2$, $x \leq a(a - 1)$, and $k \geq (c - 1)(c - 2) + 1$; when $a < c$, the target exceeds $(c - a)^2$ by at least $(2a - 3)c - 2a^2 + 2a + 3 > 0$. For the last row, its cost difference is $a - 1 + \rho_1(\text{target}) - c$. When $a < c$, the target exceeds $(c - a + 1)(c - a)$ by at least $(2a - 3)c - 2a^2 + 3a + 1 > 0$. The required surpluses are all at most their plateau caps. The remaining cases follow from the minimum positive surpluses one and two. This proves (73).

Suppose now that a strategy captured in $2q + 1$ days. Its available inspections total $2C$, so every step must be tight in the envelope inequality and every day must use all m useful inspections. A positive full departure has $d' < p$. A proper nonempty source has $d' < d + p$ on every noncapture step: its majority survivors have surplus at least zero, and its minority survivors have surplus at least two. Together with (72), these facts make a target linear branch strictly inefficient. Every tight target after departure therefore uses F , and its potential remains below C .

A positive full departure lowers j exactly once. For $p \leq B$ its rank gain is less than D , and (73) excludes no drop. For $p > B$ there is no corner clipping: a majority departure is strict because its initial F value is $C + 1$, while a tight minority departure has one drop by Lemma 12.3.

Consider a later clipped step, with proper-source deficit d and $x = d + p \leq B$. Its output equals that of a full-cohort x -inspection departure. Its source quotient is either q or $q - 1$, since $D \geq 2B - 1$ and $r \geq 2$. If both source and target quotients equal q , the strict corner inequalities give $p + F(\text{target}) > C - d \geq \Phi(\text{source})$. If the source quotient is $q - 1$, then

$$F(\text{source}) < C - d. \quad (74)$$

For completeness, when $r = 2s$ put $\alpha = \rho_1(s)$, so $m/2 = s + \alpha$; when $r = 2s - 1$ put $\alpha = \rho_0(s)$, so again $m/2 = s + \alpha$. Substitution in F bounds $C - d - F(\text{source})$ below by $\alpha - 1, \alpha, \alpha, \alpha$ for $E/\text{even}, O/\text{even}, E/\text{odd}, O/\text{odd}$, respectively. All are positive. At $q = 1$ the source quotient is zero and only even r occurs; its cardinality potential gives exactly $\alpha - 1$ and α . Thus (74) includes this endpoint. The full-corner lower bound gives $x + F(\text{target}) \geq C$, making this proper transition strict too.

Consequently every tight clipped step goes from q to $q - 1$. Every other tight noncapture step has a single quotient drop by Lemma 12.3, and capture starts at quotient zero. Each cohort therefore has exactly $q + 1$ consecutive active days, from its first positive allocation through capture. Two such intervals covering the $2q + 1$ -day horizon start on days 0 and q and overlap once. The first day spends all m on its only active cohort, which must be minority: a full-majority m -inspection departure is unclipped and strict. The first cohort spends qm before the shared day and therefore needs exactly $C - qm = m/2$ inspections on that final day.

For even $r = 2s$, the other cohort is minority on the shared day. Its necessary quotient drop requires at least $s + 1$ majority rooms to be absent from the next belief. All neighbors of those absent rooms must be inspected, requiring at least $g_0(s + 1) = L_0(s + 1) = J(r) + \delta(r) = m/2 + 1$ inspections. For odd $r = 2s - 1$, the other cohort is majority and the same argument uses s absent minority rooms and $g_1(s) = L_1(s) = m/2 + 1$. These profiles equal L_z : the complementary deficits of the absent sets are at least $m > B$, using $q \geq 1$ and the parity of q (for odd r , q is even and at least two). Thus the shared day needs more than m inspections, a contradiction.

For $b = 1$, the case $m = 2$ has $D = 1$ and $r = 0$, and follows directly from $C = F(N)$ and the

solo construction. At $m \geq 3$ the same equality argument applies; a critical remainder has $m \geq 4$, $r = m - 3 \geq 1$, and the clipped proper-source cases are immediate since $B = 1$: only a majority survivor of deficit one can clip, and its neighborhood is the full minority class, of potential C . Such a return to full is impossible after a tight departure. This proves the formula for every $m \geq B + b$.

12.5 The full range $m \geq B + 1$: a one-unit slack argument

It remains to add $B + 1 \leq m < B + b$. For $b = 1$ this interval is empty; for $b = 2$ its only budget is $m = 5$, covered by Theorem 10.7. Assume $b \geq 3$. The minimum square size and $D \leq 2B - 3$ give $q \geq 2$. Put $\eta = F(N) - C \in \{0, 1\}$. At $r = 0$, the $c = 0$ positive-discount bounds in the proof of Lemma 12.6 are nonpositive, so $\eta = 0$ and the lower time is $2q$.

A positive corner discount necessarily has $\delta(r) = 0$. To see this, suppose $\delta(r) = 1$. In an even-remainder row of the corner table, $c = \rho_1(k) = \rho_0(k + 1)$, so $k \geq (c - 1)^2$, improving the earlier lower bound by $c - 2$. In an odd-remainder row, $\rho_1(k + 1) = c + 1$, so $k \geq c(c - 1)$, an improvement by $c - 1$. The four upper bounds on $m - B$ for a positive discount become

$$v - d, \quad v, \quad v + c - a, \quad v + 2d + c - b, \quad v = 2d(c - a).$$

The first three are nonpositive since $d \geq 0$ and $c \leq a$. The last is at most $-d - 1$, since $c \leq a - 1$. The exceptional last minority band $a = b + 1$ already has upper bound -1 before this improvement. Thus a discount contradicts $m > B$ when $\delta = 1$.

For $r > 0$, Lemma 12.6 gives $C \geq qm + 1$. The lower time is at least $2q + 1$, so only a proposed upper of $2q + 2$ needs attention. Its unresolved cases are precisely

	η	$2J(r) - m$	$\delta(r)$
A	0	0	1
B	1	1	0
C	1	2	0.

We exclude $2q + 1$ days in all three cases.

Exact counters and tight proper tails. By Corollary 4.4, it suffices to study exact profile counters. More explicitly, retain any alleged physical strategy's useful allocations and evolve the minimum-profile counters, capping allocations if the smaller counter has already been reached. Monotonicity keeps these counters below the actual cardinalities, and compatible prefixes attain them with no greater daily budget. They therefore capture by the alleged deadline; pad with zero allocations if capture occurs early. All following equalities use these exact g_z updates.

The nonnegative slack of one cohort step is $p + \Phi(\text{next}) - \Phi(\text{current})$. Add each day's unused budget. Over a $(2q + 1)$ -day successful schedule the total is $(2q + 1)m - 2C$, which is 0, 1, 0 in cases A, B, C , respectively. In A and B ,

$$C - h = \frac{(q - 1)w + m + (2J - m) - 2\eta - r}{2} > 1.$$

Indeed $m - r \geq 3$, $q \geq 2$, and $w \geq 7$ suffice. Every nonempty linear branch $C - d$ is thus positive and strictly exceeds the support size.

Every tight noncapture step from a proper nonempty support is positive, remains proper, and lowers j by one. Here are the adjustments needed to extend the preceding argument to the present

budgets. A linear target is still strict because $d' < d + p$. An unclipped tight step has one quotient drop by Lemma 12.3. A return to full is strict because its potential C exceeds that of the source.

For a positive clipped step, $d + p \leq B$, so $d \leq B - 1$. Its source quotient is q or $q - 1$, because $D + r > 2B - 2$. For $b \geq 4$, in case A the even and odd bounds are, respectively, $D + r \geq 3B - 4b$ and $D + r \geq 3B - 4b + 1$; case B improves them by one. They all exceed $2B - 2$. For $b = 3$, direct substitution into J gives $J(4) = 4$, $J(5) = 5$, $J(6) = J(7) = 6$, $J(8) = 7$, and $\delta(5) = \delta(7) = 1$, $\delta(6) = 0$. Hence the possible critical residues are exactly

	m	r	$D + r$
A	10	5	18
B	11	6	21
C	10	6	19.

In particular the rows A, B needed here both exceed $2B - 2 = 16$. This strict inequality also handles the positive-remainder convention at zero. The target quotient is q or $q - 1$ since every full-corner gain is less than D .

If both source and target have quotient q , the full departure with $x = d + p$ inspections costs strictly more than $F(N)$ in A , by (73), or at least $F(N) > C$ in B . Subtracting d rules out equality. If the source quotient is $q - 1$, write $r = 2s$, $\alpha = \rho_1(s)$, or $r = 2s - 1$, $\alpha = \rho_0(s)$. Substitution gives lower bounds

$$C - d - F(\text{source}) \geq \alpha - \eta - 1, \quad \alpha - \eta, \quad \alpha - \eta, \quad \alpha - \eta$$

in the order E/even , O/even , E/odd , O/odd . They are positive: case A needs only $\alpha \geq 2$ in even parity and $\alpha \geq 1$ in odd parity. In B , $J = (m + 1)/2 \geq 6$ forces even $\alpha \geq 3$ and odd $\alpha \geq 2$. There is no cardinality exception because $q \geq 2$. The full-corner lower bound again rules out equality. Only a transition from q to $q - 1$ remains.

A zero allocation from a proper support is strict directly. A minority neighborhood adds at least two rooms. A majority neighborhood adds at least one, except at deficit one when it returns to the full minority class. Thus F increases strictly in the former cases, and the full potential C is strictly larger in the latter. The linear branch also increases, since $d' < d$. Finally, tight capture begins at quotient zero: its potential can equal the support size only in the cardinality branch. A proper tight tail consequently lasts exactly $j + 1$ consecutive days.

Case A: the phase obstruction persists. Every step and day is tight. The full-departure argument using (73) is unchanged, so each cohort has $q + 1$ consecutive active days. They overlap once, the first cohort is minority, and it requires $J = m/2$ inspections on the shared day. The fresh cohort requires $J + \delta = m/2 + 1$ by the omitted-set argument above. The small-set profiles are uncut: their complementary deficits are $(qD + w)/2$ for even r , when q is odd and at least three, and at least $(qD + w - 1)/2$ for odd r , when q is even and at least two. Both exceed B . Hence the shared day exceeds the budget.

Discounted cases: the first day uses the slack. A tight positive full departure in B or C uses the F target branch, has $p \leq B$, and crosses a quotient boundary. The linear branch is strict by $d' < p$; an uncut departure costs at least $F(N) > C$. The crossing conditions imply

	E departure	O departure
r even	$p \geq J + 1$	$p \geq J$
r odd	$p \geq J$	$p \geq J + 1.$

For even $r = 2s$, an E crossing has $p - R(p) \geq s + 1$ with $R(p) \geq Q(s)$; an O crossing has $p - Q(p) \geq s$ with $Q(p) \geq Q(s)$. For odd $r = 2s - 1$, use $R(s) \leq R(p)$ in the majority case and $R(s) \leq Q(p) - 1$ in the minority case. These prove the displayed bounds, including the surplus caps since $p \leq B$.

The first day's total slack is at least one. Unused budget already contributes one; allocating all m to one cohort costs at least one slack because $m > B$ and $C < F(N)$. If all m are divided between two positive tight departures, their minima sum to at least $2J + 1 > m$, impossible. Case C , which has zero total slack, is excluded.

In case B , day one consumes exactly the single available unit. Every later day uses all m inspections and every later step is tight. If both outputs of day one are proper, their quotients are at most q : a proper output of a full departure has rank at most N . Both tails then finish by day $q + 2 < 2q + 1$, leaving an entirely unused later day, a contradiction.

If exactly one output is proper, the other remains full. Its probes and any unused budget contribute their entire amount to slack, and so total at most one. The proper allocation is therefore $p \geq m - 1 \geq B$. Define its gain here as N minus the target rank (the initial majority rank is $N + 1$). This gain lies between $D - 3$ and D : the majority gain is $2B - 2b$ at $p = B$ and $2p - 2b - 2$ at $p > B$; the minority gain is $2p - 2b - 1$. The opposite corner is on its plateau since $n \geq w$ and $p \leq m < B + b$. Here $r \leq m - 3 < D - 3$ and $r \geq 2$, so this output has quotient $q - 1$. Both outputs cannot remain full, since that would cost slack $m > 1$.

The proper cohort has exactly q consecutive tight days left; a fresh tight departure of the other begins $q + 1$ consecutive days. They cover the remaining $2q$ days and overlap once, since no later day's budget may be unused. Immediately after day one the proper potential is $C - m + 1$. Before the overlap it receives another $q - 1$ full allocations, so its capture requires

$$C - qm + 1 = J$$

inspections. The fresh tight departure requires at least J too. But $2J = m + 1$, again exceeding the budget. This excludes B and proves Theorem 12.1 for every $m \geq B + 1$.

12.6 Removing the root functions

For Corollary 12.2, the function $H(r) = 2J(r) + \delta(r)$ is strictly increasing for $r \geq 1$: its increments alternate between positive increments of L_0 and L_1 . If $m = 2B + 1$ and $b \geq 2$, put $k = b(b - 1)$. Then

$$H(2k) = L_1(k) + L_0(k + 1) = m, \quad H(2k + 1) = L_0(k + 1) + L_1(k + 1) = m + 2.$$

The last admissible remainder is therefore $2k = m - w$. If $m > 2B + 1$, substitution at $r = m - w - 1$ and $r + 1 = m - w$ puts the relevant arguments on the plateaus. The only endpoint is $m = 2B + 2$, where $L_1(B - b) = B$; there $H(m - w - 1) = m - 1$ and $H(m - w) = m + 1$. Otherwise $H(m - w - 1) = m$ and $H(m - w) > m$. For $b = 1$, the cutoffs at $m = 3, 4$ are zero and follow directly from $H(1) = 5$; the other cases use the same plateau calculation.

13 Minimum-budget searches on seven odd widths

The exact profile recurrence can sometimes be reduced further to an affine formula in the rectangle's length. We prove such formulas at the minimum feasible budget for seven widths. The proof uses a finite rational certificate at each width together with an insertion lemma that covers every larger length.

Theorem 13.1. For $w \in \{3, 5, 7, 9, 11, 13, 15\}$ and every odd $n \geq w$,

$$T_{(w+1)/2}(P_w \square P_n) = 2wn - C_w, \quad \begin{array}{c|cccccc} w & 3 & 5 & 7 & 9 & 11 & 13 & 15 \\ \hline C_w & 8 & 20 & 44 & 84 & 136 & 208 & 288. \end{array}$$

The budget $(w+1)/2$ is the minimum feasible budget.

The width-three and width-five entries agree with the separate formulas proved for those families. This theorem does not assert a formula for every odd width or optimality of serial allocation at other budgets. The certificates in this section are checked with exact rational arithmetic; they are not Lean proofs.

13.1 A separable lower-bound certificate

Set $w = 2b + 1$, $m = b + 1$, $D = m + 1$, and let the two color classes have sizes H and $H - 1$. Use the profiles g_0, g_1 from Theorem 4.1. Choose nonnegative rational numbers $c_p(r)$, for $p \in \{0, 1\}$ and $0 \leq r \leq m$, with

$$c_0(r) + c_1(m - r) \leq 1. \quad (75)$$

Suppose nonnegative functions F_p , with $F_p(0) = 0$, satisfy

$$F_p(k) - F_{1-p}(g_p((k - r)_+)) \leq c_p(r) \quad (0 \leq k \leq H - p, \ 0 \leq r \leq m). \quad (76)$$

Here $(x)_+ = \max(x, 0)$, as elsewhere in the search recurrences.

At a current state (u, v) the sum $F_0(u) + F_1(v)$ decreases by at most one per day, by (75). It suffices to consider strategies using the full quota whenever more than m possibilities remain: additional suffix inspections preserve the prefix normal form and cannot hurt. On the final day one can pad the two quotas to sum to m , with the positive-part convention making both survivors empty. Consequently

$$T_m \geq \lceil F_0(H) + F_1(H - 1) \rceil. \quad (77)$$

The functions are indexed by current color, so this argument includes arbitrary interleaving of the two initial cohorts.

For each finite certificate below, $F_p(k)$ is computed as the shortest total charge of a path from (p, k) to an empty cohort, where a quota r costs $c_p(r)$ and sends the cohort to $(1 - p, g_p((k - r)_+))$. Multiplying by the common denominator gives a finite graph with nonnegative integer edge weights. After computing the functions, the verifier checks every inequality (75)–(76) directly with exact fractions.

13.2 Inserting an arbitrarily long affine interval

Lemma 13.2 (Affine insertion). Let $A = b(b + 1)$. Suppose a certificate at majority size H has an integer cut c such that

$$c \geq A + 2D, \quad H - c \geq A + 2D + 2, \quad (78)$$

$$F_p(k) = 2k + \alpha_p \quad (c - 2D \leq k \leq c + 2D, \ p \in \{0, 1\}). \quad (79)$$

For every integer $\Delta \geq 0$, replace H by $H + \Delta$ in the profile formulas and define

$$F'_p(k) = \begin{cases} F_p(k), & k \leq c, \\ F_p(c) + 2(k - c), & c \leq k \leq c + \Delta, \\ F_p(k - \Delta) + 2\Delta, & k \geq c + \Delta. \end{cases} \quad (80)$$

Then F'_p satisfies the same charge inequalities. The initial lower bound (77) increases by exactly 4Δ .

Proof. The pieces of (80) agree at their endpoints. For either the base or enlarged profiles, every successor $j = g_p((k - r)_+)$ satisfies

$$|j - k| \leq D. \quad (81)$$

For a nonempty survivor, its surplus lies between -1 and $b + 1$, and $r \leq m$. For an empty survivor, $k \leq m$ and $j = 0$.

If $k < c - D$, the enlarged profile agrees with the base profile at $q = (k - r)_+$. Indeed, the far-corner terms in both profile formulas exceed their plateau caps by (78). The successor is unchanged and is less than c by (81). Both potentials are unchanged, so the base inequality applies.

If $k > c + \Delta + D$, write $k_0 = k - \Delta$ and $q_0 = k_0 - r$. Since $q_0 > c + D - m = c + 1$, the low-corner terms have reached their plateau caps. The profile formulas therefore give

$$g'_p(q_0 + \Delta) = g_p(q_0) + \Delta.$$

Writing j_0 for the base successor of k_0 , we have $j = j_0 + \Delta$ and $j_0 > c$. Both potentials gain 2Δ , which cancels in their difference.

Finally suppose $c - D \leq k \leq c + \Delta + D$. The survivor $q = k - r$ is positive, with

$$q \geq c - D - m \geq A + 1, \quad H + \Delta - q \geq H - c - D \geq A + D + 2.$$

Both profiles are therefore on their affine plateaus: $j = k - r + \delta_p$, where $\delta_0 = b$ and $\delta_1 = b + 1$. By (81), both k and j lie in $[c - 2D, c + \Delta + 2D]$. Throughout that interval the corresponding potentials are $2k + \alpha_p$ and $2j + \alpha_{1-p}$. Their difference is

$$2r - 2\delta_p + \alpha_p - \alpha_{1-p}.$$

This is exactly the base potential difference at source $k = c$ with the same current color and quota r . Its survivor and successor lie in the same affine profile and potential collar, so the verified base inequality bounds the expression by $c_p(r)$.

These three cases exhaust the enlarged state space. At each full color class (80) adds 2Δ to the potential, giving the stated increase of 4Δ in the integer lower bound. $\square$

The matching upper construction has the same insertion property. A *serial strategy* directs the full budget to one initial cohort until it can be eliminated, uses any unused quota on that finishing day on the other cohort, and then directs the full budget to the other.

Lemma 13.3 (Insertion in a serial strategy). *Under (78), suppose a base serial strategy visits count c in each active cohort before an inspection. An enlarged rectangle with majority size $H + \Delta$ has a strategy taking exactly 4Δ more days.*

Proof. Under a full-budget inspection an active cohort's size never increases, because every neighborhood surplus is at most m . The certified visit to c is therefore reached from above, and the subsequent base trajectory stays at or below c . Above the cut, the enlarged active trajectory is the translation by Δ of the base trajectory, by the high profile translation identity in the preceding proof. The uninspected cohort remains its full current color class. When the first active cohort

reaches $c + \Delta$, insert 2Δ full-budget days. Throughout this inserted interval the profile plateaus give

$$g_0(k - m) = k - 1, \quad g_1(k - m) = k.$$

Every two days reduce its size by one and restore its current color. The inserted block therefore ends at c with the phase of the base schedule unchanged. Follow the base trajectory below the cut.

The first cohort's finishing-day unused quota is unchanged. The second cohort consequently begins its active phase at the translated base state. This state is above the cut: it lies within $m + 1$ of a full color class, whereas (78) places c farther from that edge. Follow its translated trajectory until $c + \Delta$, insert another 2Δ full-budget days, and then finish along the unchanged lower trajectory. Both inserted blocks have even length, so the later phases of the schedule agree with those in the base strategy. $\square$

The two insertion proofs also apply to any other compatible profile family with color-class sizes $H, H - 1$ and a fixed budget m , provided that a corner radius A has the following properties: every surplus is between -1 and m ; the central profiles are $k + m - 1$ and $k + m$; the low profiles are unchanged when the remaining tail exceeds A ; and the high profiles translate when the source count exceeds A . These are exactly the properties used in the displacement, collar, and trajectory arguments. In that formulation one sets $D = m + 1$ and uses the same margin and affine-collar conditions. This observation will also give an exact three-dimensional family below.

13.3 The finite certificates

The certificate parameters are as follows; in each row $\alpha_1 = \alpha_0 + 1$.

w	m	n_0	H	c	α_0
3	2	13	20	10	-4
5	3	13	33	16	-8
7	4	15	53	26	-15
9	5	15	68	34	-26
11	6	17	94	47	-241/6
13	7	19	124	62	-59
15	8	21	158	79	-80

The charge sequences, listed from quota zero to quota m , are

w	c_0	c_1
3	(0, 0, 1)	(0, 0, 1)
5	(0, 0, 1, 1)	(0, 0, 1, 1)
7	(0, 0, 0, 1, 1)	(0, 0, 1, 1, 1)
9	(0, 0, $\frac{1}{3}, \frac{2}{3}, 1, 1$)	(0, 0, $\frac{1}{3}, \frac{2}{3}, 1, 1$)
11	(0, 0, $\frac{1}{6}, \frac{1}{2}, \frac{5}{6}, 1, 1$)	(0, 0, $\frac{1}{6}, \frac{1}{2}, \frac{5}{6}, 1, 1$)
13	(0, 0, $\frac{1}{4}, \frac{1}{2}, \frac{3}{4}, \frac{3}{4}, 1, 1$)	(0, 0, $\frac{1}{4}, \frac{1}{4}, \frac{1}{2}, \frac{3}{4}, 1, 1$)
15	(0, 0, $\frac{1}{7}, \frac{2}{7}, \frac{1}{2}, \frac{5}{7}, \frac{6}{7}, 1, 1$)	(0, 0, $\frac{1}{7}, \frac{2}{7}, \frac{1}{2}, \frac{5}{7}, \frac{6}{7}, 1, 1$)

Proof of Theorem 13.1. The exact verifier constructs the rational potentials for every odd n between w and the displayed n_0 , inclusive: 32 finite rectangles in total. For each it checks all charge

and potential inequalities, and the equality

$$\lceil F_0(H) + F_1(H - 1) \rceil = \text{the displayed serial strategy's length} = 2wn - C_w.$$

At width eleven the unrounded bound is $2wn - 410/3$; its ceiling gives $2wn - 136$. The other six widths give the integer bound directly. At each base length n_0 , the verifier also checks every value in the affine collar (79), the margins (78), and both serial visits to the cut. The first initial cohort is the minority cohort except at widths seven and thirteen, where it is the majority cohort.

For any larger odd n , put $\Delta = w(n - n_0)/2$. The two insertion lemmas increase both the lower bound and the explicit strategy length by $4\Delta = 2w(n - n_0)$. This proves the formula for every larger length, while the finite certificates cover the remaining lengths. Finally, the majority profile's maximum surplus is b , so Corollary 5.6 gives the minimum feasible budget $b + 1 = (w + 1)/2$. $\square$

The exact checking program is `src/odd_minimum_budget_all_lengths.py` in the companion artifact. Its receipt `research/odd-minimum-budget-all-lengths.json` contains all base potential arrays and the checked quantities. The final verifier uses only integer and rational arithmetic. A numerical linear program helped discover the charge sequences, but is neither called nor trusted by this verifier. The finite checks establish the explicit premises of the two insertion lemmas; those lemmas, rather than an extrapolation of a numerical pattern, supply the infinite-length claim.

14 Eventual affine periods in odd cylinders

The interior-crossing argument extends to every fixed all-odd cross-section, in every dimension. The first step is an exact finite description of its two end regions; an affine plateau alone would not justify changing the cylinder's length.

Let Q be a fixed nontrivial box with odd side lengths. Put

$$W = |Q|, \quad S = \sum_i (\text{side}_i - 1), \quad Z = W(S + 1), \quad \beta = (W - 1)/2.$$

Consider $Q \square P_n$ with odd longitudinal length n . The one-point cross-section gives paths, already classified separately. For $n \geq S + 3$, the longitudinal coordinate is longest and is placed last in the order of Theorem 5.1. Write $H = (Wn + 1)/2$ for its majority class size and g_p^n for its exact color- p profile.

14.1 Two finite tables describe both ends

For $v \in Q_p$, let $r_p(v)$ be the rank of $(v, 0)$ in the parity order of the infinite prism $Q \times \{0, 1, \dots\}$. Every such point has weight at most S , so every preceding point has longitudinal coordinate at most S . Its rank is therefore independent of n for $n \geq S + 1$ and can be computed in the single finite slab $Q \square P_{S+1}$. At most $Z = W(S + 1)$ vertices have weight at most S , so $r_p(v) \leq Z$.

Using the transverse cost a_Q from (21), define

$$U_p(k) = \sum_{\substack{v \in Q_p \\ r_p(v) \leq k}} a_Q(v), \quad B_p(k) = |\{v \in Q_p : r_p(v) \leq k\}|.$$

Both tables are constant for $k \geq Z$.

Lemma 14.1 (Stable corner decomposition). *For $n \geq S + 3$ and $k > 0$,*

$$g_p^n(k) = k + \mathbf{1}_{\{p=1\}} + U_p(k) - |Q_p| + B_p(H - p - k), \quad g_p^n(0) = 0. \quad (82)$$

In particular, every surplus is at most $\beta + p$, and

$$k > Z, \quad H - p - k > Z \implies g_p^n(k) = k + \beta + p. \quad (83)$$

Proof. The ambient cost $a(x) - 1$ equals $a_Q(v)$ on the bottom slice, zero on an interior slice, and -1 on the top slice. This includes the origin: its ambient cost minus one is the transverse origin cost. The prefix-cost identity therefore counts the selected bottom weights through $U_p(k)$ and subtracts the number of selected top vertices.

Coordinate complement preserves parity and reverses both weight and lexicographic order. It carries the top slice to the bottom slice. Thus selected top vertices correspond to bottom vertices outside the prefix of size $H - p - k$, and their number is $|Q_p| - B_p(H - p - k)$. This proves (82). The empty prefix is separate because the extra origin-neighbor term requires a nonempty odd prefix.

The transverse full-class cost identity gives $U_p(\infty) = |Q_{1-p}| - \mathbf{1}_{\{p=1\}}$ and $B_p(\infty) = |Q_p|$. Nonnegativity of a_Q gives the asserted upper bound on surplus. Stabilizing both tables gives (83). $\square$

First slice	Every interior slice	Last slice																											
<table border="1"> <tr><td>2</td><td>1</td><td>0</td></tr> <tr><td>2</td><td>1</td><td>0</td></tr> <tr><td>1</td><td>1</td><td>0</td></tr> </table>	2	1	0	2	1	0	1	1	0	<table border="1"> <tr><td>0</td><td>0</td><td>0</td></tr> <tr><td>0</td><td>0</td><td>0</td></tr> <tr><td>0</td><td>0</td><td>0</td></tr> </table>	0	0	0	0	0	0	0	0	0	<table border="1"> <tr><td>-1</td><td>-1</td><td>-1</td></tr> <tr><td>-1</td><td>-1</td><td>-1</td></tr> <tr><td>-1</td><td>-1</td><td>-1</td></tr> </table>	-1	-1	-1	-1	-1	-1	-1	-1	-1
2	1	0																											
2	1	0																											
1	1	0																											
0	0	0																											
0	0	0																											
0	0	0																											
-1	-1	-1																											
-1	-1	-1																											
-1	-1	-1																											

Local contribution to neighborhood surplus in a $3 \times 3 \times n$ odd cylinder.

Figure 4: The local contribution $a(x) - 1$ in a cylinder with 3×3 cross-section: transverse costs on the first slice, zero in interior slices, and -1 on the last slice. The surplus of a nonempty odd prefix also has the additional 1 in (82); that term is outside the local sum.

The decomposition gives the stronger identities needed to change length. For two valid majority sizes H and $H + \delta$, with both longitudinal lengths at least $S + 3$, we have

$$H - p - k > Z \implies g_p^{H+\delta}(k) = g_p^H(k), \quad (84)$$

$$k > Z \implies g_p^{H+\delta}(k + \delta) = g_p^H(k) + \delta. \quad (85)$$

In the first identity the tail table is full; in the second the bottom weight table is full and the tail argument is unchanged. The first also holds at $k = 0$, since both profiles are zero there. Superscripts now indicate majority size rather than longitudinal length.

An all-odd box has a spanning path whose two endpoints lie in its majority color: traverse successive slices alternately forward and backward and induct on dimension. On that odd path a proper majority k -set has at least k neighbors, by omitting an unselected majority vertex and matching selected vertices toward it from both sides. A nonempty minority k -set has at least $k + 1$ neighbors, by its consecutive blocks in the spacing-two path order. The box contains these path edges, so

$$g_0(k) \geq k \ (k < H), \quad g_1(k) \geq k + 1 \ (k > 0), \quad g_0(H) = H - 1. \quad (86)$$

Corollary 14.2 (Eventual feasibility threshold). *If $H \geq 2Z + 3$, then $h(Q \square P_n) = (W + 1)/2$.*

Proof. The size condition implies $n > S + 3$. The majority surplus is at most β and attains β at $k = Z + 1$ by (83). Apply Corollary 5.6. $\square$

This is an eventual threshold; shorter cylinders can need fewer inspections. For example, a $5 \times 5 \times 5$ box needs twelve, whereas a sufficiently long cylinder with 5×5 cross-section needs thirteen.

14.2 The period and an explicit threshold

Fix a feasible eventual budget $m \geq \beta + 1$ and set

$$d = 2m - W > 0, \quad g = \gcd(W, d), \quad \Delta = \text{lcm}(W, d) = Wd/g.$$

For the explicit threshold define

$$\begin{aligned} J &= m + 1, & A &= Z + 1, & R &= \Delta + 2J, \\ B &= 8Z + 6d + 2, & L &= Z + m + 1, & \mathcal{E} &= B + 4L(B + 1), \\ H_* &= R + 2A + 4J + 2 + 2\mathcal{E}(R + 2J + 1). \end{aligned} \quad (87)$$

Theorem 14.3 (Eventual affine period in every odd cylinder). *For every odd n with $(Wn + 1)/2 \geq H_*$,*

$$T_m(Q \square P_{n+2d/g}) = T_m(Q \square P_n) + 4W/g. \quad (88)$$

At the eventual minimum budget $m = (W + 1)/2$, increasing a sufficiently large odd length by two increases the optimal time by exactly $4W$.

The threshold is not intended to be sharp. The proof uses the profile bounds, plateau and stable translations just proved, together with connectedness.

14.3 A potential with a bounded total deficit

For $n \geq S + 3$, write $H = (Wn + 1)/2$ and $M = H - 1$. Define

$$K = 2Z + 2m + 1, \quad \Gamma = (2H - 4Z - 2m - 5)_+, \quad F_p(k) = \min\{\Gamma, (2k + p - K)_+\}.$$

Lemma 14.4 (Uniform potential and time bounds). *For $0 \leq r \leq m$,*

$$F_p(k) - F_{1-p}(g_p((k - r)_+)) \leq (2r - W)_+. \quad (89)$$

Consequently

$$\lceil 2\Gamma/d \rceil \leq T_m(Q \square P_n) \leq 4 \lceil (M - m)_+/d \rceil + 2. \quad (90)$$

Proof. The potential inequality is trivial if $\Gamma = 0$. Otherwise let $q = (k - r)_+$ and $j = g_p(q)$. If $q \leq Z$, then $k \leq Z + m$ and $2k + p - K \leq 0$, so the current potential is zero. If $q \geq H - Z - 1$, then $j \geq q - 1 \geq H - Z - 2$, and

$$2j + (1 - p) - K \geq 2H - 2Z - 4 - K = \Gamma.$$

The successor potential is therefore Γ . In the remaining region both profiles are on their plateaus, so $j = k - r + \beta + p$. The difference of the unclipped affine expressions is $2r - W$. Common monotone, 1-Lipschitz clipping proves (89).

For two quota shares $r, m - r$, the right sides sum to at most $d = 2m - W$. Both full color classes have potential Γ , giving the lower bound by summing daily decreases. For the upper bound, start with the minority cohort of size M . Each pair of full-budget days reduces an uncleared cohort by at least d . Thus $2 \lceil (M - m)_+ / d \rceil + 1$ days suffice. Pad early clearing to this fixed odd phase length. The other initial cohort remains its full current color class and, after the odd phase, is also in the minority color with size M . Repeat the phase to obtain the upper bound. $\square$

For a rectangle $Q = P_w$, $w = 2b + 1$, the explicit profiles of Theorem 4.1 permit the sharper corner radius $Z = b^2$ in the preceding argument. At minimum budget, strengthen its potential slightly as follows, for every odd $n \geq w$. Take $K = 2b^2 + 2b + 2$ and $\Gamma = (2H - 4b^2 - 2b - 6)_+$. In the low region the current potential is zero for $r < m$ and at most one for $r = m$; the high-region and plateau arguments are unchanged. This gives the useful uniform estimate

$$\max\{0, 2wn - 2w^2 + 2w - 10\} \leq T_{(w+1)/2}(P_w \square P_n) \leq 2wn - 2w - 2. \quad (91)$$

Thus its leading term is $2wn$ for every fixed odd width, independently of whether a sharper correction has been determined.

Return to the general cylinder, with $Z = W(S + 1)$ and the potential of Lemma 14.4. Assume $H \geq H_*$, so $\Gamma > 0$ and $M > m$. Fix an optimal prefix strategy, using padded full quotas as in Section 13. Let Φ_t be the sum of its two current-color potentials and put

$$\eta_t = d - (\Phi_t - \Phi_{t+1}) \geq 0.$$

These are integers, and the upper bound in (90) gives

$$\sum_t \eta_t = dT_m - 2\Gamma \leq 8Z + 6d + 2 = B. \quad (92)$$

Here $d \lceil (M - m)/d \rceil \leq M - m + d - 1$ accounts for the rounding. Call a day *efficient* if $\eta_t = 0$, and *bad* otherwise. There are at most B bad days.

An efficient day gives all m inspections to one cohort. Indeed, if both positive-part losses in (89) are positive, their sum is $d - W < d$. If only one is positive but neither share is m , it is at most $d - 2$. Thus on an efficient day the active cohort loses exactly d potential and the inactive cohort's potential remains exactly constant.

14.4 From flat potentials to empty or full cohorts

Zero inspections strictly increase every intermediate potential:

$$0 < F_p(k) < \Gamma \implies F_{1-p}(g_p(k)) > F_p(k). \quad (93)$$

For a proper majority prefix, $g_0(k) \geq k$, so its unclipped expression increases by at least one after the color reversal. A full majority prefix has potential Γ and is excluded. For a nonempty minority prefix, $g_1(k) \geq k + 1$, giving the same conclusion. Clipping preserves strict growth while the old value is below Γ .

After an active cohort loses d on an efficient day, its potential is strictly below Γ . If it is still positive, (93) forces it to stay active on the next efficient day. A switch is possible only when its potential reaches zero. It cannot become active again in the same efficient run, since zero cannot lose d and its inactive value must stay constant. Hence each consecutive efficient run has at most two focused blocks, and there are at most $2(B + 1)$ such blocks in total.

Call a day *clean* when all m inspections target one cohort and the other cohort is actually empty or actually its full current color class. Flat potential alone is insufficient for this conclusion. We now bound the transient days needed to reach actual emptiness or fullness.

If a nonempty proper support X lies in one color of a connected bipartite graph without isolated vertices, then $X \subsetneq N^2(X)$. Inclusion follows by backtracking along an edge. Equality would make X closed under two-step paths; connectivity makes the two-step graph connected within each color and would force X to be full. Therefore an uninspected nonempty proper prefix grows by at least one every two days until it is full.

During an efficient focused block the inactive potential is constantly zero or Γ . At zero, its count is at most $Z + m = L - 1$; if nonempty, it cannot remain in that range for more than $2L$ days. At Γ , its count is at least $H - Z - 2$, hence within L of its full class; it becomes full within $2L$ days. Empty and full cohorts remain so under further uninspected moves. Each efficient block therefore contributes at most $2L$ nonclean days. Counting all bad days as well, the entire optimal search has at most

$$\mathcal{E} = B + 4L(B + 1) \tag{94}$$

nonclean days.

14.5 A protected band and its two individual cuts

Every one-day count change has absolute value at most $J = m + 1$: profile surplus is between -1 and $\beta + 1$, and each quota is at most m . Record both source counts from each nonclean day. A count k forbids integer cuts c with $k \in [c - J, c + R + J]$, at most $R + 2J + 1$ cuts. There are at most $2\mathcal{E}$ records. Restrict candidate cuts to

$$A + 2J \leq c \leq H - R - A - 2J - 2. \tag{95}$$

There are $H - R - 2A - 4J - 1$ candidates. By (87), this exceeds $2\mathcal{E}(R + 2J + 1)$. Choose a cut forbidden by no record.

No nonclean transition can touch the protected interval $[c, c + R]$. A cohort in that interval is therefore active on every day, receiving all m inspections; an inactive clean cohort is empty or full, both outside the interval. Its count never increases. Nor can it cross the interval upwards: a nonclean transition cannot touch it, and a clean active transition does not increase. Both cohorts begin above and finish below the interval.

For each *initial* cohort i , choose its first count x_i at most $c + \Delta + J$. The previous count is larger, so the displacement bound gives

$$c + \Delta < x_i \leq c + \Delta + J. \tag{96}$$

The entire protected interval is inside the profile plateaus, including all relevant survivors. A full-budget step there sends k to $k - m + \beta + p$, so every pair of days reduces its count by exactly d and restores its color. The next

$$\tau = 2\Delta/d$$

days therefore take x_i to $x_i - \Delta \in (c, c + J]$. This duration is even, and the trajectory stays in the protected band. Throughout the block its inactive companion remains empty or full. The two crossing blocks cannot overlap.

Define the individual endpoints

$$\ell_i = x_i - \Delta, \quad u_i = x_i.$$

They need not agree for the two cohorts. This avoids an otherwise real residue problem: when $d > 1$, a trajectory may skip a prescribed count, and the two cohorts need not have matching residues. One common protected band supplies two legitimate individual cuts. Every visit to (ℓ_i, u_i) belongs to the chosen crossing block, because active clean trajectories are nonincreasing and nonclean transitions cannot touch the larger band. Boundary stalls at the minimum budget are harmless; the even block duration preserves the needed phase.

14.6 Removing and enlarging the crossing blocks

Proof of Theorem 14.3. First contract the majority size from H to $H - \Delta$ and delete both crossing blocks of length τ . At each remaining state transform initial cohort i by

$$f_i(k) = \begin{cases} k, & k \leq \ell_i, \\ k - \Delta, & k \geq u_i. \end{cases}$$

No retained state lies strictly between these cases. At the ends of a deleted block, u_i and ℓ_i both map to ℓ_i . Its inactive companion is empty or full at both ends; a full class maps to the contracted full class. Since τ is even, the color labels also match. The two state paths therefore glue.

For a low retained source, its survivor is at most ℓ_i . The contracted far-distance satisfies

$$(H - \Delta) - \ell_i = H - u_i \geq H - (c + R) \geq A + 2J + 2,$$

so the low profile identity (84) applies. Its successor cannot enter the removed interval, by the protected-band argument. For a high retained source, the contracted survivor is at least

$$\ell_i - m > c - m > Z.$$

The stable upper profile identity (85) therefore gives

$$g_p^H(q + \Delta) = g_p^{H-\Delta}(q) + \Delta.$$

These identities verify every retained transition. Quotas remain within the same budget, and unused quota may be wasted. The result is a physical prefix search on length $n - 2\Delta/W$. Its majority size is at least $H_* - \Delta \geq 2Z + 6J + 4$, so its length exceeds $4(S + 1)$ and the longitudinal coordinate remains longest. Thus all stable profile identities still apply. The resulting search is completed in $T_m(Q \square P_n) - 4\Delta/d$ days. Hence

$$T_m(Q \square P_{n-2\Delta/W}) \leq T_m(Q \square P_n) - 4\Delta/d. \quad (97)$$

Conversely, start with an optimal search on any $H \geq H_*$ and its protected blocks. Enlarge H to $H + \Delta$. Keep lower counts at or below ℓ_i unchanged and increase upper counts at or above u_i by Δ . Replace each block from u_i to ℓ_i by a clean block from $u_i + \Delta$ to ℓ_i , lasting $4\Delta/d$ days. The affine profiles realize this trajectory explicitly, with the inactive companion empty or the enlarged full class. The extra τ days per block are even. The same low and high identities verify every other transition. Thus

$$T_m(Q \sqcup P_{n+2\Delta/W}) \leq T_m(Q \sqcup P_n) + 4\Delta/d.$$

Apply (97) at $H + \Delta$ for the reverse inequality. Since $\Delta/W = d/g$ and $\Delta/d = W/g$, this proves (88). $\square$

Let n_0 be the least odd length at or beyond the threshold. For each of the d/g odd residue classes modulo $2d/g$, the exact two-count recurrence determines the optimum and an optimal strategy at one of

$$n_0, n_0 + 2, \dots, n_0 + 2d/g - 2.$$

The theorem then determines every later time in that class and constructs its optimal strategy by insertion. In particular,

$$T_m(Q \sqcup P_n) = \frac{2W}{2m - W} n + O_{Q,m}(1).$$

At the eventual minimum budget there is one offset $C(Q)$, with $T_{(W+1)/2}(Q \sqcup P_n) = 2Wn - C(Q)$ for every sufficiently large odd n . The offset can depend on the transverse shape, not only on its number of rooms. For rectangular cross-sections $Q = P_w$, the stronger bounds in (91) give $2w + 2 \leq C(w) \leq 2w^2 - 2w + 10$. Finding smaller corner recurrences or closed expressions for the eventual offsets remains a further problem.

14.7 An exact three-dimensional family

For one cross-section the finite corner calculation gives the optimum from the shortest nontrivial cylinder, without the conservative threshold in Theorem 14.3.

Theorem 14.5. *For every odd $n \geq 3$, five inspections per day are necessary and sufficient on $P_3 \sqcup P_3 \sqcup P_n$, and*

$$T_5(P_3 \sqcup P_3 \sqcup P_n) = 18n - 36.$$

Proof. Put $H = (9n + 1)/2$. For $Q = P_3 \sqcup P_3$, the ranks of bottom-slice vertices in their respective parity orders, and their transverse costs, are

p	ranks	costs
0	(1, 2, 3, 5, 8)	(2, 1, 1, 0, 0)
1	(1, 2, 4, 6)	(2, 1, 1, 0).

These ranks already stabilize at $n = 3$. The last even bottom point $(2, 2, 0)$ is first in weight layer four, and its earlier even layers have longitudinal coordinate at most two. The last odd bottom point $(1, 2, 0)$ is sixth in its parity order: its predecessors are the three unit vectors and $(2, 1, 0)$, $(2, 0, 1)$. Enlarging the longitudinal path therefore introduces no earlier point in either list.

Let $U_p(k)$ be the cumulative listed cost through rank k , and let $B_p(k)$ count the listed ranks at most k . The cost and complement argument of Lemma 14.1 gives, for $k > 0$,

$$g_0(k) = k + U_0(k) - 5 + B_0(H - k), \quad g_1(k) = k + 1 + U_1(k) - 4 + B_1(H - 1 - k),$$

with $g_0(0) = g_1(0) = 0$. Thus the corner radius is eight, and the central profiles are $k+4$ and $k+5$. The majority surplus is at most four and equals four at $k=3$: $U_0(3) = 4$ and $B_0(H-3) = 5$ for every $H \geq 14$. Corollary 5.6 gives the minimum budget five.

For the exact time use the identical charge sequences

$$c_0 = c_1 = (0, 0, \frac{1}{3}, \frac{2}{3}, 1, 1).$$

They satisfy the daily charge bound (75). The rational shortest-charge potentials described in the preceding section are checked against every transition

$$F_p(k) - F_{1-p}(g_p(\max\{k-r, 0\})) \leq c_p(r), \quad 0 \leq r \leq 5.$$

Together with a minority-first serial strategy, the five finite certificates give

n	3	5	7	9	11
$F_0(H) + F_1(H-1)$	18	54	90	126	162
strategy length	18	54	90	126	162.

Here the strategy spends all available inspections on the first cohort until it is captured, uses any remaining inspections on its finishing day on the second cohort, and then finishes that cohort. The finite verifier checks the physical room neighborhoods and replays these inspection sets, as well as checking all 1,950 rational inequalities.

At the base $n = 11$, take $H = 50$, $c = 25$, and $D = 6$. The exact arrays satisfy

$$F_p(k) = 2k - 14 + p \quad (13 \leq k \leq 37).$$

This is the full collar $[c - 2D, c + 2D]$, and $c \geq 8 + 2D$, $H - c \geq 8 + 2D + 2$. Both active cohorts visit count 25. At the first visit the other cohort is full, since $25 > 5$ places the visit before the first cohort's finishing day. At the second visit the first cohort is empty.

Consequently Lemmas 13.2 and 13.3, in their compatible-profile formulation, apply with corner radius eight. Increasing H by Δ inserts an affine interval in each potential and raises their initial sum by 4Δ . Inserting 2Δ full-budget days at each serial cut realizes the same increase in time. Each pair of inserted days reduces the active count by one and restores its color, while the other cohort stays full or empty. All remaining transitions follow the stable low and translated high profiles above.

For every odd $n \geq 11$, choose $\Delta = 9(n-11)/2$. The matching bounds are then $162 + 4\Delta = 18n - 36$. The finite certificates cover the smaller odd lengths, completing the proof. $\square$

The exact verifier and receipt are included in the companion artifact:

```
src/three_by_three_cylinder_time.py
research/three-by-three-cylinder-time.json
```

The receipt stores the corner tables, charges, finite values, and complete base potential arrays. An independent verifier also checked actual room neighborhoods and schedules at seven lengths, including 13 and 31. These checks establish the finite premises of the insertion proof. The variable-length theorem is an ordinary proof with exact rational certificates; it is not claimed here as a complete physical Lean theorem.

15 Higher-dimensional boxes

The preceding compression theorems work in arbitrary dimension, but do not by themselves provide a closed time formula for every box. We give two unbounded applications and a fully certified three-dimensional example.

15.1 A side of length two

Let H be bipartite, with color function χ , and put $G = H \square P_2$. In either color class of G , there is exactly one vertex above each vertex v of H : its second coordinate is $s - \chi(v)$ modulo two. Under this identification,

$$\pi(N_G(R)) = \pi(R) \cup N_H(\pi(R)).$$

The vertical move preserves the projection, and an H -move changes it to an H -neighbor. This proves the equality for arbitrary sets, including boundary and empty sets.

Suppose H has an order with prefixes I_k such that every k -set has at least $c(k)$ vertices in its closed neighborhood and $I_k \cup N_H(I_k) = I_{c(k)}$. Put $A = |H|$. The two initial-color cohorts then have the exact state transition

$$(a, b) \mapsto (c(\max\{a - p, 0\}), c(\max\{b - m + p, 0\})), \quad 0 \leq p \leq m, \quad (98)$$

starting at (A, A) . A state is capturable on the current day exactly when $a + b \leq m$. Tail inspections of the prefixes attain each step. Conversely the cardinalities in every unrestricted search dominate these counters when the same allocations are used. Thus shortest-path search on $(A + 1)^2$ states returns both the exact time and an actual optimal strategy, using $O((m + 1)(A + 1)^2)$ transitions after the profile is known.

Theorem 15.1. *Under the preceding closed-neighborhood nesting hypothesis,*

$$h(H \square P_2) = 1 + \max_{0 \leq k \leq A} (c(k) - k).$$

Both this feasibility formula and recurrence (98) apply to every Cartesian box having a side of length two.

Proof. Write $B = \max_k (c(k) - k)$. If $m \leq B$, select k with $c(k) \geq k + m$. A cohort of size at least $k + m$ retains at least k possible positions after inspection and has at least $k + m$ after movement. Initially $A \geq k + m$. Since $k > 0$, the invariant prevents capture forever. If $m > B$, a canonical cohort of size $b > m$ has next size at most $b - m + B < b$. Clear it, then clear the other cohort in the same way; the untouched projected cohort stays full.

For Cartesian H , omit length-one factors, list the other side lengths in increasing order, and order vertices by increasing coordinate sum, breaking a tie by putting the larger first differing coordinate first. The classical simplicial isoperimetric theorem gives exactly the closed-neighborhood nesting hypothesis. We use the statement and definitions in Otachi–Suda [13, Theorem 2.5], where the result is attributed to Moghadam and Bollobás–Leader. For the one-vertex empty product take $c(0) = 0, c(1) = 1$. $\square$

For $a, b \geq 2$, the vertex boundary width of $P_a \square P_b$ is $\min(a, b)$, so $h(P_a \square P_b \square P_2) = \min(a, b) + 1$. If exactly one of a, b is one, the remaining nontrivial ladder has threshold two; if both are one, it is a single edge with threshold one. The theorem also covers all hypercubes. Its input is a classical *closed*-neighborhood theorem; it does not assert open-neighborhood nesting on arbitrary even rectangles.

15.2 Exact profile transfer for either longitudinal parity

There is a useful further recurrence even when global minimizing prefixes are unavailable. Let a bipartite transverse graph Q have compatible minimizing parity orders with profiles g_0, g_1 . For a one-color support in $Q \square P_n$, let k_j be its size in column $j = 0, \dots, n-1$, and let p be its global color.

Proposition 15.2. *For each prescribed column-count vector, the exact minimum neighborhood size is*

$$\sum_{j=0}^{n-1} \max\{g_{(p+j) \bmod 2}(k_j), k_{j-1}, k_{j+1}\}, \quad k_{-1} = k_n = 0.$$

This holds for every positive n , including even n .

Proof. In column j , the neighborhood is the union of the transverse neighborhood of that column and the supports in the two adjacent columns. Its size is at least the displayed maximum. Replace every column support by the corresponding transverse prefix. All three sets are now prefixes of the same opposite color order, and their union has size exactly the maximum. These replacements attain all column minima simultaneously. $\square$

Minimize the sum over vectors with $\sum k_j = k$ to obtain the exact global cardinality profile. A transfer state remembers the previous count a , current count b , and cumulative count. Appending c charges $\max\{g(b), a, c\}$ and replaces the pair by (b, c) . Zero counts at the two outside columns give the boundary conditions. Writing $W = |Q|$, this direct recurrence uses $O(n^2 W^4)$ arithmetic operations and $O(n W^3)$ memory once the transverse profiles are known. It applies in particular to every all-odd cross-section.

This computes exact neighborhood minima, not compatible global orders. The minimizing vector can depend on k without being nested. A search using only these cardinality profiles therefore supplies a lower time bound; it is not asserted to attain that bound on arbitrary even cylinders. The accompanying code checks the recurrence against all 524 288 one-color supports of the $3 \times 3 \times 4$ box.

15.3 A finite even cylinder

The transverse compression also permits an exact finite calculation when the longitudinal side is even. Global parity-prefix nesting of the whole box is not required.

Proposition 15.3. *On $P_3 \square P_3 \square P_4$, the optimal capture time with five inspections per day is 36.*

Proof. Compress each 3×3 slice using its compatible parity-prefix orders, and apply the product-lift and strategy-comparison results of the compression section. Every physical strategy is thereby replaced by a strategy no slower whose support in each slice and each current color is a prefix. Conversely, these prefix strategies are physical strategies. Thus optimization on the compressed states is exact.

Index the four slices by $j = 0, 1, 2, 3$ and the current physical colors by $p = 0, 1$. A state has eight counts a_j^p , initially

$$(a^0, a^1) = ((5, 4, 5, 4), (4, 5, 4, 5)).$$

The transverse neighborhood profiles, including the empty prefix, are

$$h_0 = (0, 2, 3, 4, 4, 4), \quad h_1 = (0, 3, 4, 5, 5).$$

After choosing survivor counts $0 \leq b_j^p \leq a_j^p$, the exact next state is

$$\tilde{a}_j^{1-p} = \max\{h_{(p+j) \bmod 2}(b_j^p), b_{j-1}^p, b_{j+1}^p\},$$

where missing end entries are zero. Indeed, the three contributions are prefixes in the same opposite transverse order, so their union has the maximum of their sizes. The inspection cost is $\sum_{p,j}(a_j^p - b_j^p)$.

The finite verifier performs breadth-first search from the displayed initial state. Whenever more than five rooms remain, it enumerates every eight-coordinate deletion vector of total five, applies the displayed recurrence, and retains the first visit to each state. It first reaches a state containing at most five rooms after 35 movements; inspecting those rooms completes capture on day 36. Using the full budget loses no strategy: extra useful inspections only decrease subsequent supports. No heuristic or neighborhood-profile lower-bound pruning is used.

The completed search examines 1,109,364 transitions, expands 1,993 states, and discovers 2,017 states before that first success. The saved predecessor path is independently replayed on the actual 36-room graph, including every intermediate support and the five-room daily budget. Exhaustion of the preceding breadth-first layers proves the lower bound, and the replay proves the upper bound. $\square$

The complete verifier is `src/three_by_three_by_four_exact.py`; its receipt and physical schedule are in `research/three-by-three-by-four-exact.json`. An independent implementation constructed the moves from actual physical neighborhoods and obtained the same breadth-first layers and optimum. This is an exact finite certificate using the proved compression theorem, not a new physical Lean theorem. It concerns this finite even cylinder; no general even-length formula is inferred.

15.4 Two complete examples with even sides

The distinction between a minimum neighborhood size and a sequence of compatible minimizing shapes is visible even on small cubes. The following two tables are exact ordinary computer-assisted theorems. Their lower certificates and physical schedules have been checked by a second implementation, independently of the search that found them.

Theorem 15.4. *The minimum capture times on P_4^3 and $P_3 \square P_4^2$ are as follows.*

$4 \times 4 \times 4$		$3 \times 4 \times 4$	
Daily budget	Minimum days	Daily budget	Minimum days
0–7	∞	0–6	∞
8	40	7	20
9	20	8	14
10	16	9	11
11	12	10	8
12	10	11	7
13–14	8	12–13	6
15	7	14	5
16–17	6	15–19	4
18	5	20–23	3
19–25	4	24–47	2
26–31	3	48 or more	1
32–63	2		
64 or more	1		

Proof. Apply Theorem 3.6 to every equal-length pair, and also compress the length-three coordinate in the second box. The one-color fixed families contain respectively 292 and 1936 sets per color. These families are enumerated without a geometric guess: order vertices by increasing $\sum_i ix_i$, and either omit each vertex or include it when all its legal predecessors are already present. The predecessor moves are (3.6), together with $x \mapsto x - 2e_i$ on an odd axis. This recursively enumerates every ideal exactly once. Direct physical neighborhoods preserve the fixed families.

Minimizing $|N(R)|$ at each cardinality in these families gives the exact unrestricted one-color neighborhood profiles, by compression. They give an impossibility trap below budgets eight and seven, respectively. On the four-cube the profile is the same in both colors and is

$$g = (0, 3, 6, 7, 9, 11, 12, 13, 15, 16, 17, 18, 19, 19, 21, 22, 23, \\ 24, 25, 25, 26, 27, 28, 28, 29, 29, 30, 31, 31, 31, 32, 32, 32).$$

For example, a cohort of size at least 21 retains at least 14 rooms after seven inspections, and then has at least $g(14) = 21$ positions again.

The two-count lower relaxation formed from the exact profiles gives every listed finite lower bound except at budget eight in either box. To verify this statement, enumerate all count pairs, all quota splits, and successive sets of pairs that can reach capture. This calculation is the same monotone finite recurrence used for the three-cube above. Actual room-coordinate schedules attain every listed bound; the verifier updates the complete physical belief set after each inspection and movement. Budget monotonicity extends endpoint schedules across each displayed range.

The two exceptional lower bounds have short certificate descriptions. For each fixed one-color set A let $F(A)$ be the nonnegative integer in the accompanying table. For every fixed survivor $R \subseteq A$, put $p = |A| - |R|$. At budget eight the tables satisfy

$$F(A) \leq c(p) + F(N(R)), \quad 0 \leq p \leq 8, \quad F(\emptyset) = 0. \quad (99)$$

On P_4^3 , the charges at $p = 0, \dots, 8$ are $(0, 0, 0, 0, 1, 2, 2, 2, 2)$; both full-color potentials are 40. Two allocations totaling at most eight have total charge at most two. Telescoping (99) therefore gives $T_8 \geq (40 + 40)/2 = 40$. There are 29 686 inequalities to check. An attaining prefix sweep for each

cohort has sizes

32, 29, 27, 25, 24, 23, 22, 22, 21, 21, 19, 19,
18, 18, 17, 16, 15, 13, 11, 8, 0.

Each sweep takes twenty days; reflection in an even axis gives the favorable starting phase for the second cohort.

On $P_3 \square P_4^2$, take $c(p) = p$ and full-color potential 54. The 654 570 inequalities give $T_8 \geq \lceil 108/8 \rceil = 14$, attained by the stored physical schedule. These certificate checks require only integer inequalities and direct finite neighborhoods. They do not trust the shortest-path procedure that generated the potential tables.

Finally, both boxes have a perfect matching. Its disjoint alternating trajectories require at least half the volume in two days; two successive inspections of one full color attain that threshold. One-day capture requires the entire volume. This proves the remaining endpoint ranges. $\square$

The independent verifier is `src/even_cube_resumed_independent_review.py`; complete certificates, room schedules, and review receipts are in the corresponding `research/even-resume-*` files in the research archive. At eight probes the four-cube’s cardinality relaxation predicts only 32 days, compared with the true 40; on $3 \times 4 \times 4$ it predicts 12 instead of 14. Thus exact neighborhood profiles alone need not determine exact capture times. The shape information retained by the compression is mathematically necessary for these lower arguments.

15.5 The complete $3 \times 3 \times 3$ example

Encode a room $(x, y, z) \in \{0, 1, 2\}^3$ by $9x + 3y + z$. Its color is $x + y + z$ modulo two. The two classes have sizes 14, 13. Their minimum open-neighborhood sizes are the following; a dash indicates a source cardinality larger than the class.

Layer $x = 0$	Layer $x = 1$	Layer $x = 2$																											
<table border="1"> <tr><td>0</td><td>1</td><td>2</td></tr> <tr><td>3</td><td>4</td><td>5</td></tr> <tr><td>6</td><td>7</td><td>8</td></tr> </table>	0	1	2	3	4	5	6	7	8	<table border="1"> <tr><td>9</td><td>10</td><td>11</td></tr> <tr><td>12</td><td>13</td><td>14</td></tr> <tr><td>15</td><td>16</td><td>17</td></tr> </table>	9	10	11	12	13	14	15	16	17	<table border="1"> <tr><td>18</td><td>19</td><td>20</td></tr> <tr><td>21</td><td>22</td><td>23</td></tr> <tr><td>24</td><td>25</td><td>26</td></tr> </table>	18	19	20	21	22	23	24	25	26
0	1	2																											
3	4	5																											
6	7	8																											
9	10	11																											
12	13	14																											
15	16	17																											
18	19	20																											
21	22	23																											
24	25	26																											

First day of the optimal 18-day, five-probe schedule.

Figure 5: The three-cube drawn as three layers. Adjacent rooms in a layer are connected; rooms at the same position in consecutive layers are also connected. Shading marks the first five inspections of the schedule.

k	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
$g_E(k)$	0	3	5	7	8	9	10	10	11	12	12	13	13	13	13
$g_O(k)$	0	4	6	7	9	10	11	12	12	13	13	14	14	14	–

Here is a finite certificate procedure for the table. For each color, list its vertices; recursively either include or exclude the next vertex, maintaining the selected count k and the union U of its physical neighbor sets. At every leaf check $|U| \geq g(k)$. This has $2^{14} + 2^{13} = 24576$ leaves. The generic recursion is proved sound in Lean before the concrete finite check is kernel evaluated. Prefixes in increasing coordinate sum with decreasing lexicographic tie breaking attain the values. Neither the lower bound nor the formal classification assumes that an arbitrary search uses these prefixes.

For a fixed budget define the relaxed count successor

$$f_p(a, b) = (g_O(\max\{b - m + p, 0\}), g_E(\max\{a - p, 0\})).$$

All arbitrary physical successors dominate one of these. The finite calculation below independently specifies the time lower bounds. It uses just $15 \cdot 14 = 210$ count pairs; W_t denotes relaxed states that can reach zero in at most t inspection rounds.

```

W = {(0, 0)}
for t = 1, 2, ...:
  Wnext = W union {(a,b): 0<=a<=14, 0<=b<=13,
                        f_p(a,b) belongs to W for some 0<=p<=m}
  if (14,13) belongs to Wnext: return t
  if Wnext == W: return infinity
W = Wnext

```

The required lower endpoint checks are 18, 10, 6, 4, 3 at budgets 5, 6, 8, 11, 12, respectively. In Lean the corresponding distance potentials are stored as finite integers and checked to be nondecreasing in both counts, zero at zero, and to drop by at most one under every relaxed transition. This proves the lower bound without trusting the search routine that found the potentials. For budgets at most four, there is a shorter proof: every same-color triple has at least seven neighbors, so a cohort of at least seven positions can never fall below seven after four inspections and movement.

For clarity, explicit upper schedules are given next. Each row is a sequence of daily inspection sets; a superscript $\times 2$ means repeat the entire listed sequence twice. The physical update $B \leftarrow N(B \setminus S)$, starting from all 27 rooms, verifies every schedule directly.

m	Inspection sets in order
5	$(\{1, 3, 5, 9, 11\}, \{4, 6, 10, 12, 18\}, \{5, 7, 11, 13, 15\},$ $\{8, 10, 12, 14, 16\}, \{9, 11, 13, 15, 17\}, \{10, 12, 14, 16, 18\},$ $\{11, 13, 15, 19, 21\}, \{8, 14, 16, 20, 22\}, \{15, 17, 21, 23, 25\})^{\times 2}$
6	$(\{8, 14, 16, 20, 22, 26\}, \{5, 7, 11, 13, 19, 25\}, \{2, 4, 10, 16, 22, 24\}, \{1, 7, 13, 15, 19, 21\},$ $\{0, 4, 6, 10, 12, 18\})^{\times 2}$
7	$(\{8, 14, 16, 20, 22, 24, 26\}, \{5, 7, 11, 13, 15, 19, 21\}, \{0, 2, 4, 6, 10, 12, 18\})^{\times 2}$
9	$\{2, 4, 8, 14, 16, 20, 22, 24, 26\}, \{1, 3, 7, 9, 11, 13, 15, 19, 21\},$ $\{5, 7, 11, 13, 15, 17, 19, 23, 25\}, \{0, 2, 4, 6, 10, 12, 18, 22, 24\}$
12	$\{2, 4, 6, 8, 10, 12, 14, 16, 20, 22, 24, 26\}, \{1, 3, 8, 9, 14, 16, 19, 20, 21, 22, 24, 26\},$ $\{1, 3, 5, 7, 9, 11, 13, 15, 19, 21\}$
13	All odd-numbered rooms, on each of two consecutive days.

Larger budgets inherit these schedules. Fewer than 27 probes cannot capture every possible initial room in one day, and 27 can. We have therefore proved the complete table stated in the introduction.

Its geometry, finite inequalities, physical schedules, and interpretation as capture of every actual target walk are all checked in Lean.

15.6 The exact five-inspection time on every $3 \times 3 \times n$ cylinder

Theorem 15.5. *For every $n \geq 3$,*

$$T_5(P_3 \square P_3 \square P_n) = 18n - 36.$$

The proof combines an explicit physical sweep with an ordinary lower certificate for arbitrary supports. The new even-length argument is not presently a complete physical Lean theorem.

The odd case is Theorem 14.5. Assume $n = 2L \geq 4$ and put $h = 9L$. Both checkerboard classes have h rooms. The compatible transverse 3×3 profiles are

$$G_0 = (0, 2, 3, 4, 4, 4), \quad G_1 = (0, 3, 4, 5, 5).$$

Compress every transverse slice to its parity prefix. This operator C is monotone and cardinality preserving, and satisfies $N(CS) \subseteq C(NS)$ by the fiber-compression theorem. For a compressed phase- p set with slice counts k_j , its exact neighborhood count in slice j is

$$\max\{G_{(p+j) \bmod 2}(k_j), k_{j-1}, k_{j+1}\}, \quad k_{-1} = k_n = 0.$$

We first establish a sharp obstruction to consecutive small boundaries.

15.6.1 A finite word certificate for every even length

For phase zero, pair consecutive counts into letters (a, b) with $0 \leq a \leq 5$, $0 \leq b \leq 4$. Give a letter the cost

$$v(a, b) = \max\{G_0(a), b\} - a + \max\{G_1(b), a\} - b,$$

and consecutive letters $s = (a, b)$, $t = (c, d)$ the edge cost

$$\begin{aligned} e(s, t) = & \max\{G_1(b), a, c\} - \max\{G_1(b), a\} \\ & + \max\{G_0(c), d, b\} - \max\{G_0(c), d\}. \end{aligned}$$

The neighborhood surplus is exactly $\sum_i v(s_i) + \sum_{i < L-1} e(s_i, s_{i+1})$; no outside edge is added. Every edge cost is nonnegative. The vertex costs are the following nonnegative matrix, with rows $a = 0, \dots, 5$ and columns $b = 0, \dots, 4$:

$$\begin{pmatrix} 0 & 3 & 4 & 5 & 5 \\ 2 & 3 & 3 & 4 & 4 \\ 3 & 3 & 3 & 3 & 3 \\ 4 & 3 & 3 & 3 & 2 \\ 4 & 3 & 2 & 2 & 1 \\ 4 & 3 & 2 & 1 & 0 \end{pmatrix}.$$

Thus a word of surplus at most four has no prefix of larger cost.

Here is the entire finite certificate specification. A state records its last letter, cost at most four, occupied count capped at four, omitted count capped at eight, whether its first $a \geq 3$, and length capped at two. Initialize with every one-letter word of cost at most four. Appending $t = (c, d)$

adds $v(t) + e(s, t)$ to the cost, $c + d$ to the occupied count, and $9 - c - d$ to the omitted count; saturate the two counts, preserve the first-letter flag, and cap length at two. Discard costs above four. Finite closure gives exactly 113 states and 346 edges. Every terminal state of capped length two satisfies:

1. Its cost is at least zero if either capped count is zero, and otherwise at least $\min\{4, K + 1, \lfloor (D + 2)/2 \rfloor\}$, where K, D are its two capped counts.
2. If its cost is four, $K = 4$, and $D = 8$, its first $a \geq 3$ and its last letter is $(a, 0)$ with $a \leq 2$.

These are finite integer checks of the stated initialization and transition rule, supplied in the accompanying certificate. Induction on word length makes them valid for every $L \geq 2$; costs above four trivially satisfy the first bound. Reflection in the even longitudinal side exchanges the colors, giving phase one too.

Consequently every monochromatic set of size k has at least $g_h(k)$ neighbors, where

$$g_h(0) = 0, \quad g_h(h) = h, \quad g_h(k) = k + \min\{4, k + 1, \lfloor (h - k + 2)/2 \rfloor\} \quad (0 < k < h). \quad (100)$$

Compression transports this bound to arbitrary supports. Call a survivor *critical* if its surplus is four and $4 \leq k \leq h - 8$. A compressed phase-zero critical set has at most two neighbors in the last slice, whereas every compressed phase-one critical set contains at least three rooms in that slice. Hence consecutive compressed critical survivors are impossible. For arbitrary critical R , (100) forces $|N(CR)| = |N(R)| = |R| + 4$. The compression inclusion therefore becomes $N(CR) = C(NR)$. If a second critical S lay in $N(R)$, monotonicity would give $CS \subseteq C(NR) = N(CR)$, the same contradiction. This proves the obstruction for arbitrary consecutive survivors.

15.6.2 A two-component integer potential

Let $e \in \{0, 1\}$ record whether the previous survivor was critical, starting with $e = 0$. A marked count satisfies $8 \leq k \leq h - 4$, and the obstruction forbids $e = e' = 1$. Use charges

$$c(0), \dots, c(5) = (0, 0, 1, 2, 3, 3), \quad c(p) + c(q) \leq 3 \quad \text{if } p + q \leq 5.$$

Define

$$\begin{aligned} P_0(k) &= (0, 0, 1, 2, 3, 3, 5, 6, 9)[k], & 0 \leq k \leq 8, \\ P_0(k) &= \min\{6k - 42, 3k + 3h - 51, 6h - 54\}, & k \geq 9, \\ P_1(k) &= 6k - 39, & 8 \leq k \leq h - 4. \end{aligned}$$

Both functions are nondecreasing in their valid count domains. For every one-cohort step with useful allocation $p \leq 5$,

$$P_e(k) \leq c(p) + P_{e'}(k'). \quad (101)$$

To specify its complete check, put $s = k - p$. At $s = 0$ the only output is $(0, 0)$. If $4 \leq s \leq h - 8$, the minimum marked output is $(s + 4, 1)$ and the minimum unmarked output is $(s + 5, 0)$; the former is forbidden when $e = 1$. Outside this interval the output is unmarked with count at least $g_h(s)$. Monotonicity handles every larger actual neighborhood. These cases include all geometric outputs.

For $h = 18, 27, 36, 45$, substitution checks respectively 183, 345, 507, and 669 inequalities, using integers only. The following three cases prove every larger parameter, so this finite verification is

not an extrapolation. Write $h = 45 + \Delta$, $\Delta \geq 0$. If $k \leq 17$, every minimum output is at most 22 and the base inequality is unchanged. If $k \geq 27 + \Delta$, subtract Δ from source and output counts: the residual is at least $22 + \Delta$, and all profiles, history endpoints, and potentials translate exactly, the latter by 6Δ . Finally, if $18 \leq k \leq 26 + \Delta$, the residual lies in $[13, 26 + \Delta] \subseteq [4, h - 8]$, and source and minimum output are in the affine region $P_e(k) = 6k - 42 + 3e$. The possible drops are $6p - 27$ for $0 \rightarrow 1$ or $1 \rightarrow 0$, and $6p - 30$ for $0 \rightarrow 0$; each is at most $c(p)$. This proves (101) for every $h \geq 45$. The four bases cover all remaining physical half-sizes $h = 9n/2$.

Both initial potentials are $P_0(h) = 6h - 54$, and both terminal potentials vanish. The combined potential decreases by at most three per day. Therefore

$$T_5 \geq \frac{2(6h - 54)}{3} = 4h - 36 = 18n - 36.$$

15.6.3 An explicit sweep attaining the bound

By Lemma 5.2, weight-then-reverse-lexicographic prefixes have prefix neighborhoods even when a box has an even side. We use this nesting only for construction. Let $U_p(k)$ be the cumulative cost through the bottom-slice ranks at most k , and $B_p(k)$ the number of such ranks. The stable bottom data are

p	ranks	costs
0	(1, 2, 3, 5, 8)	(2, 1, 1, 0, 0)
1	(1, 2, 4, 6)	(2, 1, 1, 0).

They stabilize already at $n = 3$, as in the odd-cylinder proof. Reflection now exchanges the colors. Counting the top slice gives the exact physical prefix profiles, zero at zero,

$$f_p(k) = k + U_p(k) - 4 + B_{1-p}(h - k), \quad k > 0.$$

Indeed the top slice has $4 + p$ vertices of source color p , its omitted vertices have the bottom ranks of color $1 - p$, and the origin correction is p ; the resulting constant is $p - (4 + p) = -4$.

Start one cohort in phase zero of this order and inspect its last five rooms each day. The full count follows

$$h \longrightarrow h - 2 \longrightarrow h - 3,$$

returning to phase zero. For every phase-zero count $10 \leq k \leq h - 3$, the next two days are $k \rightarrow k - 1 \rightarrow k - 1$. After $h - 12$ such pairs, the count is nine in phase zero. The final four days are

$$9 \longrightarrow 8 \longrightarrow 7 \longrightarrow 5 \longrightarrow 0.$$

All these transitions follow by direct substitution in f_p : U_0 is saturated by rank three and U_1 by rank four, while the opposite tails supply the displayed upper range. Thus the solo duration is $2 + 2(h - 12) + 4 = 2h - 18 = 9n - 18$ for every $h \geq 18$.

This duration is even. The untouched cohort is then in physical phase one. Reflect the order in the longitudinal coordinate for its sweep; reflection exchanges colors, so it starts in virtual phase zero and has the same duration. The first cohort is empty and remains empty. The two physical sweeps take $18n - 36$ days, proving the theorem. The independently reflected second order matters: the same orientation would take one additional day.

The accompanying integer certificate and an independent implementation check the finite boundary assertions and all base inequalities. Separate coordinate-neighborhood replays check the two complete sweeps at $n = 4, 6, 8, 10, 12, 30$. The displayed word induction, three-interval potential argument, and explicit sweep establish the unbounded theorem.

15.7 Six inspections on even $3 \times 3 \times n$ cylinders

Theorem 15.6. *For every even $n \geq 4$,*

$$T_6(P_3 \square P_3 \square P_n) = 6n - 8.$$

Proof. Put $n = 2L$ and $h = 9L$. We reuse the arbitrary-support geometry proved in Section 15.6; those geometric statements do not depend on the inspection budget. For a cohort of size k , let e record whether the preceding survivor had surplus four and size in $[4, h - 8]$. Initially $(k, e) = (h, 0)$. The valid domains are $0 \leq k \leq h$ for $e = 0$, and $8 \leq k \leq h - 4$ for $e = 1$. After p useful inspections, put $s = k - p$. An empty survivor gives $(0, 0)$. Otherwise the next size y satisfies $y \geq g_h(s)$, where g_h is (100). The next bit is one exactly when $4 \leq s \leq h - 8$ and $y = s + 4$. Consecutive bits equal to one are impossible. This relaxes every physical search, including searches whose supports are not global prefixes.

Assign charges $(c(0), \dots, c(6)) = (0, 0, 0, 1, 1, 2, 2)$. They satisfy $c(p) + c(q) \leq 2$ whenever $p + q \leq 6$. Define

$$\begin{aligned} (F_0(0), \dots, F_0(8)) &= (0, 0, 0, 1, 1, 2, 2, 3, 4), \\ F_0(k) &= \min \{ \lfloor 4k/3 \rfloor - 7, 4h/3 - 8 \} & (k \geq 9), \\ F_1(k) &= \lfloor (4k + 2)/3 \rfloor - 7 & (8 \leq k \leq h - 4). \end{aligned}$$

Both functions are nondecreasing. Every allowed transition satisfies

$$F_e(k) \leq c(p) + F_{e'}(y). \tag{102}$$

Here is a finite verification with an explicit extension to all lengths. For $h = 18, 27$, enumerate valid (k, e) , $0 \leq p \leq \min(k, 6)$, and every $g_h(k - p) \leq y \leq h$. Set $e' = 1$ precisely in the critical case above, reject $e = e' = 1$, and check (102); when $p = k$, check just $(y, e') = (0, 0)$. These are respectively 1,095 and 3,174 integer inequalities, all satisfied. This specification and the formulas fully determine the finite certificate; the companion artifact includes an exact enumerator.

For $h = 27 + \Delta$, where $\Delta \in 9\mathbb{N}$, it suffices by monotonicity to check the least output in each next-bit class. Split the source range:

- If $k \leq 14$, the profile and critical status are unchanged from $h = 27$, and every minimum output is at most 19. The potentials are also unchanged, so the base inequalities apply.
- If $k \geq 20 + \Delta$, subtract Δ from source, survivor, and output. Since $s - \Delta \geq 14$, the lower profile and critical upper endpoint depend only on $h - s$; the full case $s = h$ remains full. The transformed transition is therefore an allowed $h = 27$ transition. Both potentials decrease by $4\Delta/3$, including the full-state cap.
- If $15 \leq k \leq 19 + \Delta$, then $9 \leq s \leq h - 8$. The minimum branches are $0 \rightarrow 1$ with $y = s + 4$, and $0 \rightarrow 0$ or $1 \rightarrow 0$ with $y = s + 5$. The uncapped formula $\lfloor (4k + 2e)/3 \rfloor - 7$ applies throughout. The maximum drops over $k \bmod 3$, for $p = 0, \dots, 6$, are

	0	1	2	3	4	5	6
$0 \rightarrow 1, 1 \rightarrow 0$	-6	-4	-3	-2	0	1	2
$0 \rightarrow 0$	-6	-5	-4	-2	-1	0	2

Every entry is at most $c(p)$.

This proves (102) for every even length. The two initial potentials sum to $2(4h/3 - 8)$ and both vanish at capture. Their combined daily drop is at most two, giving $T_6 \geq 4h/3 - 8 = 6n - 8$.

For the upper bound, the first two moves depart slightly from a global prefix sweep. Start with the phase-one cohort, whose full slice counts are $(4, 5)^L$. Retain the transverse slice prefixes with counts

$$R_1 = ((4, 5)^{L-1}, 2, 1), \quad R_2 = ((5, 4)^{L-1}, 1, 0)$$

on the first two days. Their successive neighborhoods, calculated from the transverse profiles, are

$$((5, 4)^{L-1}, 5, 2), \quad A = ((4, 5)^{L-1}, 4, 1).$$

Each move inspects six rooms, and $|A| = h - 4$. The global weightlex prefix of size $h - 10$ is contained in A . Indeed, its only missing rooms are four final-slice cells of coordinate weight at least $n + 1$. Reflection in all three coordinates maps rooms of weight at least $n + 1$ to the seven phase-zero rooms of weight at most two. Thus all four missing cells lie among the seven largest-weight rooms and are omitted by this prefix.

Now repeatedly retain the global prefix of size $\max(k - 6, 0)$. The compatibility and exact prefix profiles established in the preceding subsection justify every later move. Starting in phase one, each pair with $k \in \{17, 20, \dots, h - 4\}$ gives $k \rightarrow k - 1 \rightarrow k - 3$. After $(h - 18)/3$ pairs the count is 14, and the final six days are

$$14 \longrightarrow 13 \longrightarrow 11 \longrightarrow 10 \longrightarrow 8 \longrightarrow 6 \longrightarrow 0.$$

One cohort therefore takes $2 + 2(h - 18)/3 + 6 = 3n - 4$ days, an even number. The other cohort stays full during these moves. Reflecting the long coordinate converts its phase zero into the virtual starting phase one, so the same construction clears it in a further $3n - 4$ days. $\square$

The finite arithmetic certificate and its all-length extension are ordinary proofs. This additional even-cylinder result is not presently a complete physical Lean theorem.

15.8 A uniform height strategy for cylinders

Theorem 15.7. *If H is any finite bipartite graph with $A \geq 1$ vertices and $n \geq 2$, then*

$$h(H \square P_n) \leq \lfloor A/2 \rfloor + 1.$$

For $m > A/2$, at most $2 \lceil A(n + 2)/(2m - A) \rceil$ days suffice.

Proof. We first clear one initial-color cohort. Write the path coordinate as $0 \leq z < n$. Bound the possible positions in each fiber by a cutoff a_v , of the current required parity, such that $|a_u - a_v| = 1$ whenever uv is an edge of H . Initially take $a_v = n - 1$ or $n - 2$ according to its parity; this bounds the full cohort. Cutoffs may later be negative or exceed the board.

At a local maximum v , lower a_v by two. Record the room at the old height if it lies on the board. Every adjacent difference stays one. A sequence of m such virtual operations records at most m distinct rooms: repeated operations in a fiber use strictly decreasing heights. Inspect those rooms simultaneously. The survivors lie below the new cutoffs a'_v . After movement their cutoffs are bounded by $a'_v + 1$: a path move increases height by at most one, and an H -move arrives from u with $a'_u \leq a'_v + 1$.

Choose the virtual operations in a fixed cyclic order. List first the initially higher color class of H , then the other class, and repeat this A -letter word. Every operation is at a local maximum;

after a whole color class has been lowered, the other is higher. Adding one to all cutoffs after a day changes no comparison. Isolated vertices of H cause no difficulty.

After t days every vertex has been lowered at least $\lfloor mt/A \rfloor$ times, so

$$\max_v a_v(t) \leq n - 1 + t - 2 \lfloor mt/A \rfloor < n + 1 - t(2m - A)/A.$$

The displayed number of days per cohort makes every cutoff negative. Since $H \square P_n$ has no isolated vertices, an empty post-movement belief certifies capture during inspection. Repeat for the other initial cohort, starting with the full bound of its then-current parity. $\square$

Theorem 15.8 (Exact feasibility on sufficiently long boxes). *Let H be a finite bipartite graph with $A \geq 2$ vertices and a Hamiltonian path. If $n \geq 2 \lfloor A/2 \rfloor$, then*

$$h(H \square P_n) = \lfloor A/2 \rfloor + 1.$$

In particular, this applies to every Cartesian box used as the transverse graph, with arbitrary even or odd side lengths.

Proof. The Hamiltonian path supplies a spanning subgraph P_A of H , so $P_A \square P_n$ is a subgraph of $H \square P_n$. An evader can restrict its moves to this subgraph; the subgraph need not be induced. The classical rectangle theorem [1, Theorem 2] therefore gives

$$h(H \square P_n) \geq \lfloor \min(A, n)/2 \rfloor + 1 = \lfloor A/2 \rfloor + 1.$$

Theorem 15.7 gives the matching upper bound. A Cartesian box has a Hamiltonian path by the usual snake construction: traverse successive slices alternately forwards and backwards along a Hamiltonian path in the lower-dimensional box. The joining endpoints differ only in the new coordinate. $\square$

Thus, for example, $4 \times 4 \times n$ requires exactly nine probes for every $n \geq 16$, and $3 \times 4 \times n$ requires exactly seven for every $n \geq 12$. The length threshold is sufficient; no claim is made that it is the first length attaining the eventual budget. This theorem settles feasibility in an unbounded family containing all side parities, without asserting an optimal-time formula.

Corollary 15.9. *For every $n \geq 3$, $h(P_3 \square P_3 \square P_n) = 5$. The remaining values are two at $n = 1$ and four at $n = 2$.*

Proof. For $n \geq 3$ the board contains the three-cube. An evader may choose to stay in that subgraph, so its hunting number is at least five. Theorem 15.7 gives the upper bound with $A = 9$. At $n = 1$ use the rectangle theorem; at $n = 2$ use Theorem 15.1 with the 3×3 base. $\square$

For a general box this construction can be applied along a longest axis, giving $\lfloor (\prod_{i \neq j} n_i)/2 \rfloor + 1$ as a sufficient budget. It can be loose. The five-inspection time of every $3 \times 3 \times n$ cylinder is determined above; the full classification at larger budgets remains open here.

15.9 The optimal growth rate for every transverse box

The exact offsets remain sensitive to geometry, but the leading term of the optimal time has a uniform answer across all side parities.

Theorem 15.10. *Let H be a finite bipartite graph with $A \geq 2$ vertices and a Hamiltonian path. For each fixed integer $m > A/2$, as $n \rightarrow \infty$ through either parity,*

$$T_m(H \square P_n) = \frac{2An}{2m - A} + O_{A,m}(1). \quad (103)$$

The height strategy in Theorem 15.7 has a bounded additive excess over the optimal time. When A is even, the following bounds hold explicitly for every $n \geq 2$:

$$\frac{2(An/2 - A^2 - m)_+}{m - A/2} \leq T_m(H \square P_n) \leq 2 \left\lceil \frac{A(n+2)}{2m - A} \right\rceil. \quad (104)$$

In particular the theorem applies to every Cartesian transverse box.

Lemma 15.11 (An interior expansion bound for even-width rectangles). *On $P_{2b} \square P_n$, $b \geq 1, n \geq 2$, put $h = bn$ and $K = 2b^2$. Every one-color set R with $K < |R| < h - K$ satisfies $|N(R)| \geq |R| + b$.*

Proof. Match consecutive short-axis rows in pairs. Contract an edge followed by the inverse matching to obtain a directed $b \times n$ grid D . It has loops, bidirectional horizontal edges, and vertical rungs directed one way in even columns and the other way in odd columns. The opposite physical color reverses all rung directions. In either case

$$|N(R)| - |R| = |N_D^+(R) \setminus R|.$$

Suppose the outside boundary $B = N_D^+(R) \setminus R$ has size $s < b$. In $D - B$, at least one entire horizontal row is undeleted. Its vertices lie in a single strongly connected component C .

Call a column bad if it contains a vertex of B . Every adjacent pair of good columns is strongly connected: the horizontal edges go both ways, and the opposite rung directions let a walk go both ways between adjacent rows. The pair meets the undeleted row, hence belongs to C . Only bad columns and isolated good columns can contain vertices outside C . There are at most s of the former and $s + 1$ of the latter, so

$$|V(D) \setminus C| \leq b(2s + 1) < 2b^2 = K.$$

This count includes deleted vertices and also covers the case with no adjacent good columns. The set R is outgoing-closed in $D - B$. It therefore either avoids C or contains C entirely. In the first case $|R| \leq K$; in the second $h - |R| \leq K$, proving the contrapositive. $\square$

Proof of Theorem 15.10. First let $A = 2b$ and consider the spanning rectangle $P_A \square P_n$. Put $C = (h - 2K - m)_+$, $L = K + m$, and

$$\psi(k) = \min\{C, (k - L)_+\}.$$

If a cohort of size k receives $p \leq m$ useful inspections, write $r = k - p$ and $k' = |N(R)|$ for its survivor size and actual next size. The matching gives $k' \geq r$ for every survivor set. If $r \leq K$, the source potential is zero. If $r \geq h - K$, the target potential is C whenever $C > 0$, while $C = 0$ is trivial. In the remaining case Lemma 15.11 gives $k' \geq k - p + b$. Since clipping is monotone and has Lipschitz constant one, all cases give

$$\psi(k) - \psi(k') \leq \max\{0, p - b\}.$$

The two cohorts' useful allocations sum to at most m , so their combined potential decreases by at most $m - b$. Their initial sum is $2C$ and their final sum zero. Thus $T_m \geq 2C/(m - b)$ on the rectangle. The rectangle is a subgraph of $H \square P_n$, giving the same lower bound there. Theorem 15.7 gives the displayed upper bound, with the same coefficient of n .

For odd $A \geq 3$, take the largest odd $\ell \leq n$. The spanning rectangle $P_A \square P_\ell$ has, by the fixed-budget odd-rectangle theorem, an eventual affine period with length increment $2(2m - A)/\gcd(A, 2m - A)$ and time increment $4A/\gcd(A, 2m - A)$. Its time is consequently $2A\ell/(2m - A) + O_{A,m}(1)$. Since $n - \ell \leq 1$, this gives the required lower bound on the whole cylinder. The same height upper bound completes the proof. $\square$

For example, at budgets nine and ten respectively, $T_9(P_4^2 \square P_n) = 16n + O(1)$ and $T_{10}(P_4^2 \square P_n) = 8n + O(1)$. At seven probes, $T_7(P_3 \square P_4 \square P_n) = 12n + O(1)$. Here the cross-section and budget are fixed while the longitudinal side grows. The next theorem strengthens this leading-term result to an exact eventual affine period for every fixed transverse box. Exact constant terms for all short boxes remain a separate question.

16 Eventual affine periods for every fixed transverse box

The optimal growth rate in Theorem 15.10 can be strengthened to an exact eventual recurrence for every fixed transverse box, without a parity restriction on its remaining side.

Theorem 16.1 (Every fixed transverse box). *Let Q be a fixed Cartesian product of finite paths, with $A = |Q| \geq 1$. For each fixed integer $m > A/2$, put*

$$D = 2m - A, \quad g = \gcd(A, D), \quad \Delta = 2D/g.$$

There is an effectively specified positive integer N , bounded by a polynomial in A and m , such that

$$T_m(Q \square P_{n+\Delta}) = T_m(Q \square P_n) + 4A/g, \quad n \geq N. \quad (105)$$

Both longitudinal parities are covered, and Δ is an even integer. For each integer $0 \leq m \leq A/2$, all sufficiently long cylinders are impossible to search with that budget.

We prove the time statement first for even-order Hamiltonian cross-sections, then for odd-order Hamiltonian cross-sections at even longitudinal lengths. The earlier odd-box theorem closes the remaining case for boxes. The argument applies to arbitrary strategies: optimality forces all but a bounded number of days to move a narrow physical boundary. A common interval untouched by exceptional behavior lets us shorten or lengthen the two boundary passages through it.

Theorem 16.2. *Let Q be a finite bipartite graph with a Hamiltonian path and $2b$ vertices, where $b \geq 1$. Fix an integer $m > b$, and put $d = m - b$. There is an effective positive integer N such that, with the even period $\Delta = 2d/\gcd(b, d)$,*

$$T_m(Q \square P_{n+\Delta}) = T_m(Q \square P_n) + \frac{2b\Delta}{d}, \quad n \geq N. \quad (106)$$

The assertion holds for both parities of n . Thus the optimal time is eventually affine on every residue class modulo Δ . The displayed period is valid but need not be the smallest; the effective threshold below is deliberately conservative.

Fix a Hamiltonian order on Q and match consecutive vertices in that order. Each checkerboard class in the cylinder identifies with a $b \times n$ array of matched pairs. Write $h = bn$. As in Lemma 15.11, following an edge by the inverse matching gives a directed graph containing loops, bidirectional edges along every longitudinal row, and alternating vertical rungs between consecutive rows. Extra edges of Q add directed edges inside columns. For a one-color support R , its ordinary neighborhood, expressed using the opposite-color matching, is $R \cup B$, where B is the directed outside boundary. In particular $|N(R)| - |R| = |B|$.

16.1 Equality in the interior expansion bound

Lemma 16.3 (A finite family of minimum-surplus frontiers). *Put $K = 4b^2$. Suppose a one-color set R in $Q \square P_n$ satisfies*

$$K < |R| < h - K, \quad |N(R)| = |R| + b.$$

In the matched array, its rows are either all prefixes or all suffixes: there are integers $1 \leq c_i \leq n-2$, $0 \leq i < b$, such that either

$$R_i = \{0, \dots, c_i - 1\} \quad \text{for every } i, \quad \text{or} \quad R_i = \{c_i + 1, \dots, n - 1\} \quad \text{for every } i.$$

Adjacent cuts satisfy $|c_{i+1} - c_i| \leq 2$. If their difference is two, the rung at the sole intermediate column points from outside R into R . The outside boundary consists exactly of the b cut points. Additional edges of Q may exclude some such patterns, but introduce no additional forms.

Proof. First use only the spanning Hamiltonian rectangle. Its boundary has size at most b , and Lemma 15.11 shows that its size is at least b . Hence its boundary B has exactly b points and also equals the boundary for the whole cylinder.

If one row misses B , the strongly connected component argument in Lemma 15.11 still applies with b deleted points. One component contains all but at most $b(2b+1) \leq K$ vertices, including the deleted points in the exceptional count. The outgoing-closed set R either contains or avoids that component, contrary to the two strict size bounds. Thus every row contains exactly one boundary point, say c_i . Horizontal bidirectionality shows that a row of R is empty, a prefix, a suffix, or the entire row except its boundary point.

For two adjacent rows, the difference of their row sets is a union of at most two integer intervals. At a column where a rung points from the first row to the second, such a difference can contain only the second row's boundary point. It therefore has at most one point of that column parity. A union of at most two intervals with this property has at most four points. Applying this in both directions gives

$$|R_i \setminus R_{i+1}| \leq 4, \quad |R_{i+1} \setminus R_i| \leq 4.$$

An empty row would now imply $|R| \leq 4b^2$, by propagating the bound through all rows. A row full except at its boundary similarly implies $h - |R| \leq b(1 + 4(b-1)) \leq 4b^2$. Both are excluded.

Suppose an adjacent prefix and suffix have lengths u and v . Their two differences have sizes $\min(u, n-v)$ and $\min(v, n-u)$, both at most four. The hypotheses give $n > 8$. If either length exceeds four, both lengths are at least $n-4$; otherwise both are at most four. Propagation through all rows again makes either R or its complement have size at most K . Hence all rows have the same orientation, and their cuts are interior.

Every column strictly between two adjacent cuts has one rung endpoint inside R and its other endpoint outside both R and B . Its rung must therefore point into R . Alternating directions

permit at most one such column, proving the cut bound and parity condition. Conversely, for the spanning rectangle these conditions give exactly the indicated boundary: horizontal edges reach every cut, and no rung reaches another outside point. Extra transverse edges impose only further restrictions on these same bounded-width patterns. $\square$

The cuts have total range at most $2(b-1)$; for subsequent spatial margins use the larger value $v = 2b + 2$. For a prefix survivor with cuts (c_i) , the next matched belief consists of prefixes of lengths $c_i + 1$. A following prefix survivor with cuts (e_i) is possible precisely when $e_i \leq c_i + 1$, and its useful inspection count is

$$p = \sum_{i=0}^{b-1} (c_i + 1 - e_i). \quad (107)$$

At $p = m$ its sum of cuts decreases by d . The suffix case is reflected and its sum increases by d . An orientation cannot change between two such bulk survivors: the next belief of a prefix survivor omits the last column in every row, whereas every nonempty suffix contains it.

16.2 Only boundedly many days can behave differently

Consider an optimal strategy on a sufficiently long cylinder, and set

$$C = h - 2K - m > d, \quad \psi(k) = \min\{C, (k - K - m)_+\}.$$

The argument proving Theorem 15.10, with the larger cutoff $K = 4b^2$, gives for a cohort receiving p useful inspections and changing from count k to k' ,

$$\psi(k) - \psi(k') \leq \max\{0, p - b\}. \quad (108)$$

The two cohorts start with total potential $2C$ and end with zero. Their combined daily decrease is at most d . The height upper bound $T_m \leq 2 \lceil b(n+2)/d \rceil$ consequently gives

$$dT_m - 2C \leq B, \quad B = 4K + 4m + 2b - 2. \quad (109)$$

Indeed the integer ceiling bound gives $dT_m \leq 2b(n+2) + 2d - 2$.

Call a day *efficient* when its total potential loss is exactly d . Every other day contributes at least one to the nonnegative integer slack in (109), so at most B days are inefficient. Potential increases contribute additional slack and cause no exception. Equality forces all m useful inspections to be assigned to one cohort, called the day's *owner*: any nontrivial split makes the sum of the two positive-part bounds in (108) strictly less than d .

On an efficient day the owner loses potential d , while its untouched mate has constant potential. The owner's source potential is positive, so its survivor count r exceeds K . Its target potential is below C , so its next count, and therefore r , is below $h - K$. The interior expansion bound gives raw count loss at most $m - b = d$. Clipping cannot increase a positive loss; equality forces exactly m useful inspections, raw count loss d , and survivor surplus exactly b . Thus every efficient owner survivor has the form in Lemma 16.3.

An untouched cohort with potential strictly between zero and C cannot have constant potential: its actual count lies in the interior range and its neighborhood grows. The mate of an efficient owner therefore has potential zero or C . A consecutive run of efficient days has at most two owner blocks. Once a cohort has owned a day its potential is below C ; if ownership switches, that cohort can

remain unchanged only at zero, from which it cannot subsequently lose another d in the run. There are consequently at most $2(B + 1)$ owner blocks.

The matching-contracted graph is strongly connected for $n \geq 2$: neighboring pairs of columns allow movement in both rung directions, and rows are bidirectional. Every proper nonempty untouched support therefore grows by at least one room per day. Put $L = K + m + 1$. An untouched nonempty mate with zero potential cannot stay in that collar for L days, while an untouched mate with potential C has deficit at most K and becomes full within K days. Mark every inefficient day and the first L days of every owner block as *nonclean*. Their total number is at most

$$E = B + 2L(B + 1). \quad (110)$$

On each remaining *clean* day the mate is globally full or empty. A maximal clean block has one owner, a fixed frontier orientation, and both its owner beliefs and survivors have bounded-width descriptions: its first belief is already the neighborhood of the preceding efficient survivor. By (107), its average cut moves at the constant speed d/b toward clearance. Individual cuts need not move monotonically.

16.3 A physical interval protected from all exceptional behavior

At the start of every consecutive block of nonclean days, mark the preceding clean frontier's b cut columns, if there is one. Initially the board is full, so no cuts need be marked. Also mark every actual inspection column on every nonclean day, and the two board ends. There are at most

$$b(E + 1) + mE + 2 \leq F, \quad F = (b + m)(E + 1) + 2$$

marked columns. Enlarge each mark by radius $\rho = E + 2b + m + 2$. Outside these intervals, each cohort is locally full or empty throughout every nonclean block. At its start, all cuts are far away. During its at most E days there are no nearby inspections, and movement propagates information by at most one longitudinal column per day. An empty region thus stays empty in the protected interior. A full region stays full using the matching inside each column. This proves the assertion for all intermediate states as well as the endpoints, even if remote shots create arbitrary holes elsewhere.

For any prescribed W , the inequality

$$n > F(2\rho + 1) + (F + 1)(W + 2v + 4) \quad (111)$$

provides an unmarked interval of length at least $W + 2v + 4$. Trim $v + 2$ columns from each end and call the retained interval J ; it has length at least W .

A clean block has its endpoint frontiers outside the untrimmed interval: the preceding nonclean block ends uniformly there, and the cuts preceding the following nonclean block were marked. Its average frontier moves strictly toward clearance, with cut range at most v . Consequently it can change its owner's status on J only from full to empty, through a single complete passage. If its endpoint statuses agree, trimming removes any partial excursion onto J . It cannot start empty and end full, since that would move the average across the interval in the wrong direction. Nonclean blocks preserve the local status. Since each cohort starts full on J and ends empty there, each has exactly one clean crossing. These crossings occur at disjoint times; during either one the other cohort is globally full or empty.

16.4 Connecting physical cutoff frontiers

A direct height descent connects any compatible endpoints once the duration exceeds a short bound. This supplies the replacement segments needed for the exact period, without enumerating frontier states.

Put $A = 2b$ and $D = 2m - A = 2d$. For a left frontier, let a_v be the last occupied physical longitudinal coordinate in the fiber over $v \in Q$. The two cutoffs in transverse matched pair i are $c_i - 1$ and $c_i - 2$, in the order required by parity. Its actual next neighborhood advances both by one. A transverse edge uv consequently forces $a_u \leq a_v + 1$, and reversing the edge gives the opposite inequality. Their parities differ, so

$$|a_u - a_v| = 1 \quad \text{on every edge } uv \text{ of } Q. \quad (112)$$

This includes every extra transverse edge. Inside the protected crossing, all fiber cutoffs are far from both longitudinal ends. Conversely, any such interior prefix configuration satisfying (112) has neighborhood cutoffs exactly $a_v + 1$: longitudinal movement supplies them and transverse edges supply nothing larger. Right frontiers use reflected coordinates.

Lemma 16.4 (Height endpoints). *Let Q be connected and bipartite, with $A \geq 2$ vertices and diameter R . Fix $m > A/2$ and put $D = 2m - A$. Let a, b be integer height configurations satisfying (112). If*

$$t \geq \lceil AR/m \rceil, \quad b_v \equiv a_v + t \pmod{2} \quad \text{for every } v, \quad \sum_v a_v - \sum_v b_v = Dt,$$

then t global additions of one, each followed by m legal local maximum flips, take a to b . At completed days the mean height moves steadily between the endpoint means. All intermediate cutoffs, including partial inspection steps, lie in

$$[\mu(b) - R, \mu(a) + R + 1], \quad \mu(x) = A^{-1} \sum_v x_v.$$

Whenever this interval lies inside the cylinder's interior, these operations are actual movements and m useful inspections per day.

Proof. First suppose two configurations a, c have the same vertexwise parity and satisfy $c \leq a$. If the current vector x is not c , choose v of maximum current height among vertices with $x_v > c_v$. A higher neighbor w must also be above its target: otherwise $x_w = c_w = x_v + 1$, while $c_v \leq x_v - 2$, contradicting the target edge condition. By maximality no such neighbor exists. All neighbors have height $x_v - 1$, so lowering x_v by two is a legal local maximum flip. It preserves $x \geq c$ and reduces the sum of gaps by exactly two. Repeatedly doing this reaches c in exactly $\frac{1}{2} \sum_v (a_v - c_v)$ flips.

For the stated endpoints put $c = b - t\mathbf{1}$. It has the same vertexwise parity as a , and $\sum_v (c_v - a_v) = -2mt$. Both a and c change by at most one along an edge, so

$$(c_v - a_v) - (c_w - a_w) \leq 2 \text{dist}(v, w) \leq 2R.$$

Averaging over w gives $c_v - a_v \leq -2mt/A + 2R \leq 0$. The descent therefore supplies exactly mt flips from a to c . Divide them into t groups of m flips, adding one globally before each group. Uniform additions change no local-maximum test. The resulting endpoint is $c + t\mathbf{1} = b$.

Each completed day lowers the mean by D/A . Every height configuration has range at most R , giving the asserted bound at these days. Movement can increase the upper bound by one;

subsequent flips descend coordinatewise to the next day's survivor. The lower bound during those flips is no smaller than that final lower bound. Inside the indicated physical margin, each flip removes an actual top room. Repeated flips in a fiber remove distinct rooms, so each group uses exactly m useful inspections. $\square$

Set $g = \gcd(A, D)$, $\ell_0 = 2A/g$, and $\Delta = 2D/g$. These are even integers satisfying $A\Delta = D\ell_0$. Suppose a clean survivor path from a to b has length $t \geq \lceil AR/m \rceil + \ell_0$. After deleting Δ columns, a left crossing requires endpoints $a - \Delta \mathbf{1}, b$. Their sum difference is $D(t - \ell_0)$ and their phase difference is $t - \ell_0$, since both shifts are even. Lemma 16.4 supplies a path of exactly that shorter length. For insertion use endpoints $a + \Delta \mathbf{1}, b$ and duration $t + \ell_0$. Reflection handles right frontiers. The mean and range bound controls each replacement inside the protected physical band, independently of the old path's details.

For the present even-order case,

$$\ell_0 = \frac{2b}{\gcd(b, d)}, \quad \Delta = \frac{2d}{\gcd(b, d)}, \quad K_0 = \left\lceil \frac{2b(2b-1)}{m} \right\rceil. \quad (113)$$

Here are conservative explicit margins that guarantee such a path well inside the protected interval. Set

$$U = K_0 + \ell_0 + 4, \\ W = 2(\Delta + v + m + 3) + \left\lceil \frac{d(U+4)}{b} \right\rceil + 2v + 2m + 10. \quad (114)$$

In a complete crossing, take the central segment after its average cut has entered J with margin $\Delta + v + m + 3$, and before it leaves at the opposite margin. Its average advances by $d/b \leq m$ per edge, so entry and exit overshoots lose at most $2m$ of spatial length. The remaining travel exceeds dU/b , providing more than U edges. Every cut is within v of the average, so the entire segment stays more than Δ from the ends of J . Its endpoint frontiers lie on opposite sides of a central interval of Δ columns. These margins also cover the entry and exit edges.

For definiteness, one may take

$$N = 4 + \left\lceil \frac{2K + m + d + 1}{b} \right\rceil + F(2\rho + 1) + (F + 1)(W + 2v + 4) + \Delta. \quad (115)$$

This ensures $C > d$, the protected-gap condition, and the same conditions after contracting by Δ . All constants depend only on b and m .

16.5 Deleting and inserting the common interval

Start with an optimal strategy on a cylinder long enough for the preceding construction. Delete Δ consecutive columns from the center of J . Outside the two central crossing segments, map rooms and inspections by ordinary column deletion. No nonclean inspection is lost. Every nonclean state is uniformly full or empty near the join, so neighborhood formation commutes with the deletion there. Other clean frontiers are entirely on one side of the join, and their transitions are either left fixed or translated as a whole by the even displacement Δ .

Within each central crossing, replace its t -edge path by the $(t - \ell_0)$ -edge path in Lemma 16.4. For a prefix frontier the new initial cutoffs are Δ lower and the final cutoffs are the old ones, matching the natural column map at both endpoints. For a suffix frontier use the reflected construction. The

mean of the replacement decreases steadily between the mapped endpoint means, and its range is at most $2b - 1$. The chosen margins therefore keep every intermediate cutoff inside the shorter board. Every new daily edge is an actual movement followed by at most m inspections.

The mate is globally full or empty during this operation. Deleting an even number of days and columns preserves its phase and state. The two crossings are disjoint in time, so their edits do not interfere. An edge is one subsequent inspection day, hence exactly ℓ_0 days have been removed in each crossing. All remaining inspections respect the original budget. We have proved

$$T_m(Q \square P_{n-\Delta}) \leq T_m(Q \square P_n) - 2\ell_0. \quad (116)$$

For the reverse inequality, prepare an optimal strategy on the shorter, still sufficiently long cylinder and insert Δ columns in its protected interval. Extend the locally full or empty regions across the inserted band and translate the far-side rooms. Within each crossing, use Lemma 16.4 to replace its t days by $t + \ell_0$ days. A prefix frontier starts shifted right by Δ and ends at the old final cutoff; reflect this for a suffix. The same local transition and margin arguments apply. The full or empty mate returns to its phase, and the total added time is $2\ell_0$. Therefore

$$T_m(Q \square P_n) \leq T_m(Q \square P_{n-\Delta}) + 2\ell_0. \quad (117)$$

Combining (116) and (117), and using $2\ell_0 = 2b\Delta/d$, proves Theorem 16.2 for every $n \geq N$ after renaming the shorter length n . The even displacement preserves either longitudinal parity throughout.

16.6 Odd-order cross-sections: geometry in paired columns

When the cross-section has odd order, there is no transverse perfect matching. At even longitudinal lengths we instead match consecutive columns. This changes the smallest bulk expansion from a constant cost to two alternating costs. A single bit recording the previous expansion will recover constant progress.

Theorem 16.5. *Let Q be a finite bipartite graph with a Hamiltonian path and $A = 2b + 1 \geq 3$ vertices. Fix an integer $m \geq b + 1$ and put $D = 2m - A$. Put $g = \gcd(A, D)$ and $\Delta = 2D/g$. There is an effective positive integer N such that*

$$T_m(Q \square P_{n+\Delta}) = T_m(Q \square P_n) + 4A/g$$

for every even $n \geq N$.

Write $n = 2L$ and let X, Y be the bipartition of Q , with $|X| = b + 1$ and $|Y| = b$. A Hamiltonian path ensures these color sizes. Match columns $2j$ and $2j + 1$. Each cohort now identifies with an $A \times L$ array of vertices (v, j) , one for each $v \in V(Q)$ and paired-column index j . The contracted graph has loops, bidirectional Q -edges within every layer, and opposite one-way longitudinal flows on the two Q colors. Choose phase zero so that X fibers flow left and Y fibers flow right; phase one reverses both flows. There is no dependence on j . The ordinary neighborhood surplus is again the cardinality of the directed outside boundary. Throughout the paired-column argument put $h = AL$ and $K = A^2$.

Lemma 16.6 (Minimum surplus in paired columns). *Every survivor R with $K < |R| < h - K$ has surplus at least b . If its surplus equals b , then in phase zero all fibers are left prefixes. Across each Q -edge xy , $x \in X, y \in Y$, the X prefix has the same length as the Y prefix or is one position longer.*

In phase one the reflected statement holds, with right suffixes. Every fiber has at least two occupied and two unoccupied positions. Two consecutive middle survivors in an actual strategy cannot both have surplus b .

Proof. Let B be the outside boundary, of size $s \leq b$. If one X fiber and one Y fiber both avoid B , all completely undeleted Q layers lie in one strongly connected component of the graph with B deleted. Each layer is strongly connected, and the two untouched, oppositely directed fibers connect these layers in both longitudinal directions. At most s layers contain a boundary point, so this component contains all but at most $As < K$ vertices, counting the deleted layers among the exceptions. The outgoing-closed set R either contains or avoids the component, contradicting its middle size. There is at least one undeleted layer: $h > 2A^2$ implies $L > 2A > b$.

If $s < b$, both colors have an untouched fiber, a contradiction. If $s = b$, there is still an untouched X fiber, so every Y fiber must contain a boundary point. Thus B consists of exactly one point (y, c_y) on each Y fiber, and no points on X fibers.

In phase zero, the boundary-free X fibers flow left and are prefixes. Bidirectional Q -edges give, for every edge xy ,

$$R_y \subseteq R_x, \quad R_x \setminus R_y \subseteq \{c_y\}.$$

Adjacent cardinalities differ by at most one. Connectivity bounds the global range of fiber sizes by $A - 1$. Both the average size and average complement size exceed A , so each fiber has at least two occupied and two unoccupied positions. If $R_x = \{0, \dots, k - 1\}$, the inclusions make R_y either that prefix or that prefix with its boundary point removed. Removing a point strictly below $k - 1$ would leave $k - 1$ occupied; its right-going edge would then make k a second boundary point. Thus

$$R_y = \{0, \dots, c_y - 1\}, \quad k \in \{c_y, c_y + 1\}.$$

This proves the prefix description. Reflection proves the phase-one statement.

The next belief of a phase-zero minimum-surplus frontier leaves its X lengths unchanged and increases each Y length by one. It still omits position $L - 1$ in every fiber. A phase-one minimum-surplus middle survivor would be a nonempty right suffix in every fiber, hence could not be contained in that belief. Reflection handles the other direction. $\square$

Lemma 16.7 (The intervening frontier). *Suppose R is a phase-zero middle survivor of surplus b , and a phase-one middle survivor $S \subseteq N(R)$ has surplus $b + 1$. Then all fibers of S are left prefixes. Across each Q -edge, its Y prefix has the same length as its X prefix or is one position longer. The reflected assertion holds with both orientations reversed.*

Proof. For any support with boundary B , a Q -edge gives $S_u \setminus S_v \subseteq B_v$. Following a simple Q path shows that any two fiber cardinalities differ by at most $|B|$: its target fibers are distinct, so their boundary counts sum to at most $|B|$. At $|B| = b + 1$, an empty fiber would give $|S| \leq A(b + 1) < A^2$, contrary to the middle-size hypothesis. Every fiber is nonempty.

The containment $S \subseteq N(R)$ ensures that all fibers omit $L - 1$. In phase one, each X fiber flows right. If such a fiber had no boundary point, its nonempty set would be right-closed and would contain $L - 1$. Hence each of the $b + 1$ X fibers meets B . They exhaust the boundary, leaving no boundary point on any Y fiber. The latter flow left and are prefixes. Applying the same bidirectional-edge and endpoint argument as in Lemma 16.6, with X, Y exchanged, proves that the X fibers are prefixes as well, with the stated length relation. Reflection gives the other orientation. $\square$

Both frontier families have finitely many relative-length patterns. Their length range is at most $A - 1$, and choosing the zero-or-one length differences along a spanning tree gives at most 2^{A-1} patterns for fixed phase and orientation. Extra edges impose consistency conditions on these patterns.

16.7 A history bit and the paired-cylinder period

For each cohort, define $e_t = 1$ precisely when its preceding day's survivor was in the strict middle range and had surplus b ; otherwise set $e_t = 0$. Initially $e_0 = 0$. If its before-inspection count is k_t , use the raw rank $2k_t + e_t$. Let p be its useful inspection count, so the current survivor has size $r = k_t - p$, surplus s , and next count $k_{t+1} = r + s$. The raw loss is

$$(2k_t + e_t) - (2k_{t+1} + e_{t+1}) = 2p - 2s + e_t - e_{t+1}. \quad (118)$$

For a middle survivor with $s = b$, Lemma 16.6 gives $(e_t, e_{t+1}) = (0, 1)$, so the loss is $2p - A$. If $s \geq b + 1$, the new bit is zero and the old bit at most one, again giving loss at most $2p - A$.

Set

$$L_0 = 2(K + m) + 1, \quad C = 2h - 4K - 2m - 1 > D, \quad \psi_t = \min\{C, (2k_t + e_t - L_0)_+\}.$$

A survivor of size at most K gives source rank at most L_0 , hence source potential zero. A survivor of size at least $h - K$ gives next rank at least $2h - 2K = L_0 + C$, hence target potential C , by the longitudinal matching. Clipping and (118) therefore give, in every case,

$$\psi_t - \psi_{t+1} \leq \max\{0, 2p - A\}. \quad (119)$$

For the two useful allocations $p + q \leq m$, the sum of these charges is at most D , with equality only if all m useful inspections belong to one cohort.

Both initial potentials equal C . The height upper bound $T \leq 2 \lceil A(n + 2)/D \rceil$ and its integer ceiling estimate imply

$$DT - 2C \leq B, \quad B = 8K + 8m + 2A.$$

Thus at most B days have total potential loss below D . On every efficient day the owner has positive source potential and target potential below C , which force its survivor into the strict middle range. Equality in the clipped loss then forces equality in (118). The only cases are

$$(s, e_t, e_{t+1}) = (b, 0, 1) \quad \text{or} \quad (b + 1, 1, 0). \quad (120)$$

The first survivor is a frontier by Lemma 16.6. In the second case the preceding survivor was a middle minimum-surplus frontier, so the actual containment in its neighborhood permits Lemma 16.7. Both efficient phases therefore have physical frontiers.

The paired contraction is strongly connected for $L \geq 2$. An untouched proper nonempty cohort grows by at least one room each day, while its history bit can drop by at most one. Its raw rank therefore increases by at least one, so its potential cannot remain constant strictly between zero and C . The owner-block argument in Section 16.2 applies unchanged. There are at most $2(B + 1)$ owner blocks. The zero-potential collar has at most $K + m$ possible rooms and the full-potential collar has deficit at most K . Marking the first $L_1 = K + m + 1$ days of every owner block, as well as every inefficient day, leaves at most

$$E = B + 2L_1(B + 1) \quad (121)$$

nonclean days. On each clean block the mate is globally full or empty, with bit zero, and the owner's surpluses alternate $b, b + 1$. The frontiers have a fixed orientation throughout this block: the bridge preserves orientation, and the next minimum-surplus phase has the same orientation as the preceding minimum-surplus phase.

There is a timing distinction when the graph vertices are survivors. At such a vertex write e for the *newly set* bit, referring to that survivor itself. If consecutive survivor sizes are r, r' and the first has surplus s , then $r' = r + s - m$. Their constant-drift rank is $2r - e$, rather than the before-inspection rank $2k + e$:

$$(2r - e) - (2r' - e') = D \quad (122)$$

in both alternating cases $(s, e, e') = (b, 1, 0)$ and $(b + 1, 0, 1)$. The ordinary average occupied length is nonincreasing and may be stationary for one round when $m = b + 1$. Its corrected version $(2r - e)/(2A)$ decreases by exactly $D/(2A)$ per edge and differs from it by at most $1/(2A)$. Use the enlarged margin $v = A + 2$ to cover this correction, the fiber range, and the direct height replacements.

These frontiers are also physical height configurations. If l_v is the occupied paired-column length and $\sigma_v \in \{0, 1\}$ its physical longitudinal parity, then

$$a_v = 2(l_v - 1) + \sigma_v.$$

In either phase, the boundary-free color has parity zero and its length equals its neighbor's length or exceeds it by one. Hence $|a_u - a_v| = 1$ across every Q -edge. The bridge lemma ensures this on the intervening phase as well. For a left frontier the sum of the parity indicators is $b + 1 - e$, so the sum of physical heights differs from $2r - e$ by a constant. The drift in (122) is therefore the same physical height drift used in Lemma 16.4.

Apply that lemma with $g = \gcd(A, D)$ and the conservative diameter bound $A - 1$. We may use

$$K_0 = \left\lceil \frac{A(A - 1)}{m} \right\rceil, \quad \ell_0 = 2A/g, \quad \delta = D/g, \quad \Delta = 2\delta. \quad (123)$$

A crossing of length at least $K_0 + \ell_0$ can be replaced by one shorter by ℓ_0 days, shifting the initial endpoint by δ paired columns, or by one longer by ℓ_0 with the opposite shift. The construction connects actual physical height endpoints, without an additional assumption about intermediate survivor shapes.

We spell out the changes to the protected-band argument to ensure that this is a physical reduction. Mark the paired-column index of every nonclean inspection and the initial frontier cuts of every nonclean block. Together with the ends there are at most

$$F = (A + m)(E + 1) + 2$$

marks. A physical move crosses at most one pair boundary, so radius $\rho = E + 2A + m + 2$ shields all nonclean blocks. Local full preservation uses the longitudinal matching *inside each pair*; it requires no perfect matching in the odd-order graph Q . Local empty preservation follows from finite propagation. Thus the argument of Section 16.3 applies in the paired coordinate. Choose a raw gap of length $W + 2v + 4$ and trim $v + 2$ at each end. The corrected mean, its nonzero speed, and the margin v give one complete clean crossing per cohort through the retained band. Their mates are globally full or empty and the crossings are disjoint in time.

Explicit constants are obtained by putting

$$\begin{aligned}
U &= K_0 + \ell_0 + 4, \\
W &= 2(\delta + v + m + 3) + \left\lceil \frac{D(U + 4)}{2A} \right\rceil + 2v + 2m + 10, \\
L_* &= 4 + \left\lceil \frac{4K + 2m + 1 + D}{2A} \right\rceil + F(2\rho + 1) + (F + 1)(W + 2v + 4) + \delta,
\end{aligned} \tag{124}$$

and taking $N = 2L_*$. The corrected mean moves by $D/(2A) \leq m$ per edge, so the entry and exit overshoots cost at most $2m$ paired positions. The stated W leaves at least U edges in a central path, with its endpoints on opposite sides of a δ -pair band and all intermediate cuts more than δ from the outer edges. The threshold ensures $C > D$, the raw gap, and the same conditions after contraction.

Delete that central band of δ pairs and replace each central crossing by the walk shorter by ℓ_0 supplied by Lemma 16.4. Start a left frontier shifted by $-\delta$ pairs, equivalently $-\Delta$ physical columns; its new endpoint is exactly the old endpoint. Reflect this for a right frontier. Every new edge is a valid physical-height transition, and the constant mean drift and bounded range keep its filled-tail configuration interior. Outside the crossings, neighborhood formation commutes with the natural column map because the protected neighborhood is uniformly full or empty, including every nonclean intermediate state. No nonclean inspection is deleted. The mate's full or empty state is preserved; the removed time is even, and paired-column translation preserves phase. As survivor edges count the following inspection, exactly ℓ_0 days are removed per cohort. This proves

$$T_m(Q \square P_{n-\Delta}) \leq T_m(Q \square P_n) - 2\ell_0.$$

Starting instead from a sufficiently long shorter board, insert the same number of pairs and use the longer endpoint connection to add ℓ_0 days to each crossing. The same filled-tail, endpoint, and mate arguments prove the reverse inequality, exactly as in Section 16.5. Since $2\ell_0 = 2A\Delta/D$, this proves Theorem 16.5 on every sufficiently long even cylinder, with the displayed effective threshold.

16.8 Combining the parity classes

Proof of Theorem 16.1. Every transverse box has a Hamiltonian path by the snake construction used in Theorem 15.8. If $A \geq 2$ is even, Theorem 16.2 already supplies the asserted recurrence for both longitudinal parities. If $A \geq 3$ is odd, every nontrivial transverse side is odd. Theorem 14.3 applies to odd longitudinal lengths, while Theorem 16.5 applies to even lengths. Both give exactly the same even period $\Delta = 2(2m - A)/\gcd(A, 2m - A)$ and the same increment $4A/\gcd(A, 2m - A)$. Taking the larger threshold proves the common recurrence on both parity classes.

If $A = 1$, the cylinder is a path. For $m = 1$ the formula $T_1(P_n) = 2n - 4$, $n \geq 3$, gives period two and increment four. For $m \geq 2$, take $\Delta = 2(2m - 1)$ and $N = m + 1$ in (1). This changes its ceiling by four and preserves both the exceptional congruence and the parity correction, giving the required increment $2\Delta/(2m - 1) = 4$.

For $A \geq 2$ and integer $m \leq A/2$, Theorem 15.8 gives impossibility whenever $n \geq 2 \lfloor A/2 \rfloor$. If $A = 1$, the only such budget is zero; no target on a nonempty path can be captured without an inspection.

For completeness the displayed sufficient time thresholds are polynomial. With $X = A + m + 1$, their definitions give $K, B, L_1 = O(X^2)$, $E = O(X^4)$, $F = O(X^5)$, and $\rho = O(X^4)$. Because $m > A/2$, we have $K_0 \leq 2(A - 1)$ and $\ell_0 \leq 2A$, so $U = O(A)$ and $W = O(X)$. The dominating

threshold term $F\rho$ is $O(X^9)$. For the earlier odd-length box theorem, $S = \sum_i (n_i - 1) \leq A - 1$, hence its corner bound satisfies $Z = A(S + 1) \leq A^2$. Substitution into (87) also gives a polynomial bound, dominated by $O(X^9)$. The path threshold is smaller. Taking the larger parity threshold thus preserves a uniform polynomial sufficient bound. $\square$

At the eventual minimum budget $m = \lfloor A/2 \rfloor + 1$, the period is two. The time increase on adding two columns is $2A$ when A is even and $4A$ when A is odd. The polynomial onset bound is sufficient, not a claim of the earliest length at which this recurrence holds.

The theorem fixes the entire cross-section and the budget while the last side grows. It does not supply the smallest period, the finite exception values, or a simple formula for every arbitrary finite box. The proof covers every transverse box; for a general odd-order Hamiltonian graph the new paired-column theorem only asserts the even-length case. These eventual-period results have ordinary proofs and independent review, rather than complete Lean formalizations.

17 A spatial transfer matrix for a fixed deadline

The preceding reductions retain geometric information about the current belief. A different choice of state removes every parity restriction: fix the number of inspection days, and retain the history of each column. The resulting matrix is independent of the longitudinal length. It gives an exact recurrence and an eventual formula for the inverse tradeoff.

Let Q be any nonempty finite simple graph, put $A = |V(Q)|$, and let $b_T(n)$ be the minimum daily budget guaranteeing capture within T days on $Q \square P_n$, where $n \geq 2$ and $T \geq 1$.

Theorem 17.1. *For fixed Q, T , there are effectively computable integers $n_0, p > 0$ and rational constants $c_0, \dots, c_{p-1}$ such that*

$$b_T(n) = \frac{An}{T} + c_{n \bmod p} \quad (n \geq n_0).$$

The period can be chosen divisible by $T/\gcd(A, T)$. The finitely many remaining values, and a winning strategy at any specified length, are also effectively computable. No bipartiteness assumption on Q is needed.

Here “effectively computable” asserts a terminating finite procedure, not a small uniform complexity bound. The state space below has $2^{2A(T-1)}$ states. In particular, this theorem does not supply a fixed matrix when the deadline itself grows with the box length.

17.1 Exact local certificates

The case $T = 1$ is $b_1(n) = An$. Assume $T \geq 2$, and use the finite alphabet $\mathcal{W} = (2^{V(Q)})^{T-1}$. A letter $b = (b_1, \dots, b_{T-1})$ records proposed survivor sets in one column. Put $b_T = \emptyset$, and let 0 denote the all-empty letter. For $a, b, c \in \mathcal{W}$ define

$$\begin{aligned} w_1(a, b, c) &= A - |b_1|, \\ w_t(a, b, c) &= |(N_Q(b_{t-1}) \cup a_{t-1} \cup c_{t-1}) \setminus b_t| \quad (2 \leq t \leq T). \end{aligned} \tag{125}$$

A word $b^1 \cdots b^n$, with boundary letters $b^0 = b^{n+1} = 0$, has day- t cost $\sum_{j=1}^n w_t(b^{j-1}, b^j, b^{j+1})$.

Lemma 17.2. *A successful T -day strategy with respective daily capacities $m_1, \dots, m_T$ exists if and only if such a word has day- t cost at most m_t for every t .*

Proof. Let R_t be the union of the word's day- t column sets, with $R_T = \emptyset$. Define nominal beliefs $B_1 = V(Q \square P_n)$ and $B_t = N(R_{t-1})$ for $t \geq 2$, and inspect $S_t = B_t \setminus R_t$. The product-neighborhood identity shows that $|S_t|$ is exactly the local cost sum in (125).

The actual belief U_t satisfies $U_t \subseteq B_t$ inductively: its survivors are contained in R_t , so their neighborhood is contained in B_{t+1} . On day T , the empty survivor envelope makes $S_T = B_T$ cover every possibility. This argument intentionally does not require $R_t \subseteq B_t$; extra envelope vertices do not compromise soundness.

Conversely, take the actual survivors of a successful strategy, discard inspections outside the actual belief, and extend an earlier capture by empty inspections and survivors. Its column sections form a word whose costs are exactly the useful inspection counts. $\square$

Use pairs $(a, b) \in \mathcal{W}^2$ as matrix states. The transition $(a, b) \rightarrow (b, c)$ has monomial weight $\prod_{t=1}^T x_t^{w_t(a, b, c)}$. Let $M(\mathbf{x})$ be this matrix; let u indicate states $(0, b)$ and v indicate states $(b, 0)$. There is exactly one n -edge walk for each n -column word, so

$$F_n(\mathbf{x}) = u^\top M(\mathbf{x})^n v$$

has a positive coefficient at $\mathbf{x}^{\mathbf{c}}$ precisely when a certificate with exact cost vector $\mathbf{c}$ exists. Thus

$$\sum_{n \geq 0} F_n(\mathbf{x}) z^n = u^\top (I - zM(\mathbf{x}))^{-1} v$$

is a rational formal series encoding the exact daily-budget tradeoff. All coefficients are nonnegative; multiplicities of envelopes do not affect the existential criterion.

For a direct algorithm, retain the accumulated cost vector and discard it if any coordinate exceeds the proposed budget m . This uses at most $O(n|\mathcal{W}|^3(m+1)^T)$ transitions. Since an optimum never needs $m > An$, searching budgets remains polynomial in n for fixed Q, T . Backtracking recovers the word and the physical inspections in the lemma.

17.2 Why the optimum has an eventual formula

A subset of $\mathbb{N}^r$ is *linear* if it is a translate of a finitely generated additive monoid, and *semilinear* if it is a finite union of linear sets. Add a length coordinate one to each edge weight. The attainable vectors $(n, c_1, \dots, c_T)$ form an effectively semilinear set L . To see this directly, eliminate the finite automaton's states to obtain a regular expression. Its additive image is computed using union, Minkowski sum and additive closure, all of which preserve semilinearity. For the only less immediate operation, if $S = \bigcup_i (b_i + \text{monoid}(P_i))$, then

$$S^* = \bigcup_I \left(\sum_{i \in I} b_i + \text{monoid} \left(\{b_i : i \in I\} \cup \bigcup_{i \in I} P_i \right) \right).$$

Here I ranges over subsets of component indices. Reserve one base for each used component; extra bases account for repeated occurrences, and all periods can be assigned to the reserved occurrence. The empty I supplies zero. This proves both containments and effectivity.

By the effective equivalence of semilinear and Presburger-definable sets [8, Theorems 1.1 and 1.3], the relation

$$E(n, m) \iff n \geq 2 \text{ and } \exists \mathbf{c} ((n, \mathbf{c}) \in L \text{ and } c_t \leq m \text{ for all } t)$$

is effectively Presburger-definable. The same holds for $E(n, m) \wedge (m = 0 \text{ or } \neg E(n, m - 1))$, which is the graph of b_T . Convert that graph to a finite union of linear sets.

Every infinite linear component of a function's graph lies on a rational affine line. Indeed, a nonzero period cannot have zero first coordinate. For two periods $(a_i, d_i), (a_j, d_j)$ with positive first coordinates, using them a_j and a_i times gives the same input, so uniqueness of the output forces $a_j d_i = a_i d_j$. All periods therefore have one slope. The component's input projection is a translated numerical semigroup; it eventually contains exactly one residue class modulo the gcd of its positive generators. A threshold is computable by shortest paths among residues modulo any generator. Taking a common multiple of the finitely many periods and a maximum of their thresholds gives an affine formula on each eventual residue class. Finite components contribute only finitely many exceptions.

It remains to identify the slopes. The longitudinal matching between columns $(1, 2), (3, 4), \dots$ supplies $2A \lfloor n/2 \rfloor$ alternating target trajectories, pairwise disjoint on every inspection day. A probe intercepts at most one of these trajectories on that day. Consequently

$$Tb_T(n) \geq 2A \lfloor n/2 \rfloor \geq A(n - 1).$$

For the upper bound put $q = \lceil n/T \rceil$. When the possible rooms lie in a suffix beginning at column L , inspect its first $q + 1$ columns, clipped to the board. After movement the suffix begins at least at $L + q$. After $T - 1$ days at most $n - (T - 1)q \leq q$ columns remain; inspect all of them on the final day. Thus

$$b_T(n) \leq A(\lceil n/T \rceil + 1) < An/T + 2A.$$

Internal Q -edges never change the column index. The two bounds force every infinite affine component to have slope A/T . Enlarging the period by $T/\gcd(A, T)$ if necessary proves Theorem 17.1, including its effective exception table and integral additive increment.

The survivor-envelope equivalence and the exact product-fiber cost sum are verified in Lean by `SurvivorEnvelope`. The matrix and semilinearity arguments above are ordinary proofs. The reference transfer implementation was compared with the independently proved path and two-row formulas in 56 cases, with physical replay of its recovered inspection sets. Neither these finite comparisons nor the formal semantic lemma are being substituted for the general periodicity proof.

18 Formal verification, reproducibility, and open problems

The formal development uses Lean 4.33.1 and its standard library. It contains no admitted proof, project-specific mathematical axiom, or trusted external solver result. Finite certificates are reduced by the Lean kernel. The dependency reports of the final theorems use only the usual logical principles of propositional extensionality, classical choice, and quotient soundness, where applicable.

Four graph classifications are complete in Lean:

- `PathClassification.can_capture_iff` and `exact_minimum` cover every positive path length and budget, including the one-room and one-probe cases.

- **LadderGame** connects the complete two-counter optimum to actual two-row rooms, daily cardinality budgets, and arbitrary avoiding walks, including the single-edge degeneration.
- **ThreeCubeTimeClassification.can_capture_iff** and **exact_minimum** cover every budget on the $3 \times 3 \times 3$ box. The statement quantifies over unrestricted room-valued inspection schedules. It does not assume a preferred order or support shape.
- **FourCubeClassification.can_capture_iff** and **exact_minimum** give the complete $4 \times 4 \times 4$ answer. The proof constructs the square-pair compression on the actual room graph, proves its normal form, verifies the shape-sensitive eight-probe certificate and all remaining budget bounds, and checks attaining inspection schedules against every legal target walk.

The abstract strategy-comparison theorem is also formalized in **StrategyCompression**. The four-cube development also formalizes product lifting, a decreasing energy, and the concrete four-square operators used in that classification. The general odd-path and arbitrary-square applications currently have ordinary proofs. The three-row and four-row classifications, five-row classifications, arbitrary-dimensional odd-box isoperimetry, general side-two application, and unbounded cylinder construction also currently have ordinary proofs, not complete Lean developments. The scalar potentials in the five-row proofs have separate Lean checks, including row-deficit consequences, the four-probe transition and equality lemmas, and the three finite high-budget bounds. The discounted potential for five or more probes and the entire physical-geometric bridge are not yet formalized. The unbounded $3 \times 3 \times n$ five-inspection formula likewise has an ordinary proof with exact finite certificates, not a complete Lean proof. The seven-row five-inspection and even-length $3 \times 3 \times n$ six-inspection formulas also have ordinary proofs with finite certificates, rather than complete Lean formalizations. Successful checking of the formal files must not be represented as formal verification of those additional theorems.

The generic profile complement identity and the one-unit difference of the maximum neighborhood surpluses for bipartite parts of sizes $M + 1, M$ are fully formalized. The module is **BipartiteProfileDuality**; its inverse arithmetic is in **ProfileInverse**. The profiles are defined as actual attained neighborhood minima over finite subsets; they are not assumed functions. This verifies the structural profile identity used in the odd-box feasibility corollary, while its geometric nesting and the published hunting-number criterion remain ordinary mathematical dependencies. The minimum-budget interval-insertion theorem and its rational certificates also currently have ordinary proofs and independent exact checks, rather than Lean proofs. The uniform minimum-budget scalar inequality is verified in **UniformOddRank**; both eventual-period theorems for cylinders, their geometric decompositions, and the wider odd-rectangle time formula currently have ordinary proofs. The module **SurvivorEnvelope** does verify the fixed-deadline certificate's equivalence to capture of actual walks, and the equivalence between local column costs and the daily budget. The subsequent semilinearity argument remains an ordinary proof.

The possible-position semantics are proved in **BeliefSemantics**: membership is equivalent to the existence of a legal walk avoiding earlier probes. The finite winning recurrence and its greatest losing trap are proved in **CaptureRecurrence**. Formal rounds start at zero, so a last-round index t corresponds to $t + 1$ inspection days. Every public classification has been separately reviewed for this indexing, its physical graph, its global daily budget, and its quantification over all walks.

The companion at <https://angelraychev.com/princess/> provides the presentation, detailed proof text, and formal sources. The research archive includes scripts for independent physical replay, finite lower certificates, exact small-instance searches, and dated research notes. The source-specific verification receipt records the compiler version, source hashes, dependency order, axiom

reports, and exit codes. The complete check runs sequentially with one compiler worker. To reproduce it, install the pinned Lean version and run `python3 src/check_lean.py` from the archive root, optionally giving `-lean` followed by the compiler path. No additional package installation is required.

18.1 What remains open in this investigation

The main unresolved target is a comparably concise account of feasibility and minimum time for arbitrary side lengths in arbitrary dimension. The exact algorithms proved here cover substantial unbounded families; they are not a closed classification of every box. Even where an exact recurrence is established, solving it in a useful explicit form remains a distinct problem. Numerical patterns are recorded as conjectures until the required geometry and all parameter ranges are proved.

The arbitrary-dimensional odd-box theorem is established by a uniform exchange argument, extending the two-dimensional base without inferring an unbounded statement from numerical checks. Boxes with even sides greater than two require further ideas: parity-open nesting already fails on some even rectangles. Also, optimal strategies often finish one initial-color cohort before working substantially on the other, with one shared transition day; an unrestricted exchange theorem that would justify this pattern in every case is not presently asserted.

18.2 Research and writing provenance

The path results originated in the author’s October 2019–March 2020 work. The present shorter proofs were reconstructed from the historical manuscripts; the original work’s credit is not re-assigned to the tool used for reconstruction. The 2026 grid and box investigation, new proof organizations, experiments, formal developments, and this draft were developed with substantial assistance from Astra 6 operating through the Codex harness. The research log distinguishes recovered results, known literature, new derivations with unverified priority, failed approaches, and unresolved claims. The human author retains responsibility for the manuscript and any eventual submission.

Acknowledgments

The author thanks Dimitar Rusev for writing the preliminary section on monotonicity (Section 3) in the 2020 student-conference version of this work. By agreement, his contribution to that version is acknowledged here; both authors remain credited in its bibliographic entry. The author also thanks his mathematics teacher and mentor Dimitar Dimitrov, who encouraged him to pursue mathematical research and guided his early work.

References

- [1] Tatjana V. Abramovskaya, Fedor V. Fomin, Petr A. Golovach, and Michał Pilipczuk. How to hunt an invisible rabbit on a graph. *European Journal of Combinatorics*, 52:12–26, 2016. doi: 10.1016/j.ejc.2015.08.002. URL <https://fedorvf.github.io/articles/2016/2016g.pdf>.
- [2] Nikolay Beluhov and Emil Kolev. Search for a moving target in a graph. *Electronic Notes*

- in *Discrete Mathematics*, 57:39–46, 2017. doi: 10.1016/j.endm.2017.02.008. URL <https://doi.org/10.1016/j.endm.2017.02.008>.
- [3] Walid Ben-Ameur, Harmender Gahlawat, and Alessandro Maddaloni. Hunting a rabbit: Complexity, approximability and some characterizations. *Theoretical Computer Science*, 1075: 115946, 2026. doi: 10.1016/j.tcs.2026.115946. URL <https://arxiv.org/abs/2502.15982>. Preprint first posted in 2025.
 - [4] Sergei L. Bezrukov and Oriol Serra. A local–global principle for vertex-isoperimetric problems. *Discrete Mathematics*, 257(2–3):285–309, 2002. doi: 10.1016/S0012-365X(02)00431-4. URL [https://doi.org/10.1016/S0012-365X\(02\)00431-4](https://doi.org/10.1016/S0012-365X(02)00431-4).
 - [5] Jessalyn Bolkema and Corbin Groothuis. Hunting rabbits on the hypercube. *Discrete Mathematics*, 342(2):360–372, 2019. doi: 10.1016/j.disc.2018.10.011. URL <https://arxiv.org/abs/1701.08726>.
 - [6] John R. Britnell and Mark Wildon. Finding a princess in a palace: A pursuit-evasion problem. *The Electronic Journal of Combinatorics*, 20(1):P25, 2013. doi: 10.37236/2296. URL <https://arxiv.org/abs/1204.5490>.
 - [7] Thomas Dissaux, Foivos Fioravantes, Harmender Gahlawat, and Nicolas Nisse. Further results on the Hunters and Rabbit game through monotonicity. *Information and Computation*, 305: 105302, 2025. doi: 10.1016/j.ic.2025.105302. URL <https://arxiv.org/abs/2309.16533>.
 - [8] Seymour Ginsburg and Edwin H. Spanier. Semigroups, Presburger formulas, and languages. *Pacific Journal of Mathematics*, 16(2):285–296, 1966. doi: 10.2140/pjm.1966.16.285. URL <https://msp.org/pjm/1966/16-2/pjm-v16-n2-p09-s.pdf>.
 - [9] John Haslegrave. An evasion game on a graph. *Discrete Mathematics*, 314:1–5, 2014. doi: 10.1016/j.disc.2013.09.004. URL <https://doi.org/10.1016/j.disc.2013.09.004>.
 - [10] Dmitry Kamenetsky. Entry A301426: Number of steps required in the worst case for three knights to find the princess in a castle with n rooms arranged in a line. The On-Line Encyclopedia of Integer Sequences, March 2018. URL <https://oeis.org/A301426>. Created 21 March 2018; conjectural formula added the same day. Accessed 11 September 2026.
 - [11] Dmitry Kamenetsky. Entry A301337: Number of steps required in the worst case for two knights to find the princess in a castle with n rooms arranged in a line. The On-Line Encyclopedia of Integer Sequences, March 2018. URL <https://oeis.org/A301337>. Created 19 March 2018; conjectural formula recorded the same day. Accessed 11 September 2026.
 - [12] János Körner and Victor K. Wei. Odd and even Hamming spheres also have minimum boundary. *Discrete Mathematics*, 51(2):147–165, 1984. doi: 10.1016/0012-365X(84)90068-2. URL [https://doi.org/10.1016/0012-365X\(84\)90068-2](https://doi.org/10.1016/0012-365X(84)90068-2).
 - [13] Yota Otachi and Ryohei Suda. Bandwidth and pathwidth of three-dimensional grids. *Discrete Mathematics*, 311(10–11):881–887, 2011. doi: 10.1016/j.disc.2011.02.019. URL <https://arxiv.org/abs/1101.0964>.
 - [14] Angel Raychev. Optimal strategies for finding a princess in a linear table. Bulgarian spring-conference manuscript, twentieth Student Section (US’20), 17 pages. Title translated into English, 2020.

- [15] Angel Raychev and Dimitar Rusev. Optimal strategies for finding a princess in a linear graph. Bulgarian student-conference manuscript, twentieth Student Conference (UK'20), 20 pages. Title translated into English, 2020.